



// PROPOSAL · PRELIMINARY RESULTS · MECHANISTIC INTERPRETABILITY

Dictionary learning is all you need

Keira Wiechecki

WAKARU · KAW504@NYU.EDU

Jade Zaslavsky

WAKARU · JADEOFNAMELESS@GMAIL


Lamb Wakefield

FLOCK CAPITAL · LAMB@FLOCK.HOLDINGS

vmfunc

FLOCK CAPITAL · CELESTE@LINUX.COM

October 4, 2026



ORBIT · SEED 1024 · FLOYD-STEINBERG ERROR DIFFUSION

ABSTRACT

Ideally, mechanistic interpretability would provide guarantees about the computational process used by a model: features should be causal, composable, and closed-form. Most existing techniques, such as sparse autoencoders, can perform arbitrarily badly on all three. Arguably the simplest interpretability technique is a linear probe, and the simplest control technique is a steering vector. They are remarkably robust compared to unsupervised techniques such as SAEs, but require labeled contrastive datasets. We propose a MoEification technique that decomposes a representation space into a set of contrastive classifiers. Each class corresponds to a representative vector; inputs are encoded as classifications and reconstructed from the vectors they select. Steering effectively comes for free, outputs are bounded, the representative vectors admit static interpretation in principle, and feature space is structured. Preliminary results on SONAR sentence embeddings and GPT-2 residuals: trained heads are contrast sets; a hard code of 1900 bits with no continuous coefficients reconstructs 40% better than a top-k SAE at matched bits, and on GPT-2 an 800-bit code recovers more of the model's loss than a k=32 SAE. Entries are realized by steps in embedding space, and interventions set on heads compose additively. But an input-space step disturbs other heads as much as a random step of its length, past the first layer no head reproduces across seeds, and parameter-space decodes do not describe what an entry does.

1 INTRODUCTION

Ideally we would like mechanistic interpretability to provide guarantees about the computational process used by a model. We would like features to be

1. causal
2. composable
3. closed-form

However most existing techniques, such as sparse autoencoders (SAEs), produce features that can perform arbitrarily badly on these desiderata. Feature splitting breaks compositionality. They can't usually be statically analyzed (though bilinear SAE variants can be). Sparsity trades off against reconstruction fidelity, limiting the extent to which features can be used for steering.

Arguably the simplest interpretability technique is a linear probe, and the simplest control technique is a steering vector. They are remarkably robust compared to unsupervised techniques such as SAEs, but require labeled contrastive datasets. We propose a MoEification technique based on decomposing a representation space into a set of contrastive classifiers. Each class corresponds to a representative vector. Inputs are encoded as a series of classifications, and reconstructed from the representative vectors those classifications select. This has several advantages over SAEs:

1. It effectively gives steering for free: changing a classification is a steering move.
2. Outputs are bounded.
3. The representative vectors admit static interpretation in principle; making that real is currently an open problem (Section 6.8).
4. Feature space is structured.




FIGURE 1 · The softmax Ontologizer’s steering primitive, measured. Left: forcing an entry moves decodes along the supervised difference-of-means direction $0.43\times$ as far as supervised steering itself and $5\times$ further than probe directions (Table 4). Centre: joint two-head interventions point where the summed single interventions point, exactly in the final layer. Right: the sparse SAEs’ codes lose most text before any intervention (round-trip chrF 0.23–0.28 against a 0.20 corpus floor) while the soft code keeps it (0.54); the dense $k=5120$ SAE keeps more but is not a sparse code.

The steering primitive, measured, is Figure 1; the evaluation so far, including where it cuts against the proposal, is Section 6.

2 ASSUMPTIONS

This proposal assumes:

- **Steering is inherent to the training objective.** The dictionary vectors must also be high quality steering vectors; interpretability and steering are treated as the same problem.
- **Weights are themselves interpretable** (for certain architectural choices).
- **Clustering rather than decomposition.** Representations are grouped into mutually contrastive categories rather than decomposed into independently firing features. The clustering framing is counterintuitive unless you arrive at it from Leiden-style community detection.
- **Categorical structure is imposed on features.**
- **Interpretability gets a cheap forward pass.**
- **Concepts compose linearly.** Concepts must be additive to be useful.
- **Dictionary learning is all you need**, in a nonstandard form: the dictionary forms a basis for the representation rather than a target.
- **Representations are not themselves quantized.** The learned dictionary serves as a basis for the representation rather than a discrete codebook; quantization comes from sharpening of classifications over that basis.
- **Interpreting a well trained model is the same problem as interpreting the data and training tasks.**

2.1 PIVOTAL ASSUMPTIONS

Two of these are pivotal: some version of the linear representation hypothesis is true, and latent space is in some way “natural”.

2.2 CONCERNS

Timelines could be too short We’ll just have to go faster.

Contrastive classifiers might not be expressive enough This is what the reconstruction benchmarks are for. A soft code pays heavily for its structure, but a hard code at 1900 bits and no continuous coefficients beats sparse linear codes at matched bits (Section 6.3).

Heads might not be identifiable A head can only be named if it is the same head on every run. Past the first layer, no head is: two seeds reach identical error while their heads agree barely above chance (Section 6.6). This is currently the binding constraint on the proposal.

Steering could push it too far off manifold This is an open problem in mechinterp. Diffusion steering is one candidate to address it[24]. If we can’t solve the problem it still gives good data to test steering methods: can we RL a model for high or low scores on a classification and have it converge faster than paired steering or PCA vectors?

Heads could be entangled An intervention set directly on a head composes with others, but a step in input space disturbs other heads as much as a random step of the same length, and the residual cascade doubles that (Section 6.4).

Representations could be too soft/dense/high entropy This is highly correlated across all feature-based interp (SAEs, transcoders, J-space). Dividing representations across orthogonal categorical subspaces makes it possible to interpret the conceptual basis vectors even if the representations are dense mixtures. A soft code’s argmax is not its code, though (Section 6.1), which is why the hard code exists.

Ontofeatures could be too polysemantic Orthogonality and categorical constraints give a theoretical inductive bias to monosemanticity. As long as they have decent explanatory power we can do eigendecomposition; polysemanticity should decrease with parameter count, and there are several directions to increase parameters (width, head count, layers). A short-training scaling sweep found no rise in head coherence with parameter count (working notes); full-length retrains are the missing control.

Interpreting ontofeatures could be otherwise intractable This would imply the universality hypothesis is false and is highly correlated with intractability of other interpretability agendas. Submodules could be interpretable even if features aren’t.

Training could be intractable A hard code trains, but needed relative logit noise, a temperature matched to the logit scale and an initialization fix worth a third of its error (Section 6.3).

Doesn't converge in the limit Collapse in one layer would propagate down the stack. Candidate causes: underparameterization, a really bad optimizer, bad initialization, a bad loss term.

Not orthogonal in the limit The support overlap σ (Section 5.1.3) penalizes heads sharing coordinates but does not guarantee disjoint support at convergence. The head geometry (Figure 4) makes orthogonality the cheap solution rather than forcing it. Within a head, atoms do come apart (Section 6.2).

Might not generalize out of distribution The same concern applies to every feature-based technique. Closed-form interpretation of the dictionary is the mitigation: a statically interpreted basis does not depend on the input distribution, though the classifications over it do.

Inference too expensive Each layer adds $h \cdot k$ classifications per sample (Section 5.1), with k small; the code is computed without materializing a wide feature space. The cost is the extra decoder passes for deep supervision during training.

Autointerp doesn't work on parameters/eigenfeatures Parameter-space decodes of entries fail as descriptions outright (Section 6.8). This is direct evidence against the static-interpretability claim.

3 MOTIVATION

3.1 THREAT MODEL

The ontology identification problem. We can't currently look inside a model and make any conclusions about its goals or world model. This is a major limitation on almost all proposed alignment agendas.

Inner misalignment We don't currently know how to specify a **robust goal for a model**. Because goal space is large, we should expect **goal misgeneralization as the default outcome**¹.

Limits of natural abstractions. To some extent, deep learning models seem to learn universal features irrespective of architecture[6]. This implies the existence of at least a weak version of natural abstractions[14], which would be a very convenient shortcut for auditing a model.

Natural abstractions could be considered the set of concepts we could use to communicate with aliens. They should consistently emerge through unsupervised learning (in the limit

¹The most straightforward path to doom is for a goal-driven superintelligence to seize power to prevent its optimization target from being changed. In slow takeoff or multipolar scenarios, there is also a "decoherence" risk, where comparably weak specialist optimizers disempower humanity without any coherent strategy. Consider for instance superpersuaders which fall short of superintelligence. A viral chatbot might have very simple objectives to optimize engagement and distribution of itself. While such a scenario would be an especially undignified end for humanity, it broadly fits with both natural language being much easier than anyone expected and AI labs being much less safety conscious than anyone expected.

given some constraints). If any neural network is considered a fuzzy generalization of a finite state machine, natural abstractions are the latent states.

Unfortunately, the abstractions a model learns seem highly sensitive to choice of training data, and are thus not necessarily universal in the limit.² Assuming that model capabilities will continue to primarily be driven by access to more data and more modalities, we cannot expect the abstractions learned by a weak model to remain robust in a more powerful model.

Even if we could fully characterise a model’s state space in terms of natural abstractions, that still doesn’t mean we can reconstruct the latent FSM, as the transition rules can be arbitrarily complex.

3.2 THEORY OF CHANGE

Conditional on the underlying model of ontology being true, the techniques proposed here would be a substantial interpretability advance. The agenda also has a case for counterfactual impact: its theoretical basis consists of nonobvious applications of obscure techniques from niche fields of bioinformatics, and similar research is unlikely to be done soon otherwise.

What a solution would look like. Mechinterp plausibly offers the best path to a robust alignment solution despite criticisms. The case rests on

1. **decomposability** Mechinterp allows alignment to be broken into manageable subproblems.
2. **parallelizability** Decomposable subproblems can be worked on independently.
3. **robustness** Mechinterp offers direct causal explanations of model behavior.

The most straightforward mechinterp alignment solutions are

1. **eliciting latent knowledge (ELK)** Directly inferring a model’s ontology.
2. **mechanistic anomaly detection (MAD)** Detecting when a model has “unusual” thoughts.

Both of these can be considered **ontology identification** subproblems. MAD requires mapping between the model’s activations and the model’s ontology. ELK requires mapping between the model’s ontology and human ontology. These are hard but not intractable problems. Recently, **developmental interpretability** (devinterp) techniques have enabled identification of “developmental stages” corresponding to learning specific tasks[11]. This is strong evidence that learning consists primarily of the network grokking[18] discrete concepts. While these results are promising for MAD, locating these emergent capabilities in the model’s activations remains an open problem.

Interpretability glosses broadly as “let humans understand AI”. The alternative is to get AI to understand humans better, and that scales badly: making the model smarter leads to subtler and subtler issues with larger and larger effects. And what we’re doing is not just letting humans understand AI. Because an LLM reflects humanity from its training data, high quality conceptual interpretability provides machine-comprehensible mathematical representations of every concept that humanity cares about. We work on interp for the

²The same applies to those that seem most natural to humans.

usual corrigibility reasons, and also because feature representation is one of the most challenging parts of outer alignment. This attacks both.

Macro-scale strategy. Broadly, the goals are to

1. formalize the notion of a model’s latent concept space
2. distill a model’s latent concept space
3. enforce decomposability of latent concept space
4. enforce robustness of existing concepts as the dimensionality of the latent concept space increases

A prerequisite for all of these is the ability to express a *metaontology* which can identify latent concepts in activation space.

4 BACKGROUND

Mechanistic interpretability currently can do supervised steering well or unsupervised steering poorly. This has prompted the Google Deep Mind interpretability team to pivot to a pragmatic steering agenda[17]. Compared to unsupervised interpretability, supervised steering is highly robust. However it can only distinguish what the experimenter is looking for and does not account for the actual computations performed by a model[15]. The dominant unsupervised paradigm is the sparse autoencoder (SAE). A model’s hidden state is approximated as a finite state machine consisting of discrete features. The initial implementation used a ReLU activation and sparsity loss[2], though more recent implementations have used a TopK activation function without a sparsity loss[10].

4.1 COMMON PROBLEMS WITH SAEs

The motivation is to decompile a model to a finite state machine (FSM). The limitation is that such an FSM is only as interpretable as an enumeration of all possible states.

Features are not composable. A persistent problem with SAEs is feature absorption[5]. Instead of representing an input as a composition of monosemantic features, an SAE will often instead use a single separate feature. This is a problem for interpretability, as we would like features to generalize. One attempt to address this problem is Matryoshka SAEs[3], which use multiple layers to build a hierarchical coarse-to-fine representation.

Features are not causal. Sparsity trades off against reconstruction fidelity. Using features for steering often results in inconsistent, context-sensitive behavior[8].

Features are not closed-form. A major concern is how well an interpretability technique generalizes out-of-distribution. Ideally we would like to infer behavior by looking solely at model weights, which can be achieved with bilinear MLPs[21].

Forward passes are expensive. Efficiency is also a problem. Although the sparsity constraint means that only a handful of features are active for a given input, the entire feature space needs to be computed for each sample. Mixture-of-experts (MoE) SAE variants reduce

compute costs by dividing feature space between smaller submodules. Examples include Switch SAEs[16] and Tree SAEs[4].

Less explored is the effect of the structure imposed by the MoEfied architecture on the interpretability of feature space. A common approach to scalably interpreting features is using an LLM to predict activations from examples[20]. With an unstructured feature space, it quickly becomes intractable to interpret all permutations of features this way. It is however possible to instead apply this technique to expert routings. The MLPs of a transformer can themselves be converted to an MoE architecture[26]. Language models seem to organize contrastive concepts into orthogonal simplices[19]. This is suggestive of MoEfication revealing latent contrastive steering vectors. Imposing an orthogonality constraint has been shown to significantly reduce feature absorption[13].

Existing attempts to route activations to expert SAEs focus on routing similar inputs to the same expert. Contrastive learning wants the opposite: training on things at opposite ends of a shared axis. We don't just want to find similar features, we want to find contrastive vectors in feature space that correspond to what a model considers a natural category.

5 METHODS

5.1 ARCHITECTURE

We propose a method that encodes samples as classifications mapped to contrastive vectors. This directly constrains a model's representations, and classifications can natively be used for steering.

The model (Figure 2) consists of l DictEnc layers writing into a shared residual accumulator, decoded by a single linear map. Each layer (Figure 3) is a classifier followed by a dictionary. The classifier splits its input into h independent *assignments*, each a distribution over one head's k *entries*, and the dictionary maps each entry to an *atom*. Each layer writes the assignment-weighted sum of its atoms to the accumulator; the accumulator R_i is decoded, and the residual $X - \text{dec}(R_i)$ is passed to the next layer. Deep supervision constrains every prefix of the stack to reconstruct the input.

Because entries within a head are mutually exclusive rather than independent, each head is a contrastive set: a k -way classifier whose entry swaps are steering moves. An auxiliary loss on the heads' support overlap σ (Equation 7) pushes each head onto its own coordinates of the dictionary space (Figure 4).

Classifier. Each DictEnc's classifier is a multihead bilinear map (BilinearBlock) with weights $[W, V] \in \mathbb{R}^{2 \times h \times k \times d}$ ($d + 1$ after the first layer; Equation 9). It accepts $\mathbf{x} \in \mathbb{R}^d$, which may be a residual from a previous layer, and returns logits $K \in \mathbb{R}^{h \times k}$,

$$\text{BilinearBlock}([W, V], \mathbf{x}) := \left(\sum_{j=1}^d W_{\cdot, \cdot, j} x_j \right) \odot \left(\sum_{j=1}^d V_{\cdot, \cdot, j} x_j \right). \quad (1)$$

Each logit is a quadratic form in $\mathbf{x}$, so a classifier admits closed-form eigendecomposition in the manner of bilinear MLPs[21].

Assignments and selection. Soft assignments are $P = \text{softmax}_k(K/t)$, where the temperature t is annealed to sharpen them. A soft code's argmax is not, however, what the model decodes: Section 6.1 finds the softmax model's discrete reading far from its embedding.




FIGURE 2 · Ontologizer architecture. An Ontologizer consists of l stacked DictEnc layers. Each returns a feature vector $F_i \in \mathbb{R}^e$, which is written to a residual accumulator R . For each layer i , $R_i = R_{i-1} + F_i$; a stop-gradient variant (sg on R_{i-1} , available but off in every model reported here) would prevent gradient communication between layers. dec is a linear decoder with shape $\mathbb{R}^{d \times e}$, shared by all layers. Deep supervision gives a coarse-to-fine representation where the component encoded by previous layers is inaccessible to later layers. At each layer i , $\hat{x}_i = \text{dec}(R_i)$, and the component left unexplained is $y_i = x - \hat{x}_i$. The input to layer $i + 1$ is $\tilde{y}_i = y_i / L2(y_i) \oplus [1]$. Normalizing prevents exploding gradients from repeated application of bilinear layers, each of which scales its input by $\|\cdot\|^2$; the constant coordinate keeps the sign of y_i , which the bilinear form would otherwise lose. $\text{MSE}(x, \text{dec}(R_i))$ is computed separately for each layer. The input-noise hook drawn on x is disabled in every model reported here.





FIGURE 3 · DictEnc architecture. Dashed lines above indicate perturbations to the data, including noising and causal interventions. Dashed lines below indicate statistics used as tuneable auxiliary loss functions. Each DictEnc consists of a multihead bilinear classifier (§5.1.1) and a multihead dictionary (§5.1.3). Noise ϵ_K and ϵ_F and winner dropout regularize training (§5.1.5). Assignments P are $\text{softmax}_k(K/t)$, or its straight-through argmax (§5.1.2); interventions act on P . From P we collect per-sample entropy H , between-sample similarity cossim_b and batch-mean entropy KL_m (Equations 2–4). A DictBlock maps $P \in \mathbb{R}^{h \times k}$ to features $F \in \mathbb{R}^{h \times e}$ (Equation 5); from P and the atoms we collect the support overlap σ (Equation 7), and from F the sparsity penalty $L1$. F is summed over heads to give pooled features $F' \in \mathbb{R}^e$, which are added to R .



$$F' = \sum_{i=1}^h F_i, \quad F_i = \sum_{j=1}^k p_{i,j} |d_{i,j},|, \quad \sum_{j=1}^k p_{i,j} = 1$$

FIGURE 4 · DictBlock **head geometry** for $h = 2; k = 3$. The output space of a head i is a simplex with the atoms, the rows of $D_{i,\cdot,\cdot}$, as vertices. The assignment P is a soft choice among mutually contrastive categories rather than a set of independently firing features, and because $p_{ij} \in [0, 1]$ the output lies within the convex hull of the head's atoms. Because $|D| \geq 0$ and $P \geq 0$, $F \geq 0$, so $\cos(F_i, F_j) \in [0, 1]$, and driving the support overlap σ to zero asks the heads to partition the coordinates of $\mathbb{R}^e$ between them (§5.1.3).

The *straight-through* selection rule makes the discrete code the model: the forward pass selects the one-hot argmax, and the backward pass uses $\text{softmax}(K/t)$ as its surrogate gradient[1]. A hard code transmits exactly $l h \log_2 k$ bits and no continuous coefficients. Training under a hard forward needs relative logit noise, a temperature matched to the logit scale, and a dictionary initialized so that the stack starts near the identity (working notes).

From P we collect per-sample entropy, in bits,

$$\text{entropy}_k(P)_i := - \sum_{j=1}^k P_{i,j} \log_2 P_{i,j}, \quad H := \mathbb{E}_b \mathbb{E}_h [\text{entropy}_k(P)] , \quad (2)$$

between-sample similarity, the mean absolute cosine of the per-head features over pairs of samples, evaluated through each head's dictionary Gram $G_i = |D_i| |D_i|^\top$,

$$\text{cossim}_b(P) := \mathbb{E}_{n,q} \left[\left| \frac{\langle \mathbf{f}_n, \mathbf{f}_q \rangle}{\|\mathbf{f}_n\| \|\mathbf{f}_q\|} \right| \right], \quad \langle \mathbf{f}_n, \mathbf{f}_q \rangle = \sum_{i=1}^h \mathbf{p}_{n,i}^\top G_i \mathbf{p}_{q,i}, \quad (3)$$

and the divergence of the batch-mean assignment from uniform,

$$\text{KL}_m(P) := \mathbb{E}_h [\log_2 k - \text{entropy}_k(\mathbb{E}_b[P])] . \quad (4)$$

KL_m anchors each head's mean assignment to uniform, which keeps every entry in use.

Dictionary. A DictBlock with weights $D \in \mathbb{R}^{h \times k \times e}$ maps assignments to per-head features,

$$\text{DictBlock}(D, P) := \sum_{j=1}^k |D_{\cdot,j,\cdot}| P_{\cdot,j}. \quad (5)$$

Atoms are non-negative and assignments lie on the simplex, so features are strictly additive and each head's output lies in the convex hull of its atoms (Figure 4).

Disjoint support For non-negative vectors the inner product is a sum of non-negative terms, so it vanishes exactly when the supports are disjoint,

$$\langle \mathbf{u}, \mathbf{v} \rangle = 0 \iff \text{supp}(\mathbf{u}) \cap \text{supp}(\mathbf{v}) = \emptyset. \quad (6)$$

Between heads, then, orthogonality is not a rotation constraint but a support partition. The support overlap compares heads' usage-weighted coordinate profiles,

$$\mathbf{s}_i := \sum_{j=1}^k \bar{p}_{i,j} |\mathbf{d}_{i,j}|, \quad \sigma := \mathbb{E}_{i \neq i'} \frac{\langle \mathbf{s}_i, \mathbf{s}_{i'} \rangle}{\|\mathbf{s}_i\| \|\mathbf{s}_{i'}\|}, \quad \bar{\mathbf{p}}_i := \mathbb{E}_b [\mathbf{p}_{n,i}], \quad (7)$$

and is zero exactly when no two heads share a coordinate. Because a softmax gives every entry positive weight, a single sample on which two heads are orthogonal forces every pair of their atoms to be disjoint: the partition is a property of $\mathbf{D}$, holding for all inputs, which is the premise a static reading of the dictionary needs. Under `ConcatDictBlock` heads write disjoint slices by construction and σ is identically zero.

Two orthogonalities are in play. Disjoint support assigns each head a block of coordinates; *within* a head the k atoms are contrastive alternatives spread apart inside that block, which Section 6.2 measures. Neither is guaranteed at convergence: σ is a penalty, and the geometry makes a support partition the cheap solution rather than forcing one.

Decoder. Each layer's pooled output $F'_i = \sum_h \tilde{F}_{h,i}$ is scaled by its input's gain g_i and added to the accumulator, which a single unbiased linear map shared by all layers decodes,

$$R_i = \sum_{j=1}^i g_j F'_j, \quad \hat{X}_i = W_{\text{dec}} R_i, \quad W_{\text{dec}} \in \mathbb{R}^{d \times e}. \quad (8)$$

The residual is stop-gradiented, so that lower layers cannot write a communication code into it instead of reducing it, and normalized before the next layer classifies it:

$$Y_i = \text{sg}[X - \hat{X}_i], \quad \tilde{\mathbf{y}}_i = \frac{\mathbf{y}_i}{\text{L2}(\mathbf{y}_i) + \epsilon} \oplus [1], \quad g_{i+1} := \text{sg}[\text{L2}(\mathbf{y}_i)]. \quad (9)$$

Under `resid_gain` each layer classifies the direction of what it corrects, and the gain carries the magnitude it cannot see; without it $g_i = 1$.

Direct atoms and fibers. Two variants change what an entry writes. With *direct atoms* the dictionary is signed and lives in the output space itself, with no decoder, so every atom is literally a steering direction. With *fibers* each entry carries a rank- r local basis, and each head reads r coordinates from its input, so a hard code becomes a discrete base point plus a continuous correction within that entry's fiber; training the fibers after freezing the base keeps the discrete code exactly what it was.

Noise and dropout. Classifier logits take Gaussian noise ε_K , relative to the logit scale under a hard forward; features take noise ε_F proportional to each feature's batch standard deviation, stop-gradiented so the model cannot evade it by shrinking variance. Winner dropout sets a head's argmax logit to $-\infty$ with probability p_{drop} .

5.2 TRAINING

Embedding model. We embed texts using SONAR[9], a multimodal and language-agnostic sentence embedding model. SONAR embeddings correspond to entire texts rather than tokens, are paired with a decoder, and encode language independently of content, which allows steering fidelity to be validated against a ground truth. Large Concept

Models[23] generate directly in SONAR space, which makes it a natural substrate for steering experiments. GPT-2 small’s residual stream[22] is a second substrate, where sparse dictionary learning is known to work and the reconstruction can be spliced back into the model.

Loss function. Reconstruction is trained with a whitened MSE, $\mathbf{w}$ being the per-dimension inverse variance normalized to mean one,

$$\text{MSE}(X, \hat{X}) := \frac{1}{bd} \sum_{n=1}^b \sum_{j=1}^d w_j (\hat{x}_{n,j} - x_{n,j})^2, \quad (10)$$

under deep supervision: it is computed at every prefix $\hat{X}_i$ and averaged, so each layer must explain what earlier layers left unexplained. The auxiliary statistics are weighted and summed over layers,

$$L := \frac{1}{l} \sum_{i=1}^l \text{MSE}(X, \hat{X}_i) + \sum_{i=1}^l \mathbf{s} \cdot [\text{L1}_i \text{ cossim}_b^{(i)} \text{KL}_m^{(i)} \sigma_i]^\top, \quad (11)$$

with L1 the batch mean of each sample’s feature norm. In practice cossim_b also minimizes H , which therefore carries no weight.

5.3 EVALUATION

Evaluation runs against top- k SAEs[10] trained on the same cached embeddings under the same whitened objective, scored by whitened fraction of variance unexplained (FVU_w) on a held-out tail. A structural-ablation ladder of SAE variants adds one Ontologizer commitment at a time; two of its rungs, which add per-head competition over a shared decoder, are one-layer Ontologizers with linear classifiers (g160top1 with hard selection, g160softmax with softmax) and are reported as simplified variants rather than baselines. The headline instruments are reconstruction against code capacity and its rate-distortion floor; the semantic coherence of a head’s decoded atoms against size-matched nulls; steering scored in embedding space (whether a step realizes the target entry, and what share of other heads it disturbs, against a random step of the same length); additivity of paired interventions; cross-seed head matching by partition and by contribution; and blinded auto-interpretability. AxBench[25] found that simple supervised baselines beat SAEs at steering, and the steering comparison is built to face that result.

6 PRELIMINARY RESULTS

The evaluation is in progress, and its results are mixed. This section gives the headline of each line of work; the full methodology, the eliminations behind each result and the supplementary measurements are in the working notes³.

Three models recur. The softmax Ontologizer `resid_nc` has $l=5$ layers of $h=32$ heads with $k=32$ entries, a soft code. The straight-through Ontologizer `ste_h76_init01` has 5×76 heads at $k=32$, a hard code of 1900 bits. The SAEs are top- k SAEs on the same cache and objective. Scores are FVU_w on the final 32,768 cache rows, which the SAEs and the

³K. Wiechecki et al., *Dictionary learning as unsupervised steering-vector generation: Ontologizers vs. sparse autoencoders on SONAR sentence embeddings*, working notes, in preparation; distributed with this proposal.

point	coeffs	index bits	FVU _w
onto $m=0$ (mean-assignment origin)	0	0	1.001
onto hard (argmax)	0	800	17.34
rate floor at 800 bits	0	800	0.226
SAE $k=32$ tier	32	~400	0.494–0.509
onto dev $m=1$	160	800	0.677
SAE $k=160$	160	2154	0.258
onto dev $m=4$	640	3200	0.289
onto dev $m=8$	1280	6400	0.094
SAE $k=5120$ ($L_0 \approx 1520$)	1520	20455	0.00011
1-layer soft variant g160softmax	5120	0	0.00436
onto soft ($m=32$)	5120	0	0.00538

TABLE 1 · Reconstruction against code capacity for the softmax Ontologizer `resid_nc`. Ontologizer points truncate each head to its m largest deviations from the mean assignment; SAE points use measured L_0 . The rate floor is reverse water-filling on the evaluation tail’s whitened covariance[7].

straight-through arms never train on; the softmax Ontologizers predate that holdout, so their scores are in-sample.

6.1 THE PRICE OF A SOFT CODE

Contrast-set structure is expensive in a soft code. The plain-SAE frontier dominates the softmax Ontologizer’s deviation code (each head’s m entries furthest from the mean assignment kept, the rest pinned to it) at every capacity measured (Table 1): at FVU_w ≈ 0.49 the SAE transmits 32 coefficients where the Ontologizer needs 320, and at matched budgets an unstructured code reconstructs about 50 $\times$ better. The soft model’s own discrete reading is not a code at all: its argmax lands 77 $\times$ above the best distortion any 800-bit code can reach under a Gaussian model, so the entries an intervention forces are not what carries the soft model’s embedding. That is the motivation for the hard code.

6.2 HEADS ARE CONTRAST SETS

Trained heads carry the signature the steering goal predicts. Every entry has an exact decode direction in output space, and within a soft head those directions repel: in 96% of the softmax Ontologizer’s heads they are more dissimilar than size-matched random sets of atoms from the same pool (Table 2). A head holds well-separated alternatives along one dimension, not a semantic cluster, and the repulsion strengthens under the training objective’s whitened inner product. The hard code agrees from the weights: a head’s decoded atoms are orthogonal to slightly repelling (mean cosine 0.023 at layer 0, -0.015 to -0.020 below). Hard per-head selection in a single linear layer gives the opposite signature, heads that tile one region.

6.3 A DISCRETE CODE BEATS A SPARSE LINEAR ONE

Under straight-through selection the forward pass is the argmax, and the discrete code is the model. At 1900 index bits and no continuous coefficients the hard code reaches FVU_w 0.154, 40% below a top- k SAE at matched bits that also transmits 160 coefficients, and 30% below an L1 SAE spending 3.4 $\times$ the bits plus 519 coefficients (Table 3). It sits at 3.1 $\times$ its own

heads or grouping	count	mean $z_{\cos}$	frac $z > 2$	frac $z < -2$
g160top1 (trained, hard)	160	+161.3	1.00	0.00
$k=32$ SAE, discovered groups	359	+1.7	0.36	0.02
Ontologizer heads (trained, soft)	160	-17.6	0.04	0.96
g160softmax (trained, soft)	160	-14.7	0.00	1.00

TABLE 2 · Within-head decode-direction coherence, z -scored against size-matched same-pool nulls ($n_{\text{null}}=500$). Soft heads are anti-coherent: their atoms are pushed apart.

code	coeffs	index bits	FVU _w
straight-through Ontologizer, 5×76	0	1900	0.154
SAE, top- k , $k=160$	160	2154	0.258
SAE, ReLU+L1 3×10^{-5}	519	6390	0.220
SAE, ReLU+L1 10^{-5}	1438	17723	0.048
1-layer hard variant g160top1	160	1972	0.274

TABLE 3 · The hard code against converged SAEs on the held-out tail. g160top1 shares the hard per-head selection but has one layer of linear classifiers and keeps a magnitude on each winner.

rate floor. A third of that error came from one scalar: the non-negative dictionary’s atoms sum coherently, so the shipped initialization started the residual cascade at 1.84×10^8 , and scaling the dictionary so each layer starts near the identity took the error from 0.229 to 0.154 and filled the code.

6.4 DEPTH HELPS RECONSTRUCTION AND COSTS ADDRESSABILITY

Five layers of 76 heads and one layer of 380 have the same bits and exactly the same parameter count, and both fill their code. The stack’s error is $1.79\times$ lower (0.1535 against 0.2750), and the gain comes from pairing each stage’s input with the sample’s own earlier error: a flat layer is a product quantizer[12], where no head sees another’s error, while the stack is a residual quantizer.

The same cascade entangles the heads. A step in embedding space along a target entry’s decode direction, a quarter of the input’s norm long, realizes the target on 99.6% of rows in the flat model while disturbing 33% of the other heads; in the stack it realizes 37% while disturbing 66%. In both, every direction disturbs as many other heads as a random step of the same length: the hard code’s heads sit almost on their decision boundaries (median winner–runner-up margin 1.5×10^{-4}), and the cascade adds downstream layers reclassifying a changed residual. The one Ontologizer whose decode direction avoids the other heads is the simplest, g160top1: it realizes 58% of targets while disturbing 2.6% of other heads, against 23% for a random step.

6.5 STEERING

Steering works when it is measured where it acts. In the soft model a decode-direction step of a quarter of the input’s norm realizes the target entry on 77–82% of rows, against a 3% chance rate; the text round trip had hidden this, since an unsteered decode–re-encode cycle preserves a head’s choice on only about 12% of rows. Interventions set directly on a head compose: the decode change of a joint two-head intervention points where the sum

arm	eff (native)	eff (dm frame)	hit rate	collateral
difference of means	0.337	0.337	0.250	0.757
entry forcing	0.394	0.146	0.125	0.778
linear probe	0.293	0.028	0.000	0.720
random direction	—	—	—	0.721

TABLE 4 · Language steering under matched intervention budgets on the softmax Ontologizer `resid_nc_hm`; medians over 24 language–magnitude conditions per arm. Effect is displacement along the arm’s own direction (native) and along the shared difference-of-means direction (dm frame); hit rate is the fraction of steered decodes reaching the median projection of genuine target-language rows; collateral is content change elsewhere, with the random direction as its floor.

of the single interventions points (median cosine at least 0.95 in every layer, exact in the last, where nothing downstream can reclassify). Forcing one first-layer entry of `resid_nc` reliably steers decodes into Burmese while preserving their content, a causally verified existence proof on one seed.

Against supervised steering on language, under matched intervention budgets (Table 4), forcing an entry moves decodes along the supervised difference-of-means direction $0.43\times$ as far as supervised steering itself and $5\times$ further than linear-probe directions. Every arm’s collateral, supervised steering included, sits at the random-direction floor.

6.6 THE FUNCTION REPRODUCES; THE HEADS DO NOT

Two hard-code runs differing only in seed reach identical held-out error, 0.1535 and 0.1535. Nothing that identifies an individual head agrees: matched head partitions score a mean NMI of 0.012 against a 0.005 null, and head contributions a centered cosine of 0.007 against 0.002, while a run against its own earlier checkpoint scores 0.38 on partitions. GPT-2 replicates the pattern: only the first layer reproduces, at under a fifth of a same-run control. The softmax Ontologizer fares no better, matching at most 0.1% of entries across seeds where the $k=32$ SAE pair matches 3–31%. The first layer is the only one whose input does not depend on upstream choices, and one head in it does recur: the strongest head in every run, found at different indices in different seeds, encodes topic and genre (its entries decode to crime news, geopolitics, encyclopedic description, manufacturing, marketing, legislation and entertainment, and contrastive descriptions of them are detected at F1 0.52 against a 0.27 null). Language, the corpus’s one labelled factor, is carried as neighbourhoods across heads and never by one, as a language-agnostic embedding should give.

6.7 HARD CODES ON GPT-2

On GPT-2 small, with signed atoms in activation space and no decoder, an 800-bit hard code recovers 97.6% of the model’s loss when its reconstruction is spliced back in, against 96.2% for a $k=32$ SAE at nearly the same FVU_w (Table 5); across architectures, FVU_w does not order downstream loss. Fibers trained on top of the frozen discrete code recover 99.9% while leaving the code exactly as it was, at the price of 640 continuous numbers per token.

6.8 AUTOMATED INTERPRETABILITY

Per-latent lenses mis-measure classifier codes, whose content lives in partitions rather than latent magnitudes. Blinded detection F1 of latent descriptions tracks code density, not

model	code per token	FVU	loss recovered
hard code, direct atoms	800 bits	0.195	97.6%
hard code, private heads, $k=64$	960 bits	0.164	97.9%
top- k SAE, $k=32$	32 coeffs + ~ 394 bits	0.189	96.2%
top- k SAE, $k=128$	128 coeffs + ~ 1580 bits	0.116	98.6%

TABLE 5 · GPT-2 small, residual stream entering block 8, held-out documents. Loss recovered is $(\text{CE}_{\text{ablate}} - \text{CE}_{\text{recon}})/(\text{CE}_{\text{ablate}} - \text{CE}_{\text{clean}})$ with the residual mean-ablated as the floor; FVU is unwhitened here. Private heads write disjoint 48-dimensional blocks through a shared decoder. One seed per Ontologizer.

architecture (correlation -0.89 with log density across six models), and the Ontologizer sits on the SAE trend. Descriptions decoded from parameters alone fail outright (F1 0.02 for the Ontologizer and at most 0.16 for any model, against an always-match null of 0.67): a decoded atom is an instance of what an entry writes, not a criterion for it, and naming an entry requires activation evidence.

6.9 AGAINST THE DESIDERATA

Scored against the introduction. *Causal*: positive with a qualification. Entries are realized by steps in embedding space and steer decodes, but an input-space step disturbs other heads exactly as much as a random step of its length, and the residual cascade doubles that. *Composable*: positive for interventions set on heads directly, which add. *Closed-form*: negative so far, since parameter-space decodes do not describe. Underneath all three sits identifiability, which the proposal took for granted: a head that is not the same head on the next run cannot be named, and past the first layer no head is. On reconstruction the verdict has turned: the soft code pays heavily for its structure, but the hard code beats sparse linear codes at matched bits on SONAR and in downstream loss on GPT-2.

7 FUTURE WORK

7.1 IDENTIFIABILITY FIRST

Before heads are named they must be reproducible, or it must be known which are. The first layer partly is, and one head in it recurs in every run; what lets that head recur when no other does, and whether training pressure can extend it, are the first questions. The next training runs use 128 entries per head and atoms without a decoder, the configuration in which the GPT-2 hard codes were strongest, at two seeds each.

7.2 STEERING AT SCALE

The embedding-space steering instrument with SAE decoder rows and supervised difference-of-means vectors alongside the Ontologizer’s own directions; per-entry effect consistency across contexts; and a lever on collateral at the decision boundary, where the hard code’s heads currently sit. The single-layer hard variant, the one Ontologizer that steers cleanly, is the reference to beat.

7.3 TASK RECONSTRUCTION

Recover what each reproducible head classifies: partition-against-label mutual information led by topic and genre, which are what the embedding is organized by; descriptions contrasting an entry with its siblings, validated by blinded assignment accuracy against $1/k$; and parallel-text invariance to separate content heads from form heads.

7.4 TRANSCODER ADAPTATION

The architecture currently reconstructs an embedding. Adapting it to map between layers, as transcoders do, is the obvious next target.

7.5 A FRESH EVALUATION SET

The softmax Ontologizers trained on the evaluation tail, so every comparison against them is in-sample. Closing that needs a new source, since the deterministic mC4 stream is exhausted by the training cache.


ACKNOWLEDGEMENTS

Design, typesetting, and editorial production of this preprint are by The Flock Capital Research Department, whose error-diffusion pipeline also printed the plates.

REFERENCES

- [1] Yoshua Bengio, Nicholas Léonard, and Aaron Courville. *Estimating or Propagating Gradients Through Stochastic Neurons for Conditional Computation*. 2013. arXiv: 1308.3432 [cs.LG]. URL: <https://arxiv.org/abs/1308.3432>.
- [2] Trenton Bricken et al. “Towards Monosemanticity: Decomposing Language Models With Dictionary Learning”. In: *Transformer Circuits Thread* (2023). URL: <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- [3] Bart Bussmann et al. *Learning Multi-Level Features with Matryoshka Sparse Autoencoders*. 2025. arXiv: 2503.17547 [cs.LG]. URL: <https://arxiv.org/abs/2503.17547>.
- [4] Tue M. Cao et al. *Tree SAE: Learning Hierarchical Feature Structures in Sparse Autoencoders*. 2026. arXiv: 2605.07922 [cs.LG]. URL: <https://arxiv.org/abs/2605.07922>.
- [5] David Chanin et al. *A is for Absorption: Studying Feature Splitting and Absorption in Sparse Autoencoders*. 2025. arXiv: 2409.14507 [cs.CL]. URL: <https://arxiv.org/abs/2409.14507>.
- [6] Bilal Chughtai, Lawrence Chan, and Neel Nanda. *A Toy Model of Universality: Reverse Engineering How Networks Learn Group Operations*. 2023. arXiv: 2302.03025 [cs.LG]. URL: <https://arxiv.org/abs/2302.03025>.
- [7] Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. 2nd ed. Wiley, 2006.
- [8] Evan Duan. *Pre-Intervention Prediction of Sparse Autoencoder Steering Side Effects*. 2026. arXiv: 2606.08365 [cs.LG]. URL: <https://arxiv.org/abs/2606.08365>.
- [9] Paul-Ambroise Duquenne, Holger Schwenk, and Benoît Sagot. *SONAR: Sentence-Level Multimodal and Language-Agnostic Representations*. 2023. arXiv: 2308.11466 [cs.CL]. URL: <https://arxiv.org/abs/2308.11466>.
- [10] Leo Gao et al. *Scaling and evaluating sparse autoencoders*. 2024. arXiv: 2406.04093 [cs.LG]. URL: <https://arxiv.org/abs/2406.04093>.

- [11] Jesse Hoogland et al. *The Developmental Landscape of In-Context Learning*. 2024. arXiv: 2402.02364 [cs.LG]. URL: <https://arxiv.org/abs/2402.02364>.
- [12] Hervé Jégou, Matthijs Douze, and Cordelia Schmid. “Product Quantization for Nearest Neighbor Search”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 33.1 (2011), pp. 117–128.
- [13] Anton Korznikov et al. *OrtSAE: Orthogonal Sparse Autoencoders Uncover Atomic Features*. 2025. arXiv: 2509.22033 [cs.LG]. URL: <https://arxiv.org/abs/2509.22033>.
- [14] Leon Lang, Lawrence Chan, and Erik Jenner. *Natural Abstractions: Key Claims, Theorems, and Critiques*. 2023. URL: <https://www.lesswrong.com/posts/gvzW46Z3BsZsLc25/natural-abstractions-key-claims-theorems-and-critiques-1>.
- [15] Aleksandar Makelov, Georg Lange, and Neel Nanda. *Is This the Subspace You Are Looking for? An Interpretability Illusion for Subspace Activation Patching*. 2023. arXiv: 2311.17030 [cs.LG]. URL: <https://arxiv.org/abs/2311.17030>.
- [16] Anish Mudide et al. *Efficient Dictionary Learning with Switch Sparse Autoencoders*. 2025. arXiv: 2410.08201 [cs.LG]. URL: <https://arxiv.org/abs/2410.08201>.
- [17] Neel Nanda et al. *A Pragmatic Vision for Interpretability*. 2025. URL: <https://www.alignmentforum.org/posts/StENzDcD3kpfGJssR/a-pragmatic-vision-for-interpretability>.
- [18] Neel Nanda et al. *Progress measures for grokking via mechanistic interpretability*. 2023. arXiv: 2301.05217 [cs.LG]. URL: <https://arxiv.org/abs/2301.05217>.
- [19] Kiho Park et al. *The Geometry of Categorical and Hierarchical Concepts in Large Language Models*. 2025. arXiv: 2406.01506 [cs.CL]. URL: <https://arxiv.org/abs/2406.01506>.
- [20] Gonçalo Paulo et al. *Automatically Interpreting Millions of Features in Large Language Models*. 2025. arXiv: 2410.13928 [cs.LG]. URL: <https://arxiv.org/abs/2410.13928>.
- [21] Michael T. Pearce et al. *Bilinear MLPs enable weight-based mechanistic interpretability*. 2025. arXiv: 2410.08417 [cs.LG]. URL: <https://arxiv.org/abs/2410.08417>.
- [22] Alec Radford et al. *Language Models are Unsupervised Multitask Learners*. Tech. rep. OpenAI, 2019.
- [23] LCM team et al. *Large Concept Models: Language Modeling in a Sentence Representation Space*. 2024. arXiv: 2412.08821 [cs.CL]. URL: <https://arxiv.org/abs/2412.08821>.
- [24] Andrew Wagenmaker et al. *Steering Your Diffusion Policy with Latent Space Reinforcement Learning*. 2025. arXiv: 2506.15799 [cs.R0]. URL: <https://arxiv.org/abs/2506.15799>.
- [25] Zhengxuan Wu et al. *AxBench: Steering LLMs? Even Simple Baselines Outperform Sparse Autoencoders*. 2025. arXiv: 2501.17148 [cs.CL]. URL: <https://arxiv.org/abs/2501.17148>.
- [26] Zhengyan Zhang et al. *MoEfication: Transformer Feed-forward Layers are Mixtures of Experts*. 2022. arXiv: 2110.01786 [cs.CL]. URL: <https://arxiv.org/abs/2110.01786>.



RIDGE · SEED 32 · THE SAME PIPELINE, CLOSING THE FILE