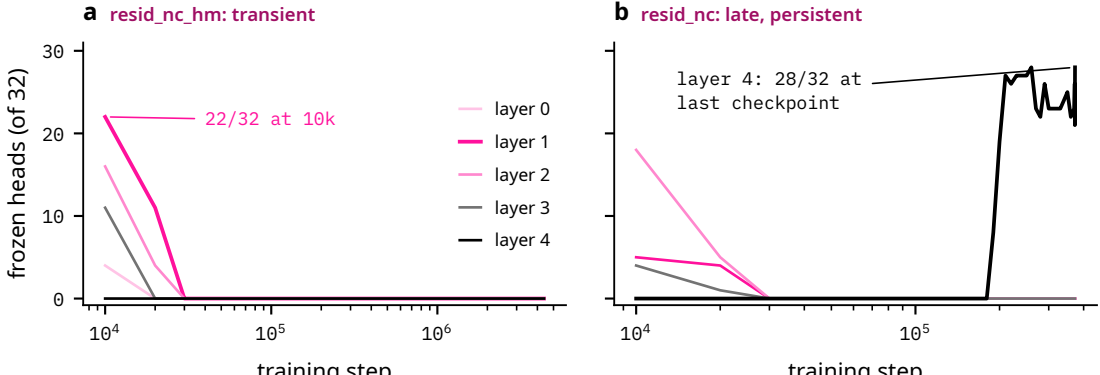
a
resid_nc_hm: transient
frozen heads (of 32)
30
20
10
0
22/32 at 10k
layer 0
layer 1
layer 2
layer 3
layer 4
$10^4$
$10^5$
$10^6$
training step
b
resid_nc: late, persistent
layer 4: 28/32 at
last checkpoint
$10^4$
$10^5$
training step