



Dictionary Learning as Unsupervised Steering-Vector Generation

Keira Wiechecki

WAKARU · KAW504@NYU.EDU

Lamb Wakefield

FLOCK CAPITAL · LAMB@FLOCK.HOLDINGS


Jade Zaslavsky

WAKARU · JADEOFNAMELESS@GMAIL

vmfunc

FLOCK CAPITAL · CELESTE@LINUX.COM

October 7, 2026—working notes, results preliminary



ORBIT · SEED 1024 · FLOYD-STEINBERG ERROR DIFFUSION

ABSTRACT

The Ontologizer is a multihead, multilayer dictionary-learning architecture whose design goal is unsupervised steering-vector generation: a stack of competitive classifier “heads” reconstructing an embedding space should discover the space’s natural contrastive clusters, MoEfyng the representation into small, nameable classification tasks whose within-head entry swaps are steering moves. We compare a softmax (soft-code) Ontologizer and a straight-through (hard-code) arm against TopK sparse autoencoders (SAEs) on SONAR sentence embeddings under the same whitened objective, and the hard code on GPT-2 small residuals. Soft heads are contrast sets — within-head decode directions repel — but their argmax is not the model’s code, and the plain-SAE frontier dominates them in reconstruction. The straight-through arm makes the discrete code the model: at 1900 index bits and no continuous coefficients it reaches FVU_w 0.154, against 0.258 for a TopK SAE at matched bits that also transmits 160 coefficients. At fixed capacity the residual cascade is worth $1.79\times$ over a flat layer, but doubles the collateral of a single-head intervention (66% of other heads against 33%), and two seeds reach identical error while no head past the first layer reproduces. The strongest head encodes topic and genre rather than language, as a language-agnostic embedding should be organized. On GPT-2, an 800-bit hard code of signed activation-space directions recovers 97.6% of the model’s loss, more than a $k_{SAE}=32$ SAE at nearly the same FVU_w .




FIGURE 1 · The softmax Ontologizer’s steering primitive, measured. (a) Forcing an entry moves decodes along the supervised difference-of-means direction $0.43\times$ as far as supervised steering itself and $5\times$ further than probe directions (`resid_nc_hm`; Table 17). (b) Joint two-head interventions point where the summed single interventions point, exactly in the final layer (`resid_nc`; Section 8). (c) The sparse SAEs’ codes lose most text before any intervention (round-trip chrF 0.23–0.28 against a 0.20 corpus floor) while the soft code keeps it (0.54); the dense k5120 keeps more but is not a sparse code (Appendix C.4).

1 EVALUATION FRAME

The comparisons below were largely built on SAE-native terms — per-feature monosemanticity, sparse probing, reconstruction at matched sparsity — and we report them unchanged. But the Ontologizer’s goal is not a monosemantic dictionary. It is an *unsupervised generator of steering vectors*: each head is a k -way classifier over its input, trained only by reconstruction pressure, and the hope is that a series of such classifiers finds the most natural contrastive clusters in the embedding space. Interpretability value is then expected from *MoEfication*: decomposing the space into small classification tasks — one per head — each nameable in isolation, with within-head entry swaps as the steering primitive (forcing head i from its argmax entry to entry j' moves the decode along exactly the dimension head i classifies).

Under this frame the evaluation questions are: (a) what does contrast-set structure cost in reconstruction (the Pareto and ladder); (b) do trained heads actually behave as contrastive classifiers (categorical form, semantic anti-coherence); (c) are the learned tasks *natural* — nameable, stable across runs, aligned with generative factors of the data (seed reproducibility, language and topic alignment); and (d) is the steering primitive good — consistent effect, low collateral, at least matching supervised diff-of-means vectors and SAE decoder-row addition (steering fidelity). Metrics that ask whether *individual latents* are concept-aligned remain informative about the SAE side but are the wrong lens for classifier codes, and we flag them as such where they appear. Figure 1 shows the softmax Ontologizer on three of question (d)’s criteria; the hard code steers less reliably than the soft one (Section 4.4).

Question (c) has a prerequisite that the original battery took for granted: a head can only be named if it is the same head on every run. Section 5.5 tests that directly and finds it holds only partly, and only in the first layer, which conditions every per-head result in this document.

2 RELATED WORK

No single line of prior work covers this combination; each commitment has a nearest neighbor.

Discrete and routed dictionaries. A DictBlock head — a softmax over k entries, summed across heads — is soft product quantization [1], and the residual layer stack is soft residual VQ, the codec lineage of SoundStream [2]; that literature treats codes as compression and does not read codebooks as features. The straight-through arm (Section 5) makes the analogy literal: with an argmax forward and a softmax surrogate gradient [3] it is a residual product quantizer with non-negative atoms. Codebook Features [4] is the closest interpretability ancestor (VQ bottlenecks finetuned into a network, hidden state a sum of a few discrete codes), but its codes are read per-code, not as classifiers, and there is no intervention grammar over within-codebook swaps. Switch SAEs [5] route activations among expert SAEs purely as a compute optimization; notably their experts collapse into redundant ensembles that duplicate features — exactly the failure mode `headstruct.py/headcoh.py` instrument, and the opposite of what trained heads show here (within-head directions repel). MoEfication [6] supplies the frame’s name but decomposes already-trained MLPs post hoc; here the partition is the training objective’s own discovery.

Structured SAEs. The field’s main structure axis is nesting by scale: Matryoshka SAEs [7] train nested dictionaries to curb feature absorption [8], and Tree SAE [9] learns explicit hierarchies. Rung 3 (`k32_p5`) imports exactly the matryoshka commitment into the ladder; the Ontologizer’s axis is instead *partition into exhaustive, mutually exclusive heads*. Both lines ask whether an architectural commitment buys interpretability that code density alone does not explain — the autointerp results below answer “not per-latent describability,” while the description-coherence view is where the partition begins to pay.

Embedding-space substrate. SAEs over dense text embeddings became a subfield in 2025–2026: discretizing retrieval embeddings [10], data-analysis toolkits [11], human-concept alignment [12]. All are single-layer TopK SAEs, and none address SONAR, where the stakes are higher — Large Concept Models generate directly in SONAR space [13] — and where the paired M2M100 decoder supplies a semantic oracle: where embedding-inversion work must train an inverter [14], `textfid.py`

and `steerfid.py` decode reconstructions and interventions through the pretrained decoder itself.

Evaluation stance. SAEBench [15] is the template for the multi-instrument battery (`langprobe.py` is deliberately SAEBench-style), and the blinded detection scoring follows the scaled autointerp line [16]. AxBench [17] found that prompting and finetuning beat every representation-based steering method, with SAEs uncompetitive even at concept detection — independent support for judging dictionaries by steering utility rather than per-latent monosemanticity, and the niche `steerfid.py`’s effect-vs-collateral-at-matched-magnitude design targets. Song et al. [18] argue seed-to-seed feature consistency should be a first-class SAE metric; `splitting.py`’s match fractions and `partition.py`’s head matching measure precisely that. Without close analog, to our knowledge: reading per-head routing as contrastive classification with a causal intervention interface; quota-free coherence tests against size-matched nulls, including over judge-written description texts; and decoding parameter directions through the paired decoder as a test of whether features can name themselves from weights alone — a test the weights fail.

3 METHODS

3.1 ARCHITECTURE

We propose a novel method which encodes samples as classifications mapped to contrastive vectors. This directly constrains a model’s representations. Moreover, classifications can natively be used for steering. Appendix A collects the notation.

The model (Figure 2) consists of l DictEnc layers writing into a shared residual accumulator, decoded by a single linear map. Each layer (Figure 3) is composed of a classifier layer followed by a dictionary layer. The classifier splits the input into h independent soft assignment vectors, each over one head’s k entries. The dictionary maps each entry to an atom. Each layer writes a weighted average of these atoms to an accumulator. After each layer the accumulator is decoded, and the residual is passed on: layer ℓ receives $\mathbf{x} - \text{dec}(\mathbf{r}_\ell)$, where $\mathbf{r}_\ell$ holds the ℓ layers before it. Deep supervision constrains every prefix of the stack to reconstruct the input, which makes the code coarse-to-fine: what earlier layers encode is gone from the residual later layers see.

Because entries within a head are mutually exclusive rather than independent, each dictionary head serves as a contrastive set. An auxiliary loss term minimizing the heads’ support overlap ω (Eq. (6)) pushes each head onto its own coordinates of the dictionary space (Figure 5). Because the dictionary returns a linear combination of representative vectors, subspaces are convex.

DictEnc. Each DictEnc consists of a multihead bilinear classifier (BilinearBlock) with weights $[\mathbf{W}, \mathbf{V}] \in \mathbb{R}^{2 \times h \times k \times d}$ ($d + 1$ with the constant coordinate of Eq. (14)) and




FIGURE 2 · **Ontologizer architecture.** Each DictEnc writes its gain-scaled output $g_\ell \mathbf{f}'_\ell$ into a shared accumulator. States count the layers applied: layer ℓ reads $\mathbf{r}_\ell$ and the residual $\mathbf{y}_\ell$, and writes $\mathbf{r}_{\ell+1}$. One decoder serves every prefix, and each prefix is scored against $\mathbf{x}$ (deep supervision). The dashed input-noise hook is off in every reported model. Section 3.1.5 gives the equations.

a multihead linear dictionary (DictBlock) with weights $\mathbf{D} \in \mathbb{R}^{h \times k \times e}$. Its input is the residual $\mathbf{y}_\ell$ left by the layers before it ($\mathbf{x}$ itself at layer 0), of which the classifier reads the shape $\mathbf{u}_\ell$ (Eq. (14)). Algorithm 1 gives one layer’s training pass, and Algorithm 3 the stack and its loss.

Classifier. A BilinearBlock accepts an input $\mathbf{x} \in \mathbb{R}^d$ (in a layer, the shape $\mathbf{u}_\ell$). It returns logits $K \in \mathbb{R}^{h \times k}$ given by

$$\text{BilinearBlock}([W, V], \mathbf{x}) := \left(\sum_{q=1}^d W_{\cdot, \cdot, q} x_q \right) \odot \left(\sum_{q=1}^d V_{\cdot, \cdot, q} x_q \right) \quad (1)$$

An optional noise term ε_K may be added. Winner dropout sets each head’s largest logit to $-\infty$ with probability p_{drop} .

Assignments. Soft assignments P are given by $\text{softmax}_k(K/t)$ where t is a temperature hyperparameter. Annealing it sharpens the assignments; each model’s schedule is in Appendix B.3. The architecture provides methods for causal interventions on assignments. From P we collect per-sample entropy H , in bits, so a head saturates at $\log_2 k = 5$ bits,

Algorithm 1: DictEnc, one layer, in training (DictEnc.withStats). Without the router, Γ is all ones and $L1_S = 0$; without fibers there are no coordinates C ; without resid_gain the input is not split and $g = 1$.

Input: accumulator $R \in \mathbb{R}^{b \times e}$; layer input $E \in \mathbb{R}^{b \times d'}$, whose row n is $[\mathbf{y}_n; \mathbf{1}_{n_{\text{const}}}]$: the varying part $\mathbf{y}_n$ (the encoded input at layer 0, the residual above) and n_{const} constant coordinates (resid_const)

Data: bilinear classifier (Eq. (1)); router (scaled); fiber readouts $A_i \in \mathbb{R}^{r \times d'}$; DictBlock (Algorithm 2)

Hyperparameters: logit noise σ_K ; fiber bound β ; the DictBlock hyperparameters

Output: accumulator R ; assignments $P \in \mathbb{R}^{b \times h \times k}$; statistics row ψ (Eq. (22))

begin

$g_n \leftarrow \text{sg}[\|\mathbf{y}_n\|]; \mathbf{u}_n \leftarrow [\mathbf{y}_n/g_n; \mathbf{1}_{n_{\text{const}}}]$ // gain and shape

$K, L1_K \leftarrow \text{classifier}(\mathbf{u}_{1:b})$ // L1 of the gate factor

$K \leftarrow \text{noise}_K(K)$ // Eq. (16) or (17)

if router **then**

$\Gamma, L1_S \leftarrow \text{router}(\mathbf{u}_{1:b}); \Gamma \leftarrow |\Gamma|$ // unless router_signed

else

$\Gamma \leftarrow \mathbf{1}; L1_S \leftarrow 0$

end

$\mathbf{c}_{ni} \leftarrow \beta \tanh(A_i \mathbf{u}_n / \beta)$ // fiber coordinates, Eq. (19)

$F', P, (L1, H, \text{cossim}_b, \mathbb{E}\kappa, \max \kappa, \text{KL}_m, \omega) \leftarrow \text{DictBlock}(K, \Gamma, C)$

$\text{KL}_{\text{pwak}}, L2_{\text{pwak}} \leftarrow \text{pwak}(P, \mathbf{u}_{1:b})$ // zero unless enabled

$\mathbf{r}_n \leftarrow \mathbf{r}_n + g_n \mathbf{f}'_n$ // row n of R : the contribution, at the input's gain

return $R, P,$

$\psi = (L1_K, L1, H, \text{cossim}_b, \mathbb{E}\kappa, \max \kappa, \text{KL}_m, \text{KL}_{\text{pwak}}, L2_{\text{pwak}}, L1_S, \omega)$

end




FIGURE 3 · DictEnc architecture. A bilinear classifier reads the layer’s shape $\mathbf{u}$ and returns logits K ; the assignments $P = \text{softmax}_k(K/t)$ select atoms from the dictionary, and the per-head outputs F are pooled into $\mathbf{f}'$. Dashed boxes above an arrow perturb it in training or carry an intervention; dashed boxes below are statistics, usable as auxiliary losses (Sections 3.1.2–3.1.4).

$$\text{entropy}_k(\mathbf{p}) := - \sum_{j=1}^k p_j \log_2 p_j, \quad H := \mathbb{E}_n \mathbb{E}_i [\text{entropy}_k(\mathbf{p}_{ni})] \quad (2)$$

between-sample similarity cossim_b , the mean absolute cosine similarity of the per-head features F over all ordered pairs of samples,

$$\text{cossim}_b(\mathbf{P}) := \mathbb{E}_{n,n'} \left[\left\| \frac{\langle \mathbf{f}_n, \mathbf{f}_{n'} \rangle}{\|\mathbf{f}_n\| \|\mathbf{f}_{n'}\|} \right\| \right], \quad \langle \mathbf{f}_n, \mathbf{f}_{n'} \rangle = \sum_{i=1}^h \mathbf{p}_{ni}^\top G_i \mathbf{p}_{n'i} \quad (3)$$

where $\mathbf{f}_{ni} := F_{ni} \in \mathbb{R}^e$ is head i ’s feature vector for sample n , $\mathbf{f}_n = \bigoplus_{i=1}^h \mathbf{f}_{ni}$ is their concatenation, and $G_i = |D_i| |D_i|^\top \in \mathbb{R}^{k \times k}$ is head i ’s dictionary Gram, with $D_i := D_{i..}$ its $k \times e$ dictionary. The implementation evaluates Eq. (3) in entry space through the G_i rather than materializing F ; the two are identical in value and gradient, but $\mathcal{O}(hk^2e + b^2hk)$ instead of $\mathcal{O}(b^2he)$.

Finally, KL_m is the divergence of the batch-mean assignment vector from uniform, in bits, averaged over heads.

$$\text{KL}_m(\mathbf{P}) := \mathbb{E}_i [\log_2 k - \text{entropy}_k(\bar{\mathbf{p}}_i)], \quad \bar{\mathbf{p}}_i := \mathbb{E}_n [\mathbf{p}_{ni}] \quad (4)$$

In practice H is not needed, as including a cossim_b loss also minimizes H . KL_m serves to anchor the batch mean of each assignment vector $\mathbf{p}_{ni}$ to a uniform distribution. This results in a uniform $\mathbf{p}_{ni}$ code corresponding to a maximally generic feature vector.

Dictionary. A DictBlock accepts assignments $P \in \mathbb{R}^{h \times k}$ and returns features $F \in \mathbb{R}^{h \times e}$, whose row $\mathbf{f}_i$ is head i ’s output, given by

$$\text{DictBlock}(\mathbf{D}, \mathbf{P})_{i\cdot} = \mathbf{f}_i := \sum_{j=1}^k p_{ij} \mathbf{a}_{ij}, \quad \mathbf{a}_{ij} := |\mathbf{D}_{ij\cdot}| \quad (5)$$

where the atom $\mathbf{a}_{ij}$ is the absolute value of head i 's dictionary row for entry j ($\mathbf{D}_{ij\cdot}$ itself when the dictionary is signed). Because atoms are absolute values and $p_{ij} \in [0, 1]$, features are strictly additive. Since each head's assignment sums to one, $\mathbf{f}_i$ lies in the simplex whose vertices are the head's atoms (Figure 5): a head's entries are mutually exclusive alternatives, not independently firing features. Algorithm 2 gives the full training pass, with the regularizing perturbations and statistics around this lookup.

From the assignments and the atoms we collect the heads' support overlap ω , the mean cosine over ordered pairs of distinct heads between their usage-weighted coordinate profiles $\boldsymbol{\pi}_i$,

$$\boldsymbol{\pi}_i := \sum_{j=1}^k \bar{p}_{ij} |\mathbf{a}_{ij}|, \quad \omega := \mathbb{E}_{i \neq i'} \frac{\langle \boldsymbol{\pi}_i, \boldsymbol{\pi}_{i'} \rangle}{\|\boldsymbol{\pi}_i\| \|\boldsymbol{\pi}_{i'}\|} \quad (6)$$

with $\bar{\mathbf{p}}_i$ the batch-mean assignment of Eq. (4). The overlap is zero exactly when no two heads share a coordinate of the dictionary space and one when every head uses the same coordinates in the same proportions (Section 3.1.4). It costs $O(hke + h^2e)$, independent of the batch. Weighting by usage keeps dead entries out of it and gives the classifier a gradient path. Figure 4 shows the assignments, the atoms and both cosine statistics for one trained head. We also collect the sparsity penalty L1 after the noise ε_F is added, $\tilde{\mathbf{F}} := \mathbf{F} + \varepsilon_F$. L1 is the batch mean of each sample's norm, so it does not change with b but scales with h and e ; its loss weight is correspondingly small.

$$\text{L1}(\tilde{\mathbf{F}}) := \frac{1}{b} \sum_{n=1}^b \sum_{i=1}^h \sum_{q=1}^e |\tilde{f}_{niq}| \quad (7)$$

L1 tends to decrease with only MSE loss, but incorporating it into the loss speeds up convergence.

Disjoint support Between heads, orthogonality of nonnegative features is not a rotation constraint but a support-partition constraint. Every $\mathbf{f}_{ni}$ lies in the nonnegative orthant of $\mathbb{R}^e$: the dictionary contributes $|\mathbf{D}| \geq 0$ and the assignments $\mathbf{P} \geq 0$. For nonnegative $\mathbf{v}, \mathbf{v}'$ the inner product $\langle \mathbf{v}, \mathbf{v}' \rangle = \sum_q v_q v'_q$ is a sum of nonnegative terms, so it vanishes if and only if every term does,

$$\langle \mathbf{v}, \mathbf{v}' \rangle = 0 \iff \text{supp}(\mathbf{v}) \cap \text{supp}(\mathbf{v}') = \emptyset \quad (8)$$

Signed vectors admit no such conclusion: there orthogonality is cancellation, and two vectors with identical support can be orthogonal. The support overlap of

Algorithm 2: DictBlock in training (DictBlock.withStats). At inference the four stochastic steps (revival, winner dropout, head dropout, fiber dropout) and the feature noise are skipped. Every statistic is stop-graduated unless its loss term is enabled.

Input: logits $K \in \mathbb{R}^{b \times h \times k}$; router scales $\Gamma \in \mathbb{R}_{\geq 0}^{b \times h}$ (all ones without a router); fiber coordinates $C \in \mathbb{R}^{b \times h \times r}$ (none without fibers)

Data: dictionary $D \in \mathbb{R}^{h \times k \times e}$; fiber bases $U_{ij} \in \mathbb{R}^{e \times r}$

Hyperparameters : temperature t ; feature noise σ_F ; winner dropout p_{drop} ; revival rate p_{rev} and threshold τ_{rev} ; head dropout p_{head} ; fiber dropout p_{fib} ; entries kept per head k_{sel} (1 under ste/argmax, k under softmax, and the rule's own value under $\text{top}k_{\text{sel}}$, e.g. 4 for top4)

Output: pooled features $F' \in \mathbb{R}^{b \times e}$; assignments $P \in \mathbb{R}^{b \times h \times k}$; statistics $(L1, H, \text{cossim}_b, \mathbb{E}_i \kappa_i, \max_i \kappa_i, \text{KL}_m, \omega)$

```

begin
   $a_{ij} \leftarrow |D_{ij}|$  // atoms;  $D_{ij}$ . itself if signed
  if  $k_{\text{sel}} < k$  then // revive starved entries
     $\text{use}_{ij} \leftarrow$  share of rows on which entry  $j$  is among head  $i$ 's top  $k_{\text{sel}}$  logits
    foreach  $(n, i)$  with probability  $p_{\text{rev}}$  do
      draw  $j$  uniformly from  $\{j : \text{use}_{ij} < \tau_{\text{rev}} k_{\text{sel}}/k\}$ , if any
      raise entry  $j$ 's logit on row  $n$  to head  $i$ 's maximum // tie the leader
    end
  end
  each  $(n, i)$ :  $\max_j [K]_{nij} \leftarrow -\infty$  with probability  $p_{\text{drop}}$  // winner dropout
   $P \leftarrow \text{select}_k(K/t)$  // softmax,  $\text{top}k_{\text{sel}}$  or straight-through
   $H \leftarrow \mathbb{E}_{n,i} \text{entropy}_k(\mathbf{p}_{ni})$  // Eq. (2)
   $\bar{\mathbf{p}}_i \leftarrow \text{sg}[\mathbb{E}_n \mathbf{p}_{ni}]$  // head origin
  foreach  $(n, i)$  do // head dropout
     $\mathbf{p}'_{ni} \leftarrow \bar{\mathbf{p}}_i$  with probability  $p_{\text{head}}$ , else  $\bar{\mathbf{p}}_i + (\mathbf{p}_{ni} - \bar{\mathbf{p}}_i)/(1 - p_{\text{head}})$ 
  end
  each  $(n, i)$ :  $\mathbf{c}_{ni} \leftarrow \mathbf{0}$  with probability  $p_{\text{fib}}$  // fiber dropout
   $\mathbf{f}_{ni} \leftarrow \gamma_{ni} \sum_j \mathbf{p}'_{nij} \mathbf{a}_{ij}$  // Eq. (5); interventions act on  $\mathbf{p}'$ 
   $\tilde{\mathbf{F}} \leftarrow \text{noise}_{\text{featvar}}(\mathbf{F})$  // Eq. (18)
   $\text{cossim}_b \leftarrow \text{cossim}_b(P)$  // Eq. (3), through the Grams
   $\kappa_i \leftarrow \mathbb{E}_{j \neq j'} \cos(\mathbf{a}_{ij}, \mathbf{a}_{ij'})$  // within-head atom cosine
   $\text{KL}_m \leftarrow \mathbb{E}_i [\log_2 k - \text{entropy}_k(\bar{\mathbf{p}}_i)]$  // Eq. (4)
   $\pi_i \leftarrow \sum_j \bar{p}_{ij} |\mathbf{a}_{ij}|$ ;  $\omega \leftarrow \mathbb{E}_{i \neq i'} \cos(\pi_i, \pi_{i'})$  // support overlap
   $L1 \leftarrow \mathbb{E}_n \sum_i \|\tilde{\mathbf{f}}_{ni}\|_1$ 
   $\mathbf{f}'_n \leftarrow \sum_i \tilde{\mathbf{f}}_{ni}$  // concatenated, not summed, under ConcatDictBlock
  if fibers then // Eq. (19), then any norm cap
     $\mathbf{f}'_n \leftarrow \mathbf{f}'_n + \sum_i \gamma_{ni} \sum_j \mathbf{p}'_{nij} U_{ij} \mathbf{c}_{ni}$ 
  end
  return  $F', P, (L1, H, \text{cossim}_b, \mathbb{E}_i \kappa_i, \max_i \kappa_i, \text{KL}_m, \omega)$ 
end

```




FIGURE 4 · One trained head, block by block. Head $i=0$ of layer 0 of the straight-through model ste_h76_init01 (step 369,500, $t=0.00015$) on $n=384$ rows drawn from the eval tail; one seed. (a) The assignments P_i , rows sorted by winning entry and entries by mean usage $\bar{p}_{ij}$ (bars). “Use” is the effective number of entries across rows, $\exp H(\bar{\mathbf{p}}_i)$, and “row” the effective number within a row, the exponential of the mean per-row entropy: 31.2 and 1.0 of $k=32$, a hard code that uses nearly every entry. (b) The neighbor graph that KL_{pwak} and L2_{pwak} (Eq. (22)) diffuse over, in the same row order: the co-assignment matrix $P_i P_i^T$ times the heat kernel $\exp((\langle \mathbf{u}_n, \mathbf{u}_{n'} \rangle - 1)/0.2)$ of the layer’s shapes without their constant coordinate, self-edges removed and each row normalized to sum to one. A hard code makes it block-diagonal by cell. This model was not trained with either statistic, so the graph is the one they would build, not one that shaped it. (c) Cosines between head i ’s atoms, whose off-diagonal mean is κ_i (0.289). (d) Cosines between the heads’ coordinate profiles π (Eq. (6)) with $\bar{\mathbf{p}}$ the mean over the n rows, whose off-diagonal mean is ω (0.967), head i ’s row outlined. (e) Head i ’s atoms on the 48 of $e=2048$ coordinates where π_i is largest, a column subset rather than a projection. Each block has its own scale, white at zero, except (d): ω is near 1, so its scale starts at the 5th percentile of the off-diagonal cosines, and the few below it show as white. Figure 15 is the same head of a softmax model.

Eq. (6) sidesteps this. It compares the absolute values of the atoms, which have disjoint supports exactly when they are orthogonal whatever the atoms' signs, so it measures disjointness under a signed dictionary too; and it is computed from the atoms and the batch-mean assignment rather than from noised features, which leave the orthant and the argument with it.

The constraint binds the dictionary, not merely one sample's features. Because the assignments are a softmax, $p_{nij} > 0$ for every entry, and expanding Eq. (5) gives

$$\langle \mathbf{f}_{ni}, \mathbf{f}_{ni'} \rangle = \sum_{j=1}^k \sum_{j'=1}^k p_{nij} p_{ni'j'} \langle \mathbf{a}_{ij}, \mathbf{a}_{i'j'} \rangle \quad (9)$$

a strictly positive combination of all k^2 cross-head atom overlaps. A single sample on which heads i and i' are orthogonal therefore forces $\langle \mathbf{a}_{ij}, \mathbf{a}_{i'j'} \rangle = 0$ for every pair of atoms (j, j') : the two heads own disjoint blocks of coordinates as a property of $\mathbf{D}$, holding for all inputs rather than on the data that happened to be measured. That is the premise a static reading of the dictionary needs, and it is stronger than what any signed-feature architecture can ask for. The budget is not tight either. In the softmax configuration (Table 9), $h=32$ heads partitioning $e=2048$ coordinates leaves 64 dimensions per head to hold $k=32$ atoms, which need not be disjoint from each other within a head; head-disjointness is most of what the width $e > d$ buys.

Two different orthogonalities are in play and only the between-head one is architectural. Disjoint support assigns each head a block of coordinates; *within* a head the k atoms are contrastive alternatives spread apart inside that block (Figure 5). Section 4.3 measures the second. Neither is guaranteed at convergence: ω is a penalty, and the geometry is built to make a support partition the cheap solution rather than to force one. Measured on the trained models, the cheap solution was not taken: the within-head orthogonality holds, and no model partitions its coordinates between heads (Appendix G.2). Under ConcatDictBlock heads write disjoint slices by construction and ω is identically zero.

We add noise ε_F to F , then sum over the head dimension to obtain pooled features $\mathbf{f}' \in \mathbb{R}^e$, which are added to the accumulator $\mathbf{r}$.

$$\mathbf{f}' = \sum_{i=1}^h \tilde{\mathbf{f}}_i \quad (10)$$

Decoder. Each DictEnc output $\mathbf{f}'_{\ell}$, scaled by its input's gain g_{ℓ} , is added to an accumulator $\mathbf{r}$, whose state after ℓ layers is

$$\mathbf{r}_{\ell} = \sum_{\ell'=0}^{\ell-1} g_{\ell'} \mathbf{f}'_{\ell'}, \quad g_{\ell} := \text{sg}[\text{L2}(\mathbf{y}_{\ell})] \quad (11)$$




FIGURE 5 · DictBlock **head geometry** for $h=2, k=3$. Each head's output $\mathbf{f}_i$ lies in the simplex spanned by its atoms $\mathbf{a}_{ij}$, at the point its assignment picks. The support overlap ω compares heads: driving it to zero would give them disjoint coordinates (Section 3.1.4).

Under `resid_gain` a layer classifies only the direction of its input, so the gain carries the magnitude it cannot see; without it $g_\ell = 1$. The empty sum makes $\mathbf{r}_0 = \mathbf{0}$, so layer ℓ 's input is the encoded input itself, $\mathbf{y}_0 = \mathbf{x}$, whose gain is identically 1 on unit-norm SONAR embeddings.

A single-layer unbiased linear decoder `dec` projects $\mathbf{r}_\ell$ back to the input space $\mathbb{R}^d$. Weights $W_{\text{dec}} \in \mathbb{R}^{d \times e}$ are shared for all `DictEnc` $_\ell$.

$$\hat{\mathbf{x}}_\ell = W_{\text{dec}} \mathbf{r}_\ell \quad (12)$$

The residual $\mathbf{y}_\ell$, layer ℓ 's input, is given by $\mathbf{x} - \hat{\mathbf{x}}_\ell$, stop-graduated so that lower layers cannot write a communication code into the residual instead of reducing it. Layer ℓ splits it into the gain g_ℓ (Eq. (11)) and a shape $\mathbf{u}_\ell$, its unit direction with a constant coordinate appended, which is what the classifier, the router and the fibers read; ϵ is the floating-point epsilon of the working dtype.

$$\mathbf{y}_\ell = \text{sg}[\mathbf{x} - \hat{\mathbf{x}}_\ell], \quad \text{L2}(\mathbf{y}) := \sqrt{\sum_{q=1}^d y_q^2} \quad (13)$$

$$\mathbf{u}_\ell = \frac{\mathbf{y}_\ell}{\text{L2}(\mathbf{y}_\ell) + \epsilon} \oplus [1] \quad (14)$$

Normalizing fixes the classifier's input scale: a bilinear classifier's logits scale with the square of its input's norm, so they would explode on large inputs (raw GPT-2 activations) and vanish on the small residuals deep layers see; the gain carries the magnitude instead. A constant coordinate [1] is appended, giving a dimension of $d + 1$. This prevents the sign information of $\mathbf{y}_\ell$ from being lost by the bilinear transformation. Effectively, it adds a bias to an otherwise unbiased linear form while preserving eigendecomposition.

Noise. Noise is added to X , K , and F , controlled by hyperparameters σ_X , σ_K , and σ_F . Each draws a standard normal of the same shape as its target,

$$\varepsilon \sim \mathcal{N}(0, 1) \quad (15)$$

and scales it three different ways. Classifier logits take unscaled Gaussian noise,

$$\text{noise}_{\text{normal}}(\mathbf{K}) := \mathbf{K} + \sigma_K \varepsilon \quad (16)$$

inputs take noise proportional to their own norm, so the scale is a fraction of signal magnitude rather than an absolute quantity,

$$\text{noise}_{\text{batchnorm}}(X) := X + \sigma_X \varepsilon \odot \frac{\text{sg}[\text{L2}(X)]}{\sqrt{d}} \quad (17)$$

and features take noise proportional to each feature’s standard deviation across the batch.

$$\text{noise}_{\text{featvar}}(\mathbf{F}) := \mathbf{F} + \sigma_F \varepsilon \odot \text{sg}\left[\sqrt{\text{var}_n(\mathbf{F}) + \epsilon}\right] \quad (18)$$

Here $\text{L2}(X)$ is the per-sample norm of Eq. (13), broadcast over d . Both scales are stop-gradiented: without sg the model could evade its own noise by shrinking feature variance, and the ϵ inside the root guards the $0/0$ gradient where a feature’s batch variance underflows to exactly zero. Because σ_K enters the softmax as σ_K/t , holding it constant under a falling temperature would grow effective exploration rather than hold it, so models with an annealed temperature anneal σ_K with it. Each model’s noise settings are in Appendix B.3.

Router and fibers. Two optional branches read the same shape $\mathbf{u}$ as the classifier (Eq. (14), with the layer index dropped) and meet it in the dictionary lookup (Figure 6). Both are optional; the models that use them are listed in Appendix B.3.

The *router* (scaled) is a second bilinear map, from $\mathbf{u}$ to one scale per head, $\gamma \in \mathbb{R}^h$, made non-negative by an absolute value (`router_signed` drops it). Its scale γ_i multiplies head i ’s whole output, so the classifier decides *what* a head says and the router *how much*; its own L1, L1_S , prices the router alone, where L1 on the output can be paid down by either factor.

Fibers (`fiber_rank` r) give every entry a local basis $U_{ij} \in \mathbb{R}^{e \times r}$ (in its head’s slice under `ConcatDictBlock`), and each head reads r coordinates $\mathbf{c}_i = \beta \tanh(A_i \mathbf{u} / \beta)$ from the input, so

$$\mathbf{f}_i = \gamma_i \sum_{j=1}^k p_{ij} (\mathbf{a}_{ij} + U_{ij} \mathbf{c}_i). \quad (19)$$

Under a hard code the selected entry is a discrete base point and the fiber a continuous correction within the plane its basis spans (Figure 7). Unconstrained, the correction can take over and leave the base meaningless, so the base is kept honest by fiber dropout, a norm cap ($\|U_{ij} \mathbf{c}_i\| \leq \rho$ per head, or summed over heads), a base-only deep-supervision term (`base_aux`, which also decodes the fiber-free




FIGURE 6 · DictEnc with router and fibers. The gain–shape split hands the unit input $\mathbf{u}$ to three branches: the classifier (which entry, P), the router (how much, γ , one non-negative scale per head) and the fiber read (where within the entry’s fiber, C , r coordinates per head). They meet in the lookup (Eq. (19)), and the combined heads are scaled by the measured gain g . Dashed boxes above an arrow act on it in training; dashed boxes below are statistics of the fiber-free base output.

output), or training the fibers after freezing the base. A radial fiber ($r=1$, no basis) only rescales the selected entry. Statistics describe the base output. Head dropout (`p_head_drop`) is the third addition here: in training it replaces a head’s classification by the batch mean, its origin, so heads cannot co-adapt and ablating one to its origin stays in distribution.

3.2 TRAINING

Embedding model. We embed texts using SONAR [19], a multimodal and language-agnostic sentence embedding model. SONAR embeddings have several desirable properties:

- They correspond to entire texts rather than tokens.
- They are bidirectional.
- Language is encoded independently of text.

This allows us to validate steering fidelity against a ground truth. Large Concept Models [13] generate directly in SONAR space, which makes it a natural substrate for steering experiments.

Loss function. Reconstruction fidelity is trained with MSE, where $\mathbf{w} \in \mathbb{R}^d$ is a fixed per-dimension weight vector.

$$\text{MSE}(X, \hat{X}) := \frac{1}{bd} \sum_{n=1}^b \sum_{q=1}^d w_q (\hat{x}_{nq} - x_{nq})^2 \quad (20)$$




FIGURE 7 · **Fiber geometry** for one head, extending Figure 5. The head selects the entry with atom $\mathbf{a}_{i2}$; its fiber is the plane spanned by U_{i2} at that point, and the coordinates c_i place the output within it. The dashed circle is the per-head norm cap; the dotted ray is the radial variant, which moves the output only along the atom's own direction.

The total objective adds the ghost-gradient reconstruction and the auxiliary statistics, each weighted by a scalar. Statistics are *summed* over the l layers rather than averaged.

$$L := \text{MSE}(X, \hat{X}) + s_g \text{MSE}(\hat{X}_g, X - \hat{X}) + \sum_{\ell=0}^{l-1} \mathbf{s} \cdot \boldsymbol{\psi}_\ell \quad (21)$$

Each layer contributes the statistics (Algorithms 2 and 1)

$$\boldsymbol{\psi}_\ell = \left[\text{L1}_K \text{ L1 } H \text{ cossim}_b \mathbb{E}_i \kappa_i \max_i \kappa_i \text{KL}_m \text{KL}_{\text{pwak}} \text{L2}_{\text{pwak}} \text{L1}_S \omega \right]_\ell^\top \quad (22)$$

where L1_K is the L1 of the classifier's gate, κ_i the mean within-head cosine between head i 's atoms, KL_{pwak} and L2_{pwak} the two partition-consensus statistics, and L1_S the router's L1. The max of κ_i is logged for a setpoint controller and never itself penalized. Appendix B.4 gives the implementation's row layout.

Each model's weights are in Appendix B.3. The ghost path is disabled throughout, since $\hat{X}_g$ fires only on features that are exactly zero across the whole batch, which bilinear classifiers with softmax selection and absolute-valued dictionaries never produce.

Reconstruction is trained with deep supervision: $\text{MSE}(X, \text{dec}(R_\ell))$ is computed separately for every prefix $\ell = 1, \dots, l$ and averaged, so each prefix carries weight $1/l$ and each layer must explain what previous layers left unexplained. A stop-gradient variant (deepsup_sg, off in every reported model) stops each prefix's gradient at the accumulator it was handed, so that a layer trains only against its

Algorithm 3: Ontologizer in training (Ontologizer.withStats) and the loss it feeds (Hyperparams.loss), for the residual forward (forward="resid") every configuration here uses. Without an encoder, enc is the identity; with direct atoms, dec is.

Input: batch $X \in \mathbb{R}^{b \times d}$

Data: encoder enc ; layers $\text{DictEnc}_0, \dots, \text{DictEnc}_{l-1}$ (Algorithm 1); decoder dec (Eq. (12)); MSE weights $\mathbf{w}$; loss weights $\mathbf{s}, s_g$

Hyperparameters: input noise σ_X ; deep supervision; deepsup_sg; resid_norm, resid_const, resid_gain

Output: loss L ; per-layer statistics $\psi_0, \dots, \psi_{l-1}$

begin

$E \leftarrow \text{enc}(\text{noise}_X(X))$ // Eq. (17)

$R \leftarrow \mathbf{0} \in \mathbb{R}^{b \times e}$ // state 0: no layers applied

$E \leftarrow E \oplus [1]$ // if resid_const

for $\ell \leftarrow 0$ **to** $l-1$ **do** // layer ℓ reads state ℓ , writes state $\ell+1$

$R_{\text{prev}} \leftarrow R$

$R, \mathbf{P}_\ell, \psi_\ell \leftarrow \text{DictEnc}_\ell(R, E)$

if *deep supervision* **then** // a decode per prefix

$\hat{X}_{\ell+1} \leftarrow \text{dec}(R - R_{\text{prev}} + \text{sg}[R_{\text{prev}}])$ // dec(R) without deepsup_sg

end

if $\ell < l-1$ **then**

$Y \leftarrow \text{sg}[X - \text{dec}(R)]$ // Eq. (13)

$E \leftarrow Y \oplus [1]$ // if resid_const; the next layer splits it into gain and shape, Eq. (14) (normalized here instead under resid_norm alone)

end

end

if not *deep supervision* **then** $\hat{X}_l \leftarrow \text{dec}(R)$

$\text{MSE} \leftarrow \mathbb{E}_\ell \text{MSE}_\mathbf{w}(X, \hat{X}_\ell)$ // Eq. (20); over the prefixes computed, each weighted $1/l$

$L \leftarrow \text{MSE} + s_g \text{MSE}_g + \mathbf{s} \cdot \sum_{\ell=0}^{l-1} \psi_\ell$ // Eq. (21); $\text{MSE}_g = 0$ without ghosts

return $L, (\psi_\ell)_{\ell=0}^{l-1}$

end

own residual. The MSE is whitened: $\mathbf{w}$ in Eq. (20) is the per-dimension inverse variance, normalized to mean one, the same whitening the evaluation applies. The implementation applies $\sqrt{\mathbf{w}}$ to X and $\text{dec}(R_\ell)$ before the squared error, which is equivalent. Four statistics are active as loss terms in the reported models (Figure 3; Appendix B.3): between-sample similarity cossim_b , computed from P through the dictionary Gram so it measures feature-space similarity of the per-head reconstructions, and which in practice also minimizes per-sample entropy H ; the batch-mean term KL_m , anchoring each head’s mean assignment vector to a uniform distribution so that a uniform code corresponds to a maximally generic feature; an L1 penalty on features; and a between-head disjointness term, the support overlap ω , which measures the extent to which heads use nonoverlapping coordinates of the dictionary space. Most reported models trained before ω existed and used its predecessor (Appendix B.4). L1 tends to fall under MSE alone; including it in the loss speeds up convergence. Classification sharpness is controlled by the temperature. Winner dropout and the noise terms ε_K and ε_F regularize classifications and features respectively.

3.3 EVALUATION

Evaluation runs against TopK SAEs trained on identical cached SONAR embeddings under an identical whitened reconstruction objective. The battery is deliberately SAE Bench-style [15]: a reconstruction-versus-capacity Pareto; categorical-form and semantic-coherence tests of head structure; cross-model feature splitting and seed stability; downstream text fidelity through the paired M2M100 decoder; sparse language probing; shrinkage decomposition; steering fidelity; and blinded autointerp scoring. A structural-ablation ladder of SAE variants, each adding one Ontologizer-style commitment at a time, attributes effects to specific structure. AxBench [17] found that simple supervised baselines beat SAEs at steering; the steering comparison here is built to face that result.

3.4 DATA, MODELS AND EVALUATION

This section is an overview; Appendix B is the full description, with every evaluation’s defining equation or algorithm.

Data. The SONAR models train on a fixed cache of 4,000,000 unit-norm SONAR sentence embeddings ($d=1024$) from an interleaved 86-language mC4 sample, with a language label per row and text recoverable by row index. Sections 7.1 and 5.5 also use GPT-2 [20] layer-8 residual-stream activations ($d=768$), a substrate where sparse dictionary learning is known to work; Section 6’s separate GPT-2 setup compares only within that section (Appendix B.1).

Objective and score. Every model minimizes the whitened MSE of Eq. (20), and we report whitened FVU (Eq. (23)) on the final 32,768 cache rows, which SAEs and straight-through arms never train on. The softmax Ontologizers predate that holdout, so their scores are in-sample. Every score uses the run’s own final temperature.

Models. The reference softmax Ontologizer `resid_nc` has $l=5$ residual layers of $h=32$ heads with $k=32$ entries; a selection sweep varies only its selection rule, and the straight-through arm (Section 5) replaces it with hard codes. The SAEs are TopK SAEs after Gao et al. [21], plus a structural-ablation ladder that adds one Ontologizer constraint at a time. Two of its rungs, `g160top1` and `g160softmax` (160 heads of hard or softmax competition over a shared decoder), are each a one-layer Ontologizer without the bilinear classifier, residual stack or non-negative dictionary, and are treated as simplified variants rather than baselines (Appendix B.2; configurations in Appendix B.3).

Evaluations. Each is a standalone script on the shared eval split (Appendix B.5): reconstruction against code capacity and the rate-distortion floor (`pareto.py`); whether SAE latents form exhaustive, exclusive head-like groups (`headstruct.py`); the semantic coherence of a head’s atoms against size-matched nulls (`headcoh.py`); feature splitting and cross-seed head matching by partition and by contribution (`splitting.py`, `partition.py`, `headcontrib.py`); text fidelity through the M2M100 decoder (`textfid.py`); language probing by latent and by partition (`langprobe.py`, `headlang.py`); shrinkage by least-squares refitting on a frozen support (`refit.py`); steering fidelity in text and in embedding space (`steerfid.py`, `steerembed.py`, `steergeom.py`); additivity of paired interventions (`compose.py`); and autointerp (`autointerp.py`).

4 RESULTS: THE SOFTMAX ONTOLOGIZER AGAINST SAEs

4.1 THE PRICE OF STRUCTURE: RECONSTRUCTION-CAPACITY PARETO

The SAE frontier dominates the soft Ontologizer’s code at every capacity measured (Figure 8b). Truncating the reference model’s code to the v largest deviations per head from its mean assignment (Appendix B.5), it needs 320 coefficients to reach the $FVU_w \approx 0.49$ that the $k_{SAE}=32$ SAEs reach with 32, and its coarsest code, one deviation per head (160 coefficients, $FVU_w 0.677$), is already worse than the SAE’s 32-coefficient point, so the very sparse regime is closed to it. Unstructured reconstruction is better still: the dense control reaches $FVU_w 1.1 \times 10^{-4}$ with $3.4\times$ fewer active coefficients than the soft code, about $50\times$ lower error at a matched budget. The truncations are applied after the fact to a model never trained to truncate, so they may overstate what the structure itself costs. The code’s




FIGURE 8 · Reconstruction against code size on the shared evaluation tail, by index bits (a) and by continuous coefficients (b): the reference softmax model’s truncated codes (resid_nc), the SAEs (the L1 SAEs at all five weights of Table 13), the single-layer variants and the straight-through stack (ste_h76_init01, Section 5). A code appears only where its size is nonzero, so the soft codes are in (b) alone, the softmax argmax is in (a) alone, and the straight-through stack is a level in (b). The dashed curve is the Gaussian reference (Eq. (25)). The softmax model saw the tail in training and the others hold it out. One seed per point, except the $k_{SAE}=32$ SAE’s seed-43 replica.

information is spread across heads: keeping whole heads instead of deviations is far weaker at matched coefficients (FVU_w 0.50 for 8 of 32 heads per layer, against 0.094 at $v=8$).

The soft model’s own discrete reading is not a code. Its argmax reconstructs at FVU_w 17.3, 77 $\times$ the Gaussian reference for 800 bits (0.226; Figure 8a), the least distortion an 800-bit code could reach if the embeddings were Gaussian with their measured covariance (Eq. (25)). Real embeddings are no harder to compress than that, so the argmax is at least 77 $\times$ from optimal, and on a selection-sweep run of the same shape the gap is nine orders of magnitude (8.7×10^6 against 0.0007 for the soft code). The entries that interventions on a softmax model act on are therefore not what carries its embedding, which motivates the straight-through arm of Section 5. The reference does not say that 800 bits are too few, since non-Gaussian data can be compressed below it.

4.2 STRUCTURE DECOMPOSITION (THE LADDER)

rung	FVU_w	L0	note
k32 baseline	0.499	32	0 dead
g160top1	0.274	160	1434/5120 dead (28%)
g160softmax	0.0044	5120	beats resid_nc (0.0054)
k32_p5	0.509	32	nested prefixes, $\sim 2\%$ cost
k32_bl	0.602	32	1848/5120 dead (36%)

TABLE 1 · Converged structural rungs, $m=5120$, 144.3k steps each.

Per-head competition sets most of the gap to unstructured codes: the single-layer soft variant costs about 30 $\times$ the dense control (0.0044 against 1.1×10^{-4}). Everything the reference Ontologizer adds to it (bilinear classifiers, five residual layers, a shared 2048-dimensional dictionary space) adds nothing to reconstruction, at 2.2 $\times$ the parameters. A selection-sweep run of the same shape does reach 0.0007, six times below the single-layer variant, but both softmax runs saw the evaluation tail in training, so whether the stack beats it out of sample is open. Reconstruction alone does not justify the stack; steering has to.

Hard per-head selection costs only about 6% over an unstructured code of equal sparsity (0.274 for the single-layer hard variant against 0.258 for the $k_{SAE}=160$ SAE). The softmax model’s catastrophic argmax is therefore a training artifact, not a property of discrete head codes: one-hot selection inflates its output norm 3.7 $\times$, because winning atoms carry about 4 $\times$ the soft mixture’s norm. A bilinear encoder costs about 20% in FVU_w and leaves a third of the dictionary dead.

Categorical structure appears where it is trained in as hard selection. The hard variant’s partition is head-like, while soft per-head codes converge to dense, near-uniform mixtures and the SAE’s spontaneous groups are co-firing topical clusters (Appendix C.1). The rest of the softmax-versus-SAE battery (feature




FIGURE 9 · Within-head decode-direction coherence: each head’s, or discovered group’s, ζ_{cos} against a size-matched null drawn from the same pool ($n_{\text{null}}=500$), on a scale linear between the dotted lines at ± 2 and logarithmic beyond them; bars are means. The soft heads are those of `resid_nc` and `g160softmax`, the discovered groups those of the $m=11264$, $k_{\text{SAE}}=32$ SAE. One seed per model.

reliability, text fidelity, language localization and shrinkage) is in Appendix C.

4.3 SEMANTIC COHERENCE: HEADS ARE CONTRAST SETS

Hard-trained heads are coherent, and soft-trained heads are contrast sets (Figure 9). Every head of the single-layer hard variant decodes to directions far more similar than a random set of the same size: winner-take-all carves the data so that each head’s atoms tile one region. An SAE’s post hoc co-firing groups are mildly coherent, as topical clusters would be. Both soft architectures show the opposite: in 96–100% of heads, the entries’ decode directions are more dissimilar than random sets. Softmax competition pushes a head’s atoms apart, so its entries are well-separated alternatives, and swapping one for another is a large, clean move along the head’s own dimension. The repulsion does not depend on the inner product: under the whitened metric the objective optimizes, it strengthens for the Ontologizer (mean $\zeta_{\text{cos}} -24.4$ against -17.6) and is essentially unchanged for the single-layer variants (Appendix C.2).

The hard-code Ontologizer shows the same in its weights. Its within-head atom cosine, decoded into output space, is 0.023 at layer 0 and -0.015 to -0.020 at layers 1–4, so a head’s decoded atoms are orthogonal and slightly repelling; read in the dictionary space instead, the statistic is dominated by the decoder’s null space (Appendix G.1). Of the two orthogonalities in Section 3.1.4, the within-head one holds in both code types and the between-head one holds in neither: no trained model partitions its dictionary space between heads. The support overlap ω sits at or above a null that shuffles each head’s coordinates, and each head’s usage covers 52–98% of the coordinates, where a partition into 32 heads needs about 3%

(Appendix G.2). Heads are contrast sets that share their coordinates.

4.4 STEERING FIDELITY

Steering is measured in embedding space: a step from the input along a direction, scored by whether the target entry becomes selected (its realization) and by how many other heads change (its collateral; Algorithm 4). The text round trip cannot resolve an intervention, because even an unsteered round trip preserves a head’s argmax on only about 12% of rows (Appendix C.7).

direction	hard code @0.25 / 1.0	softmax @0.25 / 1.0
decode	0.364 / 0.607	0.824 / 0.569
grad	0.304 / 0.103	0.453 / 0.477
margin	0.158 / 0.275	—
adjoint, oriented	0.001 / 0.001	0.427 / 0.501
adjoint, oriented per sample	—	0.446 / 0.512
random	0.018 / 0.024	0.020 / 0.037
collateral	0.66 / 0.88	0.58 / —

TABLE 2 · Realization rate (target entry newly selected) by direction and strength; 60–64 entries $\times$ 32 held-out rows, chance 0.031. Collateral agrees within 0.01 across directions at a given strength, so rows compare at an equal budget. Hard code: the 5×76 stack, whose seed-43 twin agrees within 0.015 in every decode, grad and random condition. The adjoint direction’s orientation is in Appendix B.5.

In embedding space steering works. The decode direction realizes the target entry on 82% of rows for the softmax model at strength 0.25. The hard code is the less steerable of the two: it realizes 36% while already disturbing more of the other heads (0.66 against 0.58), because a hard argmax has margins a small step cannot cross. Swept over strength, its decode direction saturates near 0.65 (Figure 10b), so a third of its interventions cannot be realized at any strength, and it passes the soft model only beyond collateral 0.85, where most of the other heads have moved too. Steering by target layer reverses the aggregate ranking of the decode and gradient directions (Appendix E.4).

Against supervised steering vectors. The same measurement scores the two standard supervised steering vectors for each target: the difference of means between rows that select the target and rows that do not (activation addition [22, 23]), and the weight vector of a class-balanced logistic probe, both fitted on 8,192 reference rows with the steered rows held out. Their labels are the code’s own selections, so they are an oracle baseline: how well a supervised method that already knew the partition would steer into it.

Only the softmax stack’s own direction beats the supervised vectors, reaching its categories on 75% of rows against 46–47%. The hard-code stack’s categories

model	own decode	diff. of means	probe	random
softmax stack	0.75 / 0.57	0.46 / 0.46	0.47 / 0.51	0.02 / 0.56
hard-code stack	0.37 / 0.66	0.91 / 0.63	0.94 / 0.64	0.02 / 0.65
hard-code flat	1.00 / 0.33	0.97 / 0.27	1.00 / 0.31	0.01 / 0.37
g160top1	0.58 / 0.03	0.94 / 0.13	0.95 / 0.14	0.04 / 0.23
g160softmax	1.00 / 0.30	0.99 / 0.27	1.00 / 0.31	0.01 / 0.34
TopK SAE, $k_{\text{SAE}}=32$ ($m=11264$)	1.00 / 0.21	0.92 / 0.21	0.97 / 0.23	0.00 / 0.25
SAE, bilinear encoder	1.00 / 0.07	0.97 / 0.12	0.99 / 0.11	0.01 / 0.13

TABLE 3 · Each model’s own steering direction against supervised vectors fitted to the same targets (realized / collateral), at a step of $0.25\|x\|$; 64 targets $\times$ 32 held-out rows per model, one seed. An SAE’s collateral counts its other active latents, so compare within a row.




FIGURE 10 · Realization against collateral as the step grows from $0.25\|x\|$ to $4\|x\|$, left to right along each curve, on the draw of Table 3. (a) The softmax stack (sweep_softmax_shm). (b) The hard-code stack (ste_h76_init01). (c) The flat hard code (ste_l1_h380_i01_hm1e4). (d) The single-layer hard variant g160top1. At each strength the models’ own directions disturb as many other heads as the random step, to within about 0.05, except g160top1’s decode direction, which disturbs far fewer; what a direction buys is its height above the random curve. The logistic probe tracks the difference of means (within 0.03 realized, 0.05 collateral) and is not drawn.

are reachable, but its own direction misses them: the supervised vectors reach the same entries on 91–94% of rows, where its decode direction reaches 37%. The single-layer hard variant trades reach for precision (58% realized against 94%, at a fifth of the collateral); every other model’s own direction disturbs the other heads as much as a random step does. Knowing the partition buys the supervised vectors little precision (the difference of means sits at most 0.1 below random’s collateral), and an SAE’s decoder direction is already as good as the oracle. At each method’s own step size the difference of means realizes more in most models, partly by taking longer steps (Appendix E.5). This tests reaching the model’s own categories; steering toward an external concept, against supervised vectors fitted to that concept’s labels, remains to be run.

A single entry can still steer content: on the reference softmax model, forcing entry 17 of head 8 in layer 1 reliably yields Burmese decodes of otherwise preserved content.

4.5 DESCRIPTION SPECIFICITY (AUTOINTERP)

Individual SAE latents are about twice as describable as individual entries. Under the same plain summarization prompt, 51% of $k_{\text{SAE}}=32$ latents receive a specific description (20% naming a language or script), against 23% of the reference model’s entries (8%). Blinded detection over the full campaign agrees:

model	density	freq	params	eig	acts	cacts	features
onto (resid_nc)	100%	10.80%	0.020	—	0.473	0.392	225
k32	0.6%	0.61%	0.158	—	0.645	0.603	250
k32_bl	0.6%	2.62%	0.061	0.021	0.578	0.470	75
g160top1	3.1%	5.08%	0.039	—	0.646	0.594	150
g160softmax	100%	0.00%	0.025	—	0.449	0.364	256
k5120 (dense ctrl)	45.5%	17.97%	0.001	—	0.470	0.332	256

TABLE 4 · Mean per-feature detection F1 by description mode (Appendix B.5). The always-match null is about 0.67, and a generic description that matches every item scores the null, so generic descriptions inflate toward it. Density is nominal code density; freq is the firing rate at the $2/k$ threshold. k32_bl’s pool is only 75 live features; no null control was recorded for this campaign.

Describability tracks code density, not architecture. Activation-based descriptions rank the sparse codes (the $k_{\text{SAE}}=32$ SAE and the single-layer hard variant, 0.65) above the bilinear SAE (0.58), the reference model and the dense control (0.47) and the single-layer soft variant (0.45), and density alone carries about 80% of that variance (Figure 11a). The reference Ontologizer sits on the SAEs’ density trend, while the hard variant lies 0.12 above it, unusually describable for its density (Figure 11b; Appendix D.6). Per-latent detection F1 therefore cannot compare architectures at unmatched density, and every soft Ontologizer is dense




FIGURE 11 · Mean detection F1 per run of Table 4 against code density. (a) Activation descriptions against nominal density, the share of a code’s coefficients that are nonzero. (b) Contrastive descriptions against the firing rate at the $2/k$ threshold, with a line fitted in log rate to the four sae.py runs that fire; gl60softmax never crosses the threshold and is not drawn, and the Ontologizer (resid_nc) is held out of the fit. Detection sets are half positives, so a description that matches every item scores the always-match null, $2/3$ (dotted). One seed per run.

by construction. The contrastive descriptions (cacts) give the honest number: they score below plain summaries everywhere, least for the sparse codes (by 0.04–0.05) and most for the dense control (by 0.14), so part of every plain-summary score is generic matching.

Parameter-space decodes are not descriptions. Decoding the parameters scores at most 0.06 everywhere except the plain SAE (0.158), and decoding bilinear eigenfeatures fails the same way (0.021). Part of the failure is how the decode is taken: an atom’s cosine to its head’s mean atom is 0.849, so a raw decode returns one template sentence with the details shuffled. Subtracting the origin decode makes the decodes legible but not detectable (on one head, 0.052 against a 0.014 null, where contrastive activation descriptions score 0.479 against 0.273). A decoded sentence is an instance of what an entry writes, and detection needs a criterion, so naming requires activation evidence.

5 A HARD-CODE ARM: STRAIGHT-THROUGH SELECTION

The softmax model’s discrete reading is not its code (Section 4.1), so its argmax entries are not what interventions act on. Under straight-through selection the forward pass is the argmax, and no such gap can open: at every checkpoint scored, the hard and soft-forward reconstructions agree to four decimals.

5.1 CAPACITY, AND WHY THE HEAD COUNT RISES

A hard code transmits exactly $l \cdot h \log_2 k$ index bits. The Gaussian reference (Eq. (25)) gives the bits a code would need to reach a given error if the embeddings were Gaussian; real embeddings may need fewer:

FVU _w target	bits	heads at $k=32, l=5$
0.226	800	160 (the softmax architecture)
0.05	1,898	380 (this arm, $h=76$)
0.0023	4,167	833

Heads are the cheap axis: bits and parameters both grow linearly in h , and heads run in parallel where layers run in sequence. The arm has 5×76 heads of 32 entries, 1900 bits and 51.93M parameters. Training a classifier under a hard forward takes three settings, all because the argmax is scale-invariant and nothing pushes the logits up: logit noise relative to the logit scale, a surrogate temperature matched to that scale, and a rescaled s_{KL_m} (Appendix E.1; every arm’s configuration is in Table 10).

5.2 INITIALIZATION: A THIRD OF THE ERROR

The dictionary is non-negative, so a layer’s h rows add coherently and its output grows with h . The per-layer gain compounds that, so at the default initialization the residual grows geometrically with depth: the first arm started at 1.84×10^8 where it should start near 1. Rescaling each dictionary at initialization so that a layer writes about a tenth of its input (`-dict-init-scale 0.1`) divides h out and starts the cascade as a near-identity:

arm	shape	s_{KL_m}	realized bits	FVU _w	params
ste_h76 (shipped init)	5×76	10^{-4}	1714 (90.2%)	0.2293	51.93M
ste_h76_init01	5×76	10^{-4}	1895 (99.7%)	0.1535	51.93M
ste_h76_i01_s43 (seed 43)	5×76	10^{-4}	1895 (99.7%)	0.1535	51.93M
ste_k2_h380	$5 \times 380, k=2$	5×10^{-4}	1887 (99.3%)	0.1844	17.67M
ste_l1_h380_i01_hm1e4	1×380	10^{-4}	1898 (99.9%)	0.2750	51.93M
ste_l1_h380_i01	1×380	5×10^{-4}	1898 (99.9%)	0.2870	51.93M
ste_l1_h380_i01_s43	1×380	5×10^{-4}	1898 (99.9%)	0.2871	51.93M

TABLE 5 · Straight-through arms at the fixed initialization (except the first row), held-out tail. Realized bits are the summed entropy of each head’s argmax usage over the tail, against the nominal $l h \log_2 k$; the sum bounds the information the code carries per sample from above, with equality only when heads are independent (Eq. (34)).

One scalar takes a third off the held-out error and has every head using its entries almost uniformly, the largest single effect in the straight-through work. It also closes the utilization gap that rescaling the entropy pressure and whitening

the input each closed in part (Appendix E.1). The hard-code results below use the fixed initialization, except the steering-by-layer and head-health measurements in Appendix E.4.

A purely discrete code beats a sparse linear one. On one frontier of converged SAEs and the single-layer hard variant, all trained for about 144k steps (Figure 8a), the stack’s 1900 index bits reach FVU_w 0.154 with no continuous coefficients. At matched index bits (about 2,000) that is 40% below the TopK SAE that additionally transmits 160 continuous coefficients (0.258), and 30% below an L1 SAE spending $3.4\times$ the bits plus 519 coefficients (0.219). It is also 44% below the single-layer hard variant (0.274), which shares its hard per-head selection but has one layer of linear classifiers and keeps a ReLU magnitude on each winner. For a hard model the capacity axis is $l \cdot h \log_2 k$ itself: truncating toward the origin only adds mass the model was never trained to decode (0.259 at $v=1$). The stack is less efficient per bit than a single layer, $3.08\times$ the Gaussian reference at its rate against $1.53\times$ for one layer’s 380-bit prefix, and bits are not the whole capacity: 380 binary heads, at the same bits, reach 0.1844 (Appendix E.2).

5.3 DEPTH AT FIXED CAPACITY

Five layers of 76 heads and one layer of 380 have the same nominal capacity (1900 bits) and exactly the same parameter count (51.93M; the dictionary hke and the bilinear classifier $2hkd$ depend on l and h only through $l \cdot h$). At the fixed initialization both use their entries almost uniformly (summed head entropy 99.7% and 99.9% of nominal), and depth is worth $1.79\times$ the reconstruction error (0.1535 against 0.2750). Seed does not explain the gap: replicas agree to four digits. Summed head entropy bounds what each code carries from above, and whether the two codes carry different amounts jointly is not measured. Neither arm has converged, and the ratio is narrowing slowly, but the stack leads at every checkpoint (Appendix E.3).

What depth buys is not usage, input conditioning or the per-layer gain. A flat arm with a whitened input uses every entry uniformly and still reconstructs $1.7\times$ worse, and pinning the gains costs the stack only 2.1% (Appendix E.3). Blending each layer’s subtracted reconstruction toward another row’s, without retraining, shows instead that what matters is pairing each residual with its own sample: the error rises well before the pairing is fully severed. A flat layer is a product quantizer, in which no head sees another’s error; the stack is a residual quantizer, each stage coding what earlier stages missed.

5.4 ENTANGLEMENT: THE COST OF THE CASCADE

A single-head intervention in embedding space realizes its target on almost every row of the flat model (99.6%) while disturbing a third of the other heads; in the

stack it realizes 37% while disturbing 66%. Restricting the stack to its layer 0, so that the target is as close to the input as in the flat model, widens the gap (realized / collateral at strength 0.25, and at each model’s native step for the last row; 12 layer-0 targets in the stack, 64 in the flat arm):

direction	stack, layer 0	flat
decode	0.049 / 0.680	0.996 / 0.328
grad	0.281 / 0.656	0.391 / 0.380
margin	0.034 / 0.661	0.076 / 0.379
random	0.008 / 0.655	0.005 / 0.368
native decode	0.044 / 0.678	0.169 / 0.043

Collateral is not a property of the intervention. In both models every direction disturbs as many other heads as a random step of the same length, to within 0.05 (Figure 10b, c), because the heads sit almost on their decision boundaries: the median margin between winner and runner-up logits is 1.5×10^{-4} , at layer 0 and in the flat arm alike, so a step of a quarter of the input’s norm crosses many of them. Linearizing each head’s logits with the upstream codes held fixed separates two sources. The stack’s layer-0 heads flip directly about as often as the flat arm’s (0.22 against 0.24), so what the stack adds is downstream layers reclassifying a changed residual: the cascade roughly doubles collateral but is not its only source. At the model’s own step size (the forced entry’s output change, unscaled) the flat arm’s step is small (a median 0.028 of $\|x\|$) and moves 4% of the other heads, while the stack’s (0.13) realizes 8% of targets across layers and moves 54%, as a random step of that length does. This is the steering counterpart of the composition result, in which non-additivity decays with the number of layers below an intervention (Section 8). Steering by target layer, the head-health levers tried against entanglement (only a row-collinearity setpoint bought steerability), and a head-independence penalty that turned out to measure estimator bias are in Appendices E.4 and E.6.

5.5 SEED REPRODUCIBILITY: THE FUNCTION IS DETERMINED, THE HEADS ARE NOT

Two runs differing only in seed, both at the fixed initialization, reach identical held-out error: 0.1535 and 0.1535 for the stack, 0.2870 and 0.2871 for the flat arm (Table 5). Nothing that identifies an individual head agrees:

what is compared	stack	flat	null
head partitions (mean matched NMI)	0.0121	0.0177	0.0048
head contributions (centered cosine)	0.0074	—	0.0015
per-layer subspace, rank 8	= data’s	—	see text
per-head atom Grams	1.05–1.11× chance	—	within-layer pairing




FIGURE 12 · The GPT-2 concatenated-head seed pair by layer: what each head separates (a, partition NMI on 32,768 rows) and what it writes (b, centered contribution cosine on 8,192 rows, log scale), across the two seeds and against the first run’s own checkpoint at step 350,000. Dotted lines are the nulls; the band in (b) runs from the null’s mean to its maximum. One seed pair.

Partitions agree barely above the null, where a run scores 0.38 against its own earlier checkpoint. Contributions agree no better. Per-layer subspaces agree exactly as much as each model agrees with the data’s own principal directions, and per-head Grams match at chance even against the same run’s earlier checkpoint (Appendix E.7). What reproduces is aggregate or layer-level, and what identifies an individual head, with one exception, does not. Claims about individual heads, including the Burmese entry and the per-head steering tables, therefore describe one seed. Non-negativity and the layer-dependent gain put structure into the floor of every such measure, and several of them, among them uncentered contributions and subspace overlap, read as agreement until compared against a null of the same shape.

Layer 0 is the exception in every head-level measure, as the only layer whose input does not depend on upstream choices, and the one head that reproduces is there: the topic head of Section 5.6, the strongest layer-0 partition in both seeds at different indices, matched across them at NMI 0.41 against a 0.01 median for other heads. The softmax Ontologizer fares no better: a seed replica at matched training matches at most 0.1% of entries, where the $k_{\text{SAE}}=32$ SAE pair matches 3–31% (Figure 21).

Replication on GPT-2. A GPT-2 pair differing only in seed (5×20 heads, $k=256$, concatenated heads at $d_{\text{head}}=128$, 371,900 steps) gives the same picture. Each measure is run across seeds and, as a positive control, against the first run’s own checkpoint at step 350,000 (Figure 12).

Only layer 0 reproduces, at under a fifth of the control, and each of its heads

matches into the other seed's layer 0 on both measures. Layer 1 keeps its identity (19 and 20 of its 20 heads land in the other seed's layer 1) while agreeing on almost nothing it separates or writes; layers 2–4 lose even that. The partition null is high, because plug-in NMI between two 256-way partitions of 32,768 rows is dominated by bias, so it cannot resolve the deep layers; against the tight contribution null they sit two to four times above it, at about 1% of the control, which is real and negligible. The excess over the null is larger than on SONAR (24× against 5–11×), almost all of it in layer 0.

What layer 0's lead is made of. Three findings, detailed in Appendices E.8 and E.9. Layer 0's larger contributions are all gain: it receives the raw input, while its atoms are no larger than layer 1's (Table 18), and at equal atom spread it still reproduces 14× better on GPT-2 and 23× on SONAR. The stack costs reproducibility at every layer: against flat-arm heads of the same size, which see the same input with nothing downstream, the stack's layer 0 reproduces about 8× worse and its deeper layers 19–49× worse, the reproducibility side of entanglement. Within layer 0 the atoms decide: layer-0 entries are unequal by atom norm rather than usage (Table 19). On SONAR one head (54) holds 8 of the layer's 11 atoms above three times the median, reproduces at 0.844 against 0.249 for the next, and drives the within-layer correlation between agreement and largest atom (Table 20); without it the layer-0 mean falls only from 0.080 to 0.070. GPT-2's layer 0 has no such head: every head carries a graded tail of large atoms, and agreement is narrowly spread.

5.6 HEADS AGAINST A KNOWN FACTOR: LANGUAGE, AND WHAT THEY ENCODE INSTEAD

Language is the corpus's one factor with free labels, which is why it is measured, not because a head was expected to encode it: SONAR maps translations to nearby points, so language is a weak, residual factor of the embedding. No head encodes it outright. The best reaches a partition NMI of 0.18 against an attainable 0.875 (0.165 on the second seed), and no head-health lever changes that (Table 21). Partitions do carry the factor well above chance, and the hard code carries more of it: four times the softmax reference on null-corrected best-head NMI, and 0.42 against 0.28 in a 32-head probe. Single cells and conjunctions of cells detect language neighborhoods rather than identities (Finnish with Estonian, Japanese with Korean), which matches an embedding in which languages are weakly separated regions rather than directions (Appendix E.10).

What the strongest head holds: topic and genre. A semantic, language-agnostic embedding should be organized by meaning, and its strongest head is. Ranking layer-0 heads by η^2 and matching partitions across runs (65,536 held-out rows):

run	top head	η^2	next head	match across runs
ste_h76_init01 (seed 42)	h54	0.117	0.062	—
ste_h76_i01_s43 (seed 43)	h14	0.122	0.068	NMI 0.413 (others 0.010)
ste_h76 (unfixed init, seed 42)	h54	0.111	0.080	NMI 0.355 (others 0.024)

The topic head is the one head found in every run, the strongest partition in each by nearly a factor of two. A different seed grows it at a different index, and the same seed keeps its index across the initialization fix. It is also the head that reproduces on contributions (Appendix E.9), the exception to Section 5.5’s finding that heads do not reproduce, though at NMI 0.41 against a 0.87 same-run control it is the same variable carved differently, not the same partition. It mixes languages freely (NMI 0.155). Decoded as deviations from the head’s mean atom, its entries read as topic and genre: crime news, geopolitics and parliament, encyclopedic place description, manufacturing, marketing, legislative text and entertainment. Contrastive autointerp agrees, with every category of description above its null (Appendix E.10). Testing whether heads are natural variables therefore needs topic or genre labels, which this corpus does not have.

6 HARD CODES ON GPT-2 SMALL: ATOMS IN ACTIVATION SPACE

A second line of straight-through runs, from the headline branch, takes the hard code to GPT-2 small and adds what an embedding substrate lacks: the fraction of GPT-2’s loss recovered when the reconstruction is spliced back into the model, $(CE_{\text{ablate}} - CE_{\text{recon}})/(CE_{\text{ablate}} - CE_{\text{clean}})$, with the residual mean-ablated as the floor. These runs model the residual stream entering block 8, one block earlier than the GPT-2 runs of Section 5.5, centered and scaled to unit mean square, on FineWeb, and their FVU_w is unwhitened; they compare only with each other and with SAEs trained on the same cache. Held-out evaluation uses documents disjoint from training (520k tokens for the 40M-token runs, 65k for one-epoch runs).

The design. Relative to Section 5, three things change. Atoms are direct: signed vectors in activation space, with no latent dictionary space and no decoder, so every atom is literally a steering direction. Every layer, layer o included, classifies the unit direction of what it corrects, and each layer’s contribution is scaled by the observed residual norm, clipped at $\|x\|$, which makes the model positively homogeneous ($f(\xi x) = \xi f(x)$ for $\xi > 0$). The straight-through temperature is annealed early in training (Table 12). With 5×32 heads of 32 entries the code is 800 bits per token, with no continuous coefficients.

A discrete 800-bit code recovers more of GPT-2’s loss than a $k_{\text{SAE}}=32$ SAE at nearly the same FVU_w : 97.6% against 96.2%, at 0.195 against 0.189. Across architectures, then, FVU_w does not order downstream loss, while within the Ontologizer family on SONAR, text fidelity was monotone in FVU_w (Appendix C.4). With 64 entries

model	atoms	code per token	FVU _w	loss recovered
direct + gain + STE	signed, in $\mathbb{R}^{768}$	800 bits	0.195	97.6%
private heads + gain + STE	48-dimensional blocks, shared decoder	800 bits	0.186	97.5%
private heads, $k=64$	48-dimensional blocks, shared decoder	960 bits	0.164	97.9%
TopK SAE, $k_{SAE}=32$, $m=5120$	—	32 coeffs + ~394 bits	0.189	96.2%
TopK SAE, $k_{SAE}=128$, $m=5120$	—	128 coeffs + ~1580 bits	0.116	98.6%
soft Ontologizer (1 epoch)	—	4960 mixture weights	0.016	99.85%

TABLE 6 · GPT-2 small, residual entering block 8. Ontologizers: one pass over 40M tokens, one seed. SAEs: 10 epochs over 10M tokens with `sae.py`’s unwhitened objective. Clean cross-entropy 3.806, mean-ablated 8.283; the direct code reaches 3.914.

per head (960 bits) the code passes the SAE on both measures. Every layer contributes, entry usage stays near uniform, and the observed gain is close to the best per-sample scale for each layer; one-epoch ablations show that the gain-shape split and direct atoms each help, direct atoms most, and that annealing matters (Appendix F.1).

Fibers. A rank- r local basis per entry, with coordinates read linearly from the unit residual, turns the code into a discrete base point plus r numbers per head. “Base only” is the same model with the fibers removed, which measures whether the discrete code still means what it did:

variant (1 epoch, direct + gain + STE)	extra numbers	FVU _w (base only)	loss recovered
base	0	0.217	96.95%
rank-4 fiber, fiber dropout 0.5	640	0.059 (0.274)	99.4%
rank-8 fiber, fiber dropout 0.5	1280	0.014 (0.304)	99.9%
rank-4 fiber, joint norm cap 0.2	640	0.180 (0.227)	97.6%
rank-1 radial fiber (gain per entry)	160	0.148 (0.669)	98.5%
frozen base, rank-4 fibers trained after	640	0.016 (0.215)	99.9%
unconstrained fibers	—	— (~10)	—

Fibers recover far more than the SAE’s advantage in FVU_w, but they carry hundreds of continuous numbers per token, more than the SAE’s 32 coefficients, so this is not a matched-capacity comparison. The question they raise is what keeps

the discrete base honest. Unconstrained, the fibers take over (base-only FVU_w about 10); fiber dropout holds the base at 0.23–0.30 whatever the rank; and a joint norm cap guarantees it at the cost of most of the fibers’ value. Training the fibers after freezing the base is best on both counts: the base is untouched by construction (0.215), and the fibers reach the reconstruction of unconstrained joint fibers. The mechanisms used in joint training suppress the fibers rather than protect the base.

Seed stability, replicated. Matching heads across two seeds by partition NMI, the partitions agree only in layer 0 (NMI 0.09 against 0.05 for unrelated heads within one run) and not at all in layers 1–4. With direct signed atoms, the gain-shape split from layer 0, a different site and a different objective, this repeats Section 5.5: which entries share a head depends on the seed past the first layer.

7 DICTIONARY GEOMETRY AND THE NON-NEGATIVITY CONSTRAINT

Two geometric facts frame this section (Appendix G). Row collinearity measured in the dictionary space is mostly gauge: 74–78% of each atom’s energy lies in the decoder’s null space, so dictionary statistics should be computed after decoding (Appendix G.1). And the selection rule shapes the dictionary less than training duration does, while the last layer becomes inert once the error before it is small (Appendix G.4).

7.1 THE ORTHANT FORCES A COLLINEAR DICTIONARY

Non-negativity binds through the Gram, not the null space. Two non-negative vectors have a non-negative inner product exactly, so for independent $|N(0, 1)|$ coordinates the chance pairwise cosine is $2/\pi = 0.637$ (minimum +0.58 over 4000 draws), where signed draws give 0.000. Dropping the absolute value (–signed) on GPT-2 layer 8, with the straight-through arm’s shape, crossed with dictionary width and trained to convergence (371,900 steps):

arm	$\bar{\kappa}$ (row cosine)	eff. rank	FVU_w	at step 90k
signed, $e=1536$	+0.0009	27.37	0.1540	0.2910
signed, $e=768$	+0.0019	26.95	0.1543	0.2942
abs, $e=1536$	+0.3038	9.15	0.1934	0.3962
abs, $e=768$	+0.4543	5.41	0.2927	0.4631
concat, $d_{\text{head}}=32$	+0.5353	3.30	0.4155	0.8055

TABLE 7 · Non-negativity $\times$ width on GPT-2 layer-8 residuals, held-out tail: geometry from the weights, FVU_w scored from each checkpoint. No collinearity penalty is active in any arm.

Forced collinearity costs reconstruction. At $e=1536$ the non-negative dictionary runs at a third of the signed effective rank and 26% more held-out error; at




FIGURE 13 · The orthant arms over training: held-out FVU_w scored from each checkpoint (a) and mean within-head row cosine (b). The dotted line marks step 90k. The signed arms coincide at both widths; abs is behind at both, and further behind at the square decoder.

$e=768$, 90% more. Across the five arms $\bar{\kappa}$ orders the error exactly at convergence (Spearman +1.00; +0.90 at step 90k, where concat is the single inversion). The cost is an interaction: narrowing e costs 0.2% with signed atoms and 51% with non-negative ones, so the wide dictionary buys room for the constraint rather than capacity (Appendix G.1). That undercuts Section 3.1.4’s claim that head-disjointness is what $e > d$ buys, and no model attains head-disjointness in any case (Appendix G.2). The non-negative arm’s gap at $e=1536$ grew from 8% at step 20k to 50% at 200k and has since narrowed to 26% (Figure 13), so it may be partly slower, but it has not caught up by the end of a full run.

Removing the absolute value buys a decorrelated, near-full-rank dictionary and lower error, roughly what the collinearity setpoint exists to buy. Signed atoms give up two things: a subtractive feature is harder to read as “this concept is present”, and the $\|F\|_1 \geq h$ floor under `norm_rows`, which rests on non-cancellation, disappears. Concatenated heads, which face a collapsed tradeoff rather than a raised floor, are the worst arm (Appendix G.3).

8 FORMAL READING: HEADS AS CAUSAL VARIABLES

The steering frame has a natural formalization in the causal-abstraction vocabulary [24, 25, 26, 27], and stating it makes several instruments in the battery estimators of named quantities rather than ad hoc scores.

Definitions. Read the Ontologizer as a structural causal model: each head (ℓ, i) is a candidate high-level *variable* with value space $\{1, \dots, k\}$ (soft state: its simplex),

the residual stack is the causal graph (with residual forwarding, layer $\ell+1$'s classifier reads the residual written by layers $\leq \ell$, so every earlier head is a parent of every later one), and withArgs set-interventions are the do-operator: $\text{do}[(\ell, i)=j']$ forces the head's assignment while leaving all non-descendants untouched by construction. The abstraction map assign sends an embedding x to its per-head assignments. Under this reading: `steerembed.py`'s realization rate is an estimator of whether an *input-space* intervention implements $\text{do}[(\ell, i)=j']$, and its collateral of how far it also intervenes on other variables; `steerfid.py`'s cycle-consistency gain estimates *interchange-intervention accuracy* [28] in the closed loop $\text{assign} \circ \text{encode} \circ \text{decode}$, bounded above by the $\sim 12\%$ round-trip survival of Section 4.4; and seed-to-seed matching is *empirical identifiability*. What the softmax leaves unidentified in principle is the per-head entry order (permutation) and any logit gauge the temperature absorbs; what must therefore be reproducible for the variable reading to hold is the partition itself and the decode directions up to relabeling. Section 5.5 measured exactly that on the hard code, on SONAR and GPT-2. Past layer 0 the variables are identified at the level of the layer, not the head — and on GPT-2 not even that beyond layer 1. Layer 0 is partially identified, with SONAR's identification concentrated in a single large-atom head.

Why the cascade costs identifiability. The stack's reproducibility deficit at matched head size (Section E.8) and its intervention collateral (Section 5.4) are one fact under this reading. A head's value is a function of its own input, and every head past layer 0 takes as input the residual that upstream heads leave; a different upstream solution changes what any downstream variable is *about*. Even layer 0, whose input is the data, is pinned only partly: its heads are trained jointly with everything downstream of them, so which partition of the raw input is best depends on what the later layers will do with the error. A flat layer has neither dependency: its heads are more reproducible at equal size and realize interventions far more reliably, with half the collateral. They are still not independently addressable — the hard code's heads sit almost on their decision boundaries in either shape — so the cascade compounds an entanglement that the code already has.

Metric hygiene. Claims about decode-direction geometry should not depend on an arbitrary choice of inner product. The contrast-set result does not: under the training objective's inverse-variance metric the Ontologizer's within-head repulsion strengthens (-24.4 vs. -17.6 raw) and both rivals' signatures are unchanged. Two further rules emerged. Measure in output space, not in the dictionary space, whose null-space component is unconstrained (Section G.1). And separate a layer's gain from a head's atoms before attributing anything to head size: across layers, size is the gain, and the gain is the layer (Section E.8).

Composition. A variable reading earns little if interventions on two heads interact uncontrollably. Measured additivity of decode deltas on `resid_nc` (`compose.py`, whitened metric, medians over pairs):

stratum	pairs	cos	rel. err	later layers
same-layer L0	28	0.959	0.350	4
same-layer L1	28	0.947	0.354	3
same-layer L2	28	0.967	0.259	2
same-layer L3	28	0.996	0.092	1
same-layer L4	28	1.000	0.019	0 (numerical floor)
cross-layer (earlier = L0/L1)	42	0.97–0.98	0.20–0.23	
cross-layer (earlier = L2/L3)	18	1.00	0.02–0.04	

Three observations. (i) Joint interventions point where the summed singles point everywhere (median cosine ≥ 0.95 in every stratum): `withArgs` interventions compose approximately additively. (ii) The non-additive component is entirely attributable to downstream reclassification: it decays monotonically with the number of layers below the intervention and hits the numerical floor when there are none (the measured 0.019 is float32 rounding amplified by the decoder cancellation, $\sim 3 \times 10^{-6}$ in float64). (iii) Cross-layer pairs compose *better* than same-layer pairs at matched depth: a set-intervention is immune to upstream perturbation, so cross-layer interaction is one-way. Note the contrast with Section 5.4: `withArgs` sets a head directly and composes well, while an *input-space* push reaches every head and disturbs most of them. The downstream reclassification that is a 5–35% correction to a set-intervention is the dominant effect of an input-space one.

9 PRELIMINARY CONCLUSIONS

1. **Contrast-set structure is real, and it is what the architecture buys.** Trained softmax heads’ decode directions repel (96% of heads below null), and hard-code heads’ decoded atoms are orthogonal to slightly repelling: each head holds well-separated alternatives. This is invisible to reconstruction metrics and to per-latent probes.
2. **The other designed orthogonality did not emerge: heads share their coordinates.** No trained model partitions its dictionary space between heads; the support overlap sits at or above a coordinate-shuffle null in every model, each head’s usage covers 52–98% of the coordinates, and this holds for the part of each atom the decoder sees as well as for its null-space part. Only very long training (8.27M steps) moved it, by sparsifying each head’s profile rather than separating heads. A static reading of the dictionary cannot assume disjoint heads; concatenated heads are the only route to them, by construction.

3. **A discrete ontology is possible, but not with the softmax architecture's readout.** The softmax model's argmax is not its code (at least $77\times$ from the best 800-bit code for `resid_nc`, nine orders of magnitude on `softmax_shm`). A straight-through arm with $2.4\times$ the heads makes the discrete code the model, and at 1900 index bits and no coefficients beats converged SAEs at matched index bits that also send 160 coefficients (0.154 vs. 0.258–0.274), and an L1 SAE spending 519 coefficients and $3.4\times$ the bits (0.220). It sits $3.1\times$ above the Gaussian reference for its rate, so at least that far from the best achievable. It needed a temperature matched to the logit scale, relative logit noise, a rescaled entropy bonus, and an initialization fix worth a third of its error.
4. **The discrete code holds up downstream.** On GPT-2 small, an 800-bit hard code with signed atoms in activation space and gain-shape on every layer recovers 97.6% of GPT-2's loss, more than a $k_{\text{SAE}}=32$ SAE at nearly the same FVU_w (96.2%), and 64 entries per head pass it on both measures. Across architectures FVU_w does not order downstream loss. Fibers on top of a frozen discrete base recover almost all the rest (99.9%) while leaving the base exactly as it was, so a discrete ontology and a near-lossless reconstruction need not trade against each other — at the price of hundreds of continuous numbers per token.
5. **Reconstruction is the price sheet, not the verdict, for soft codes.** Unstructured codes reconstruct $\sim 50\times$ better at matched coefficient budgets, and the per-head simplex constraint alone costs $\sim 30\times$; text fidelity is a monotone function of FVU_w and adds no independent axis.
6. **The residual cascade is measured on both sides of the ledger.** For: $1.79\times$ lower error than a flat layer at equal bits, parameters and near-uniform head usage on both sides, sourced in the pairing of each stage's input with the sample's own error (not utilization, conditioning, or the gain side-channel). Against: an input-space intervention disturbs 66% of the stack's other heads against 33% of a flat layer's, and realizes its target on 37% of rows against 99.6% — in both, collateral equals a random step's, because the hard code's heads sit almost on their decision boundaries; and at equal head size the stack's heads reproduce across seeds several times worse than flat heads at every layer, layer 0 included. The single-layer softmax variant `g160softmax` matches `resid_nc` on reconstruction, but a config-matched softmax Ontologizer reaches $6\times$ lower error (in-sample), so the reconstruction verdict against the stack was checkpoint-specific.
7. **Steering works in embedding space and the text cycle hides it.** The decode direction realizes a target entry on 82% of rows for the softmax model at strength 0.25, against a 3% chance rate, while the text round trip can observe at most $\sim 12\%$ of perfect interventions. The soft code is the more steerable at

low collateral; steerability does not track reconstruction, and the classifier’s own gradient is the right direction only where the input reaches the classifier directly.

8. **Individual heads past layer 0 are not reproducible, and layer 0 only partly.** Seeds agree on the error to four digits. Past layer 0 they agree on nothing that identifies a head — partitions, contributions, subspaces beyond the data’s, or Grams — each tested against a same-shape null, on SONAR and on GPT-2, and again under the direct-atom design of Section 6. Layer 0 reproduces at under a fifth of a same-run control. Its lead over the deeper layers is not atom magnitude: its atoms are no larger, and its larger contributions are its gain, which is a property of the layer. Within layer 0, heads with outsized atoms are the reproducible ones — on SONAR almost entirely one head. This remains the most consequential result for the steering frame: a head that is not the same head on the next run cannot be named, and every per-head claim here inherits that.
9. **Entries contribute near-equally except where the input is raw.** In layers 1–3 every entry carries close to its share (Gini 0.06–0.15, no dead entries); layer 0’s entries are unequal (Gini 0.33–0.59) because its atom norms vary widely inside heads, while usage stays near-uniform. The flat arm shows the same pattern, so it belongs to classifying the raw input.
10. **Heads follow what the embedding encodes: meaning before language.** The strongest head encodes topic and genre, as a language-agnostic semantic embedding should be organized, and it is the one head that recurs across seeds. Language, the one factor with free labels, was a probe rather than an expected variable: it is carried as neighborhoods across heads, never by one head, consistent with languages being weakly separated regions of SONAR space. On that probe the hard code carries four times what the softmax reference does, so partitions do align with known factors; a test of whether heads are natural variables needs topic or genre labels.
11. **Per-latent lenses mis-measure classifier codes.** Sparse probing ranks the Ontologizer worst while partition-level analysis finds verified causal language heads; detection F1 is ~80% explained by code density, and the Ontologizer sits on the SAE density trend. Parameter-space decodes are instances, not descriptions, and must be taken as deviations from the head mean to be legible at all.
12. **Non-negativity is expensive in geometry and in error.** It forces a collinear, low-effective-rank dictionary (a third of the signed rank) and costs 26% of held-out FVU_w at the wide dictionary, 90% at a square decoder; row collinearity orders all five GPT-2 arms by error at convergence. Signed atoms

decorrelate the dictionary and reconstruct better at either width. Dictionary statistics computed in the dictionary space rather than after decoding are mostly gauge.

13. **Every measure in this architecture needs a null of the same shape.** Non-negativity and the layer-dependent gain put structure into every floor. Three of four seed-reproducibility measures, the HSIC head dependence, the null-space readings, and the apparent size effect on reproducibility each reversed or dissolved under the right control.

10 FOLLOW-UP INVESTIGATIONS

Ordered by relevance to the steering goal.

- **Head identifiability first.** Before naming heads, make them reproducible or establish which are. Candidates: partition-level seed stability on the softmax Ontologizer and `g160softmax`; whether large-atom heads' reproducibility can be extended by training pressure; whether the $k=2$ arm's one-bit heads reproduce better than 32-way ones; what lets the topic head recur when no other head does; and a GPT-2 flat arm, so the matched-size comparison of Section E.8 has a second dataset.
- **Bring the GPT-2 hard code onto this branch, then to SONAR.** Section 6's runs need `direct`, `resid_first` and `gain_clip`, which only the headline branch has; porting them makes those results reproducible here and lets the same design run on SONAR, where its 800-bit code can be compared with the straight-through arm's 1900 bits. Direct atoms also make steering literal — a set-intervention adds $g_\ell(\mathbf{a}_{ij'} - \mathbf{a}_{ij})$ at one layer — and that has not been measured.
- **Fibers on a frozen base.** The two-stage recipe kept the base exact and matched unconstrained fibers' reconstruction; it should be the default, scored at matched continuous capacity against SAEs (rank-1 radial fibers already sit at 160 numbers per token).
- **Remeasure on the fixed initialization.** Entanglement and the gain and blend ablations ran on the shipped initialization, which was 1.84×10^8 off at step 0 and left 10% of the code unused.
- **The steering overlay at scale.** `steerembed.py` with SAE decoder directions, bilinear eigenfeatures and supervised diff-of-means vectors alongside Ontologizer directions; per-entry effect consistency across contexts; and the flat arm's near-perfect realization, halved collateral and better reproducibility weighed against its $1.79\times$ reconstruction cost. Collateral at the boundary

itself needs a lever: the hard code's winner–runner-up margins are $\sim 10^{-4}$, and binary heads, with margins a hundred times larger, flip the fewest other heads of any Ontologizer.

- **Task reconstruction.** Recover what each head classifies: partition-vs-label-library MI, led by topic and genre (from a labeled corpus or unsupervised topic clusters) since those are what the embedding is organized by, with length, formatting/noise and script as secondary factors; *sibling-contrastive* descriptions (positives where entry j wins within head i , background where its siblings win), validated by blinded assignment accuracy vs. $1/k$; and FLORES parallel-text invariance to split content heads from form heads. Run identically on Ontologizer heads and g160softmax's heads, and only on heads that pass the identifiability bar.
- **Hierarchy as MoE routing.** Conditional MI between successive layers' partitions tests whether depth builds compositional tasks; `forward="labels"` at the fixed initialization tests the pairing account from the other side.
- **The collinearity setpoint's mechanism.** The one lever that buys steerability acts on a statistic that is mostly gauge; repeat it on the decoded statistic, and on signed dictionaries.
- **Train with the support overlap.** No model attains between-head disjoint support (Appendix G.2), and none has trained with ω as a loss term rather than its predecessor; whether it can buy the partition, and at what reconstruction price, is open. Concatenated heads are the comparison that already has it.
- **Logit-scale control for straight-through training.** Normalize logits per head, or control t against a spread/ t setpoint, in place of the hand-matched temperature.
- **Fresh-encoded eval set.** New runs hold the tail out; the softmax Ontologizers do not, and every comparison against them inherits that asymmetry.
- **A TopK sweep at matched bits.** The converged frontier has TopK points only at $k_{\text{SAE}}=32$ and 160; a k_{SAE} sweep through the 1,900-bit region would pin the discrete-vs-linear comparison at exactly matched rate.
- **Winner dropout for the one-layer variants.** g160top1's 28% dead latents are within-head winner-take-all starvation.
- **Trained truncatability.** A deviation-sparsity penalty at training time would show whether the soft Pareto curve can shift left.

- **Rung composition and frontier fill.** groups+prefixes (2b+3), the one-layer variants at $m=11264$, the mid-range TopK sweep ($k_{\text{SAE}} \in \{512, 2048\}$); random-support and nonnegative-LS refit controls.
- **Closed.** HSIC head independence (heads are already independent), decoder sparsification and bias (no structure; $< 1\%$), code-forwarding alongside the residual, input whitening (subsumed by initialization), and entropy-pressure rescaling across shapes (likewise). The resid_nc layer-4 freeze is explained by saturation; the two anomalously coherent heads (L2h25, L3h28) remain to inspect.

REPRODUCIBILITY

All code is available at <https://github.com/WakaruLabs/ontologize>. The evaluation protocols are root-level scripts (sae.py, pareto.py, headstruct.py, headcoh.py, splitting.py, textfid.py, langprobe.py, refit.py, steerfid.py, compose.py, autointerp.py), driven by the sae_ladder.sh, sae_evals.sh, and autointerp_run.sh runners (idempotent, crash-resumable systemd user units; judge calls rate-metered). The straight-through arm and its diagnostics live in experiments/ste-arm/ (train_ste.py, steerembed.py, blendablate.py, gainablate.py, inputgeom.py, dictgeom.py, lastlayer.py, partition.py, headcontrib.py, entryshare.py, bigatoms.py, headlang.py, headscript.py, headeta.py, embedgeom.py, decodehead.py, make_zca.py); its notebook experiments/ste-arm/notes.md records every run’s command and the superseded readings. The GPT-2 hard-code runs of Section 6 live in experiments/gpt2/ (train_onto.py, diagnose.py, seedstab.py; notes in DIRECT_GAIN_STE.md and results/README.md, with the two headline runs’ configs, diagnostics and loss curves) and need the headline branch’s code. The curated test suite pins each protocol’s statistics against planted-structure synthetics. Artifacts live under data/out/sonar/ and data/out/gpt2_18/. At most three straight-through runs fit concurrently on one 24 GB device.

ACKNOWLEDGEMENTS

Design, typesetting, and editorial production of this preprint are by The Flock Capital Research Department, whose error-diffusion pipeline also printed the plates.


REFERENCES

- [1] Hervé Jégou, Matthijs Douze, and Cordelia Schmid. “Product Quantization for Nearest Neighbor Search”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 33.1 (2011), pp. 117–128.

- [2] Neil Zeghidour et al. *SoundStream: An End-to-End Neural Audio Codec*. 2021. arXiv: 2107.03312 [cs.SD]. URL: <https://arxiv.org/abs/2107.03312>.
- [3] Yoshua Bengio, Nicholas Léonard, and Aaron Courville. *Estimating or Propagating Gradients Through Stochastic Neurons for Conditional Computation*. 2013. arXiv: 1308.3432 [cs.LG]. URL: <https://arxiv.org/abs/1308.3432>.
- [4] Alex Tamkin, Mohammad Taufeeque, and Noah D. Goodman. *Codebook Features: Sparse and Discrete Interpretability for Neural Networks*. 2023. arXiv: 2310.17230 [cs.LG]. URL: <https://arxiv.org/abs/2310.17230>.
- [5] Anish Mudide et al. *Efficient Dictionary Learning with Switch Sparse Autoencoders*. 2025. arXiv: 2410.08201 [cs.LG]. URL: <https://arxiv.org/abs/2410.08201>.
- [6] Zhengyan Zhang et al. *MoEfication: Transformer Feed-forward Layers are Mixtures of Experts*. 2022. arXiv: 2110.01786 [cs.CL]. URL: <https://arxiv.org/abs/2110.01786>.
- [7] Bart Bussmann et al. *Learning Multi-Level Features with Matryoshka Sparse Autoencoders*. 2025. arXiv: 2503.17547 [cs.LG]. URL: <https://arxiv.org/abs/2503.17547>.
- [8] David Chanin et al. *A is for Absorption: Studying Feature Splitting and Absorption in Sparse Autoencoders*. 2025. arXiv: 2409.14507 [cs.CL]. URL: <https://arxiv.org/abs/2409.14507>.
- [9] Tue M. Cao et al. *Tree SAE: Learning Hierarchical Feature Structures in Sparse Autoencoders*. 2026. arXiv: 2605.07922 [cs.LG]. URL: <https://arxiv.org/abs/2605.07922>.
- [10] Seongwan Park, Taeklim Kim, and Youngjoong Ko. *Decoding Dense Embeddings: Sparse Autoencoders for Interpreting and Discretizing Dense Retrieval*. 2025. arXiv: 2506.00041 [cs.IR]. URL: <https://arxiv.org/abs/2506.00041>.
- [11] Nick Jiang et al. *Interpretable Embeddings with Sparse Autoencoders: A Data Analysis Toolkit*. 2026. arXiv: 2512.10092 [cs.AI]. URL: <https://arxiv.org/abs/2512.10092>.
- [12] Wonseok Shin and Songkuk Kim. *Aligning Sentence Embeddings to Human Concepts via Sparse Autoencoders*. 2026. arXiv: 2607.00023 [cs.IR]. URL: <https://arxiv.org/abs/2607.00023>.
- [13] LCM team et al. *Large Concept Models: Language Modeling in a Sentence Representation Space*. 2024. arXiv: 2412.08821 [cs.CL]. URL: <https://arxiv.org/abs/2412.08821>.
- [14] John X. Morris et al. *Text Embeddings Reveal (Almost) As Much As Text*. 2023. arXiv: 2310.06816 [cs.CL]. URL: <https://arxiv.org/abs/2310.06816>.

- [15] Adam Karvonen et al. *SAEBench: A Comprehensive Benchmark for Sparse Autoencoders in Language Model Interpretability*. 2025. arXiv: 2503.09532 [cs.LG]. URL: <https://arxiv.org/abs/2503.09532>.
- [16] Gonalo Paulo et al. *Automatically Interpreting Millions of Features in Large Language Models*. 2025. arXiv: 2410.13928 [cs.LG]. URL: <https://arxiv.org/abs/2410.13928>.
- [17] Zhengxuan Wu et al. *AxBench: Steering LLMs? Even Simple Baselines Outperform Sparse Autoencoders*. 2025. arXiv: 2501.17148 [cs.CL]. URL: <https://arxiv.org/abs/2501.17148>.
- [18] Xiangchen Song et al. *Position: Mechanistic Interpretability Should Prioritize Feature Consistency in SAEs*. 2025. arXiv: 2505.20254 [cs.LG]. URL: <https://arxiv.org/abs/2505.20254>.
- [19] Paul-Ambroise Duquenne, Holger Schwenk, and Benoît Sagot. *SONAR: Sentence-Level Multimodal and Language-Agnostic Representations*. 2023. arXiv: 2308.11466 [cs.CL]. URL: <https://arxiv.org/abs/2308.11466>.
- [20] Alec Radford et al. *Language Models are Unsupervised Multitask Learners*. Tech. rep. OpenAI, 2019.
- [21] Leo Gao et al. *Scaling and evaluating sparse autoencoders*. 2024. arXiv: 2406.04093 [cs.LG]. URL: <https://arxiv.org/abs/2406.04093>.
- [22] Alexander Matt Turner et al. *Activation Addition: Steering Language Models Without Optimization*. 2023. arXiv: 2308.10248 [cs.CL]. URL: <https://arxiv.org/abs/2308.10248>.
- [23] Nina Rimskey et al. *Steering Llama 2 via Contrastive Activation Addition*. 2024. arXiv: 2312.06681 [cs.CL]. URL: <https://arxiv.org/abs/2312.06681>.
- [24] Paul K. Rubenstein et al. “Causal Consistency of Structural Equation Models”. In: *Proceedings of the Thirty-Third Conference on Uncertainty in Artificial Intelligence (UAI)*. 2017. URL: <https://www.auai.org/uai2017/proceedings/papers/11.pdf>.
- [25] Sander Beckers and Joseph Y. Halpern. “Abstracting Causal Models”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 33. 2019, pp. 2678–2685. DOI: 10.1609/aaai.v33i01.33012678. URL: <https://ojs.aaai.org/index.php/AAAI/article/view/4117>.
- [26] Atticus Geiger et al. “Causal Abstractions of Neural Networks”. In: *Advances in Neural Information Processing Systems*. Ed. by M. Ranzato et al. Vol. 34. Curran Associates, Inc., 2021, pp. 9574–9586. URL: <https://proceedings.neurips.cc/paper/2021/hash/4f5c422f4d49a5a807eda27434231040-Abstract.html>.
- [27] Atticus Geiger et al. “Causal Abstraction: A Theoretical Foundation for Mechanistic Interpretability”. In: *Journal of Machine Learning Research* 26.83 (2025), pp. 1–64. URL: <https://jmlr.org/papers/v26/23-0058.html>.

- [28] Atticus Geiger et al. “Inducing Causal Structure for Interpretable Neural Networks”. In: *Proceedings of the 39th International Conference on Machine Learning*. Ed. by Kamalika Chaudhuri et al. Vol. 162. Proceedings of Machine Learning Research. PMLR, 2022, pp. 7324–7338. URL: <https://proceedings.mlr.press/v162/geiger22a.html>.
- [29] Michael T. Pearce et al. *Bilinear MLPs enable weight-based mechanistic interpretability*. 2025. arXiv: 2410.08417 [cs.LG]. URL: <https://arxiv.org/abs/2410.08417>.
- [30] Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. 2nd ed. Wiley, 2006.
- [31] Maja Popović. “chrF: character n-gram F-score for automatic MT evaluation”. In: *Proceedings of the Tenth Workshop on Statistical Machine Translation*. Ed. by Ondřej Bojar et al. Lisbon, Portugal: Association for Computational Linguistics, Sept. 2015, pp. 392–395. DOI: 10.18653/v1/W15-3049. URL: <https://aclanthology.org/W15-3049/>.
- [32] Simon Kornblith et al. “Similarity of Neural Network Representations Revisited”. In: *Proceedings of the 36th International Conference on Machine Learning*. 2019.
- [33] Le Song et al. “Feature Selection via Dependence Maximization”. In: *Journal of Machine Learning Research* 13 (2012), pp. 1393–1434.
- [34] J. G. Wendel. “A Problem in Geometric Probability”. In: *Mathematica Scandinavica* 11 (1962), pp. 109–111.



RIDGE · SEED 32 · THE SAME PIPELINE, CLOSING THE FILE

A NOTATION

This appendix collects the symbols used throughout, with where each is defined.

Conventions. Lowercase italic letters are scalars (t, g_ℓ), bold lowercase letters vectors ($\mathbf{x}, \mathbf{a}_{ij}$), uppercase italic letters matrices (W_{dec}, D_i), and sans-serif letters tensors of rank three or more ($\mathbf{D}, \mathbf{P}$). A subscripted object takes the font of what it is: p_{nij} is an entry of $\mathbf{P}$, $\mathbf{p}_{ni}$ a vector slice of it, and $D_i := D_{i\cdot}$ a matrix slice; where a lowercase entry would read as a count, it is bracketed, as in $[K]_{nij}$. Batching

adds a rank: one sample’s assignments are $P \in \mathbb{R}^{h \times k}$, a batch’s $\mathbf{P} \in \mathbb{R}^{b \times h \times k}$. Vectors are columns, so a matrix acts on the left ($W_{\text{dec}} \mathbf{r}_\ell$, with $W_{\text{dec}} \in \mathbb{R}^{d \times e}$ mapping $\mathbb{R}^e$ to $\mathbb{R}^d$); a batch stacks its samples as rows ($X \in \mathbb{R}^{b \times d}$), so a batched product carries a transpose, $\hat{X} = R W_{\text{dec}}^\top$, or is written as function application, $\text{dec}(R)$. Sets are calligraphic ($\mathcal{S}_n, \mathcal{G}$). Named statistics, losses and operators are roman ($\text{FVU}_w, \text{MSE}, \text{KL}_m, \text{dec}, \text{sg}$), as are H , entropy, and L , the loss; the paragraph around Eq. (34) follows information theory in writing random variables in uppercase. A hat marks a reconstruction ($\hat{x}_\ell$), a tilde a noised quantity ($\tilde{F}, \tilde{f}_i$), a bar a mean ($\bar{\mathbf{p}}_i, \bar{\kappa}$), and a prime on an index a second index of the same kind (i', j'). $\odot$ is the elementwise product and $\oplus$ concatenation.

Layers and states. A layer is indexed by its position, $\ell = 0, \dots, l - 1$, and a state of the stack by the number of layers applied, $0, \dots, l$: layer ℓ reads state ℓ and writes state $\ell + 1$. So $\mathbf{r}_0 = \mathbf{0}$ and $\mathbf{y}_0 = \mathbf{x}$, the final reconstruction is $\hat{\mathbf{x}}_l$, and the deep-supervision prefixes are $\hat{\mathbf{x}}_1, \dots, \hat{\mathbf{x}}_l$. Tables and prose name layers (“layer o ”), so a table’s row for layer ℓ reports the prefix through it, $\hat{\mathbf{x}}_{\ell+1}$.

Indices and sizes.

index	ranges over	size	
n, n'	samples	b	batch size
ℓ, ℓ'	layers, $0, \dots, l - 1$	l	layers
i, i'	heads of a layer	h	heads per layer
j, j'	entries of a head	k	entries per head
q	coordinates	d, e	input dimension, dictionary width
		r	fiber rank
		d_{head}	a head’s slice of e under ConcatDictBlock
		k_{sel}	entries kept by the $\text{top}k_{\text{sel}}$ selection rule
		h'	heads kept, probed or conjoined, out of h
		n_{const}	constant coordinates appended to a layer’s input
		m, k_{SAE}	SAE width, and active latents under TopK

k alone always counts Ontologizer entries.

Symbols. Grouped by the part of the paper that introduces them.

symbol	meaning	defined
<i>The model</i>		
$\mathbf{x}, X; \bar{\mathbf{x}}$	an input (d), a batch ($b \times d$); the data mean	Eqs. (20), (23)
enc	the encoder, the identity unless stated	Alg. 3

symbol	meaning	defined
$\mathbf{u}_\ell, \mathbf{u}$	layer ℓ 's shape: its input's unit direction, with the constant coordinate appended; what the classifier, router and fibers read ($\mathbf{u}$ within a layer)	Eq. (14)
$\mathbf{W}, \mathbf{V}$	the bilinear classifier's weight tensors	Eq. (1)
$K, \mathbf{K}$	logits, one sample's ($h \times k$) and a batch's	Eq. (1)
t	classifier temperature, $P = \text{softmax}_k(K/t)$	§3.1.3
$P, \mathbf{P}; \mathbf{p}_{ni}, p_{nij}$	assignments, one sample's and a batch's; head i 's distribution, one entry's probability	§3.1.3
$\bar{\mathbf{p}}_i$	head i 's mean assignment over the data, its usage	Eq. (4)
p_{drop}	winner-dropout probability	§3.1.2
$\mathbf{D}, D_i$	the dictionary ($h \times k \times e$), and head i 's slice ($k \times e$)	Eq. (5)
$\mathbf{a}_{ij}$	an atom: $ \mathbf{D}_{ij} $, or D_{ij} . when signed	Eq. (5)
G_i	head i 's dictionary Gram	Eq. (3)
$\mathbf{f}_i; F, \mathbf{F}$	head i 's output (e); stacked over heads, and over heads and samples	Eq. (5)
$\mathbf{f}'_\ell, F'$	layer ℓ 's pooled output, and a batch's	Eq. (10)
γ, γ_i, Γ	router scales: per head, one head's, and a batch's	Eq. (19)
$U_{ij}, \mathbf{c}_i; C, \mathbf{C}$	fiber basis ($e \times r$) and coordinates; one sample's coordinates, and a batch's	Eq. (19)
$\mathbf{A}, A_i; \beta, \rho$	fiber readout, and head i 's; coordinate bound, norm cap	§3.1.7
g_ℓ, g_n	gain: the norm of layer ℓ 's input, or of one sample's	Eq. (11)
$\mathbf{r}_\ell, R$	accumulator after ℓ layers, and batched	Eq. (11)
$W_{\text{dec}}, \text{dec}$	the decoder shared by all layers	Eq. (12)
$\hat{\mathbf{x}}_\ell, \hat{\mathbf{X}}$	reconstruction after ℓ layers; the final one, batched	Eq. (12)
$\mathbf{y}_\ell$	residual after ℓ layers, which is layer ℓ 's input	Eq. (13)
$\text{sg}[\cdot]$	stop-gradient	Eq. (13)
W_{flat}	the code-to-output map of a model's flattened code, $\hat{\mathbf{x}} = W_{\text{flat}}\mathbf{p}_{\text{flat}}$; its columns are the entries' decode directions	§B.5
<i>Noise, statistics and loss</i>		
$\varepsilon; \varepsilon_K, \varepsilon_F$	a standard normal draw; the logit and feature noise terms	§3.1.6
$\sigma_X, \sigma_K, \sigma_F$	noise scales for inputs, logits and features	§3.1.6
$\tilde{\mathbf{F}}, \tilde{\mathbf{f}}_i$	noised per-head outputs	Eq. (7)
ϵ	floating-point epsilon of the working dtype	Eq. (14)
τ_{rev}	dead-entry revival threshold	Alg. 2
$\text{entropy}_k, H$	entropy of a head's assignment in bits, and its mean over samples and heads	Eq. (2)

symbol	meaning	defined
cossim_b	between-sample similarity of per-head outputs	Eq. (3)
KL_m	divergence of each head's batch-mean	Eq. (4)
$\kappa_i, \bar{\kappa}$	assignment from uniform, averaged over heads mean cosine between head i 's atoms, and its mean over heads (row collinearity)	Eq. (22); §G.1
π_i, ω	head i 's usage-weighted coordinate profile; the support overlap between heads	Eq. (6)
$\text{L1}, \text{L1}_K, \text{L1}_S$	feature L1; the classifier gate's; the router's	Eqs. (7), (22)
$\text{KL}_{\text{pwak}}, \text{L2}_{\text{pwak}}$	partition-consensus statistics	Eq. (22)
ψ_ℓ	layer ℓ 's statistics row	Eq. (22)
$\mathbf{w}$	per-dimension MSE weights: inverse variance, normalized to mean one	Eq. (20)
$\mathbf{s}; s_g, s_\omega, s_{\text{L1}}, \dots$	loss weights: the vector paired with ψ_ℓ , and individual weights	Eq. (21)
MSE, L	whitened reconstruction error; the total loss	Eqs. (20), (21)
<i>Evaluation and baselines</i>		
FVU_w	whitened fraction of variance unexplained	Eq. (23)
$\nu, \mathcal{S}_\nu; \mathbf{o}$	deviations kept per head, and their entries; the origin, the mean assignment	Eq. (24)
$\text{FVU}_{w,G}, n_{\text{bits}}; \lambda_j, \theta$	the Gaussian reference at a bit budget; covariance eigenvalues, the water level	Eq. (25)
H_{real}	realized bits: the summed per-head entropy of usage	Eq. (34)
$\mathcal{G}; \text{exh}, \text{exc}$	a head or latent group; its exhaustiveness and exclusivity	Eq. (26)
$\kappa(\mathcal{G}), \zeta_{\text{cos}}; \mathbf{v}_j$	a group's mean decode-direction cosine, and its standardized score; a decode direction	Eq. (27)
$(\cdot)^{(1)}, (\cdot)^{(2)}$	the two models in a pairwise comparison	Eq. (28)
$\mathbf{q}_i^{(1)}; [\text{M}^{\text{part}}]_{ii'},$ $[\text{M}^{\text{contr}}]_{ii'}$	head i 's entry choices over the rows; the head-matching matrices	Eq. (28)
ϕ_{ni}, size_i	head i 's output contribution on sample n ; its share of the variance	Eqs. (28), (29)
$\mathbf{a}_{ij}^{\text{dec}}, \bar{\mathbf{a}}_i^{\text{dec}}$	a decoded atom; head i 's usage-weighted mean	Eq. (29)
recovered	fraction of the next-token loss recovered	Eq. (30)
$\eta^2; C, n_{\text{rows}}$	variance explained by a partition; its cells, the rows scored	Eq. (31)
k_{probe}	the sparse-probing budget	§B.5
$\mathbf{z}_n^*, \mathcal{S}_n; W, \mathbf{b}$	refitted coefficients on a frozen support; a linear code's decode directions (the columns of W) and offset	Eq. (32)
$\mu_{ni}; j_1, j_2; J_i; \alpha$	a head's logit margin between two entries; the Jacobian of head i 's logits; step strength	Eq. (33)
j^*, dir, ξ_n	a steering target, a direction, the random control	Alg. 4

symbol	meaning	defined
assign	the abstraction map from an input to its code	§8
$\mathbf{z}, z_{nj}$	SAE latent activations	Eq. (26)
$W_{\text{enc}}, \mathbf{b}_{\text{enc}}; W_{\text{dec}}, \mathbf{b}_{\text{dec}}$	SAE encoder weights and bias; its decoder, whose columns are the latents’ directions, and the pre-subtracted decoder bias	§B.2
$L0, s_{L1}, \text{aux_k}$	active SAE latents per sample; an L1 SAE’s weight; dead-latent revival	§B.2; Table 13

B SUPPLEMENTARY METHODS

B.1 DATA AND OBJECTIVE

Embedding space and data. All SONAR models are trained on a fixed cache of 4,000,000 SONAR sentence embeddings ($d=1024$, unit-norm) computed from an interleaved 86-language mC4 sample, with a language label sidecar. The stream order is deterministic, so raw text is recoverable by row index (verified against the language sidecar per row). The cloud is unit-normed but not centered: $\|\mathbb{E}[\mathbf{x}]\| = 0.564$, so it is a spherical cap and 0.318 of every squared norm is one shared direction. It is dense without being isotropic — effective dimension 108 of 1024 (participation ratio 268 on the whitened eval tail), yet 943 principal components for 99% of variance — and samples average a pairwise cosine of 0.31. Sections 7.1 and 5.5 additionally use GPT-2 [20] layer-8 residual-stream activations ($d=768$, harvested by `encode_acts.py`) as a substrate where sparse dictionary learning is known to work. Section 6 uses a separate GPT-2 setup from the headline branch — the residual stream entering block 8 on FineWeb, unwhitened, scored also by GPT-2 loss recovered — whose numbers compare only within that section.

Evaluation. All models minimize the whitened MSE of Eq. (20). We report FVU_w , whitened MSE divided by the whitened variance of a constant predictor,

$$\text{FVU}_w := \frac{\sum_n \sum_q w_q (\hat{x}_{n,q} - x_{n,q})^2}{\sum_n \sum_q w_q (x_{n,q} - \bar{x}_q)^2}, \quad \bar{\mathbf{x}} := \mathbb{E}_n[\mathbf{x}_n], \quad (23)$$

with $\mathbf{w}$ the per-dimension inverse variance normalized to mean one, on the final 32,768 cache rows. This tail is excluded from SAE training and from the training of every straight-through arm. *Caveat: the softmax Ontologizers (resid_nc and the selection sweep) predate the holdout and trained on the full cache including the tail, so their scores are in-sample; a fresh-encoded evaluation set is future work.* Where activations are needed, Ontologizer features are the flattened per-head classification probabilities $P \in \Delta^{k \cdot l \cdot h}$ at the run’s final temperature (entry j of head i in layer ℓ is feature $(\ell h + i)k + j$; fire threshold $2/k$), and SAE features are latent activations (fire threshold 0; runs with softmax heads use the above-uniform threshold $2/(m/h)$ for h heads,

matching the entry convention). Scoring a softmax run at the default temperature 1.0 rather than its annealed 0.03 puts a healthy checkpoint at $FVU_w \approx 0.91$, which is indistinguishable from a broken one; every number here uses the run’s own temperature.

B.2 MODELS AND BASELINES

Every model’s configuration is tabulated in Appendix B.3.

Softmax Ontologizers. The reference Ontologizer (`resid_nc`) is the architecture of Section 3.1 with softmax selection: $l=5$ residual layers of $h=32$ heads with $k=32$ entries each. Layers 0–3 trained healthily (prefix FVU_w $0.129 \rightarrow 0.0051$); layer 4 froze 22/32 heads late in temperature annealing and is mildly harmful (Section G.4 explains why). A later *selection sweep* holds that shape fixed and varies only the DictBlock selection rule — `top1`, `top2`, `top4`, `top8`, `top16`, `softmax` — each with and without a raised mean-entropy bonus (`_shm`).

The straight-through arm. `train_ste.py` imports `sonar.py`’s configuration and overrides `select="ste"`: the forward pass selects the argmax one-hot, and the backward pass uses $\text{softmax}(K/t)$ as its surrogate. Its shape, calibration and initialization are described with its results in Section 5.

SAEs. SAEs follow Gao et al. [21]: pre-activation $\text{relu}(W_{\text{enc}}(\mathbf{x} - \mathbf{b}_{\text{dec}}) + \mathbf{b}_{\text{enc}})$, reconstruction $\hat{\mathbf{x}} = W_{\text{dec}}\mathbf{z} + \mathbf{b}_{\text{dec}}$, with TopK selection of k_{SAE} latents, unit-norm decoder directions (the columns of W_{dec}), tied initialization and `aux_k` dead-latent revival, trained to convergence; ReLU+L1 SAEs replace TopK with an L1 penalty and neuron resampling.

The structural-ablation ladder. To attribute differences to specific architectural commitments rather than to whole architectures, we train SAE variants that each add one Ontologizer-style constraint (all $m=5120$, same data and objective):

rung	run	added structure	analogue
1	k32	none (TopK, $k_{\text{SAE}}=32$)	baseline
2a	g160top1	160 heads, hard winner per head	heads (hard)
2b	g160softmax	160 heads, softmax per head	heads (soft)
3	k32_p5	5 nested prefix losses (matryoshka)	deepsup layers
—	k32_b1	bilinear encoder, $[\mathbf{x} - \mathbf{b}_{\text{dec}}; 1]$ factors	bilinear classifier

Rungs 2a and 2b go far enough that they are no longer SAEs in any established sense. Per-head competition over a shared decoder is a single layer of linear classifiers selecting dictionary entries: `g160top1` is a one-layer hard-code Ontologizer and

shared by every softmax run

input	SONAR mC4 embeddings, $d=1024$, unit norm
shape	$l=5$ layers $\times$ $h=32$ heads, $k=32$ entries per head
classifier	bilinear (order 2), no gate
forward	residual; resid_norm, resid_const, resid_gain; joint deep supervision
optimizer	Adam, lr 5×10^{-5} , gradient clip 1.0, batch 256, 24 epochs
temperature	$t=1 \rightarrow 0.03$, geometric over 50k steps, then held
noise	logits normal, $\sigma_K=0.02 \rightarrow 0.006$; features featvar, $\sigma_F=0.1$; input off
winner dropout	$p_{\text{drop}}=0 \rightarrow 0.1$ over the anneal
loss	whitened MSE; $s_{\text{cossim}_h}=10^{-5}$; ghost path off
seed	42

run	selection	e	s_{KL_m}	other	steps
resid_nc	softmax	4096	$10^{-6}^\dagger$	$s_{\text{L1}}=10^{-9}$, $s_{\text{cossim}_h}=0 \rightarrow 10^{-5}^\dagger$	372.6k
resid_nc_hm	softmax	2048	10^{-6}	$s_{\text{L1}}=10^{-9}$, $s_{\text{cossim}_h}=10^{-5} \rightarrow 10^{-6}^\dagger$	8.27M [‡]
sweep_softmax	softmax	2048	10^{-6}	$s_{\text{cossim}_h}=10^{-6}$	369.5k
sweep_topk_sel	top- k_{sel} , $k_{\text{sel}} \in \{1, 2, 8, 16\}$	2048	10^{-6}	$s_{\text{cossim}_h}=10^{-6}$	369.5k
sweep_top8_relu	top-8	2048	10^{-6}	ReLU gate on the classifier	369.5k
sweep*_shm	softmax, top- k_{sel} , $k_{\text{sel}} \in \{2, 4, 8, 16\}$	2048	3×10^{-5}	$s_{\text{cossim}_h}=10^{-6}$	369.5k

TABLE 9 · Softmax Ontologizer configurations. [†]Changed at a resume: resid_nc acquired s_{KL_m} and its s_{cossim_h} mid-run, and resid_nc_hm lowered s_{cossim_h} . [‡]Latest checkpoint; results attributed to resid_nc_hm in Appendix D use its 4.38M-step checkpoint. cossim_h is the retired between-head cosine (Appendix B.4).

shared by every SONAR straight-through arm

input	SONAR mC4 embeddings, $d=1024$; eval tail held out
selection	straight-through (ste): argmax forward, softmax backward
classifier	bilinear (order 2), no gate
dictionary	non-negative, $e=2048$, shared linear decoder
forward	residual; resid_norm, resid_const, resid_gain; joint deep supervision
optimizer	Adam, lr 10^{-5} , gradient clip 1.0, batch 256, 24 epochs
temperature	$t=1.5 \times 10^{-4}$, constant
noise	logits batchnorm, $\sigma_K=0.3$; features featvar, $\sigma_F=0.1$; input off
winner dropout	$p_{\text{drop}}=0.1$
loss	whitened MSE; $s_{L1}=10^{-9}$, $s_{\text{cossim}_b}=10^{-5}$, $s_{\text{cossim}_h}=10^{-6}$; ghost path off
seed	42 unless noted

arm	shape	k	bits	s_{KL_m}	init scale	other	steps
ste_h76	5×76	32	1900	10^{-4}	—		369.5k
ste_h76_init01	5×76	32	1900	10^{-4}	0.1		369.5k
ste_h76_i01_s43	5×76	32	1900	10^{-4}	0.1	seed 43	369.5k
ste_h76_hm1e6	5×76	32	1900	10^{-6}	—		369.5k
ste_h76_kcos	5×76	32	1900	10^{-4}	—	row-cosine setpoint 0.5	377.3k
ste_h76_hsic3	5×76	32	1900	10^{-4}	—	head CKA 10^{-3} (biased)	380.2k
ste_k2_h380	5×380	2	1900	5×10^{-4}	0.1	batch 256 $\rightarrow$ 128 [†]	739.0k
ste_l1_h380	1×380	32	1900	10^{-4}	—		369.5k
ste_l1_h380_hm5	1×380	32	1900	5×10^{-4}	—		369.5k
ste_l1_h380_zca	1×380	32	1900	5×10^{-4}	—	frozen ZCA encoder	353.3k
ste_l1_h380_enc	1×380	32	1900	5×10^{-4}	—	trained linear encoder	369.5k
ste_l1_h380_kcos	1×380	32	1900	10^{-4}	—	row-cosine setpoint 0.5	929.4k
ste_l1_h380_i01	1×380	32	1900	5×10^{-4}	0.1		369.5k
ste_l1_h380_i01_s43	1×380	32	1900	5×10^{-4}	0.1	seed 43	369.5k
ste_l1_h380_i01_hm1e4	1×380	32	1900	10^{-4}	0.1		369.5k
ste_k128_h54	5×54	128	1890	10^{-4}	0.1	$s_\omega=10^{-6}\ddagger$	369.5k
ste_k128_h54_direct	5×54	128	1890	10^{-4}	0.1	signed direct atoms, $s_\omega=10^{-6}\ddagger$	369.5k

TABLE 10 · SONAR straight-through configurations. Init scale is -dict-init-scale: each dictionary rescaled at initialization so a layer writes that fraction of its input (— leaves the default initialization). [†]Changed at a resume. [‡]Trained after cossim_h was retired, so s_ω takes its place.

shared by every GPT-2 straight-through arm on this branch

input	GPT-2 small, residual leaving block 8, $d=768$; document-level holdout
selection	straight-through (ste)
classifier	bilinear (order 2)
forward	residual; resid_norm, resid_const, resid_gain; joint deep supervision
optimizer	Adam, lr 5×10^{-5} , gradient clip 1.0, batch 256, 24 epochs
temperature	$t=1.0$, constant
noise	logits normal, $\sigma_K=0.02$; features featvar, $\sigma_F=0.1$; input off
winner dropout	$p_{\text{drop}}=0.1$
loss	MSE weighted by gpt2_l8.mse_weights_matched; $s_{L1}=10^{-9}$, $s_{\text{cosim}_b}=10^{-5}$, $s_{\text{cosim}_h}=10^{-6}$, $s_{KL_m}=10^{-6}$

arm	shape	k	dictionary	gate	init scale, seed	steps
ste_h76	5×76	32	non-negative, $e=1536$	none	1.0, 42	371.9k
ste_h76_e768	5×76	32	non-negative, $e=768$	none	1.0, 42	371.9k
ste_h76_sgn	5×76	32	signed, $e=1536$	none	1.0, 42	371.9k
ste_h76_sgn768	5×76	32	signed, $e=768$	none	1.0, 42	371.9k
ste_h76_cat32	5×76	32	non-negative, concat, $d_{\text{head}}=32$	none	1.0, 42	371.9k
ste_h20_cat128	5×20	256	signed, concat, $d_{\text{head}}=128$	sigmoid	—, 42	371.9k
ste_h20_cat128-43	5×20	256	signed, concat, $d_{\text{head}}=128$	sigmoid	—, 43	371.9k
ste_h20_cat128-sc	5×20	256	as above, plus router	sigmoid	—, 42	371.75k
ste_h20_cat128-sc43	5×20	256	as above, plus router	sigmoid	—, 43	371.75k

TABLE 11 · GPT-2 straight-through configurations on this branch: the orthant arms of Table 7 and the concatenated seed pairs of Section 5.5. The router (scaled) is sigmoid-gated and non-negative.

shared by the headline-branch runs of Section 6

input	GPT-2 small, residual entering block 8, $d=768$; mean-centered, scaled to unit mean square; FineWeb
shape	$l=5$ layers $\times$ $h=32$ heads, $k=32$ entries per head (800 bits) unless noted
selection	straight-through, temperature $t=1 \rightarrow 0.03$ over the first 60k steps
classifier	bilinear (order 2)
forward	residual; resid_first, resid_gain, gain_clip, resid_norm, resid_const; joint deep supervision
optimizer	Adam, lr 2×10^{-4} , batch 256
noise	logits normal, $\sigma_K=0.02 \rightarrow 0.006$; features featvar, $\sigma_F=0.1$
winner dropout	$p_{\text{drop}}=0 \rightarrow 0.1$
loss	unwhitened MSE; $s_{L1}=10^{-9}$, $s_{\text{cossim}_b}=10^{-5}$, $s_{\text{cossim}_h}=10^{-6}$, $s_{KL_m}=10^{-6}$
seed	42 unless noted

run	dictionary	data
direct + gain + STE	signed atoms in $\mathbb{R}^{768}$, no decoder	40M tokens, one pass (156k steps)
private heads + gain + STE	non-negative, $e=1536$ in 48-dimensional blocks, shared decoder	40M tokens, one pass
one-epoch ablations	as in Appendix F.1	10M tokens

TABLE 12 · Configurations of the headline-branch GPT-2 runs (experiments/gpt2/results/D40_*/config.json). Fiber variants add fiber_rank r with the dropout, cap or freezing given in Section 6’s table.

shared by every sae.py run

input	the SONAR cache and eval tail of the Ontologizers, $d=1024$
objective	whitened MSE, as the Ontologizers
architecture	pre-subtracted decoder bias, tied initialization
decoder	unit-norm directions, the columns of W_{dec}
dead latents	aux_k 512 (TopK runs) or neuron resampling every 5000 steps (L1 runs)
optimizer	Adam, batch 4096

run	m	sparsity	encoder	structure	lr	steps
m5120_k32	5120	top-32	linear	—	4×10^{-4}	144.3k
m5120_k32_s43	5120	top-32	linear	seed 43	4×10^{-4}	144.3k
m5120_k32_b1	5120	top-32	bilinear	—	4×10^{-4}	144.3k
m5120_k32_p5	5120	top-32	linear	5 nested prefixes (Matryoshka)	4×10^{-4}	144.3k
m5120_l1*	5120	L1, five weights s_{L1} (caption)	linear	—	4×10^{-4}	144.3k
m11264_k32	11264	top-32	linear	—	10^{-4}	336k
m11264_k160	11264	top-160	linear	—	$10^{-4\dagger}$	336k
m11264_k5120	11264	top-5120	linear	— (dense, $L0 \approx 1520$)	10^{-4}	22k
m5120_g160top1	5120	160 heads, top-1	linear	one-layer hard variant	4×10^{-4}	144.3k
m5120_g160softmax	5120	160 heads, softmax	linear	one-layer soft variant	4×10^{-4}	144.3k

TABLE 13 · SAE baselines (above the rule) and the single-layer variants (below it; Section B.2). The L1 sweep’s five weights s_{L1} are 3×10^{-6} , 10^{-5} , 3×10^{-5} , 10^{-4} and 3×10^{-4} ; * is the weight as the run name spells it (m5120_l13e-05). $\dagger$ Not recorded; inferred from the run sharing m11264_k32’s step count.

entries a sample selected, so a hard code’s unselected entries were unconstrained. The support overlap ω measures the property directly, for every used entry, under either sign, at $\mathcal{O}(hke + h^2e)$.

Flattened row cosine $\text{cossim}_{\text{flat}}$ was the mean off-diagonal cosine over every pair of atoms in the layer, between heads as well as within. It was added to catch heads whose atoms are individually diverse but parallel to another head’s, which the within-head κ_i cannot see. Its mean is dominated by between-head pairs ($k - 1$ of a row’s $hk - 1$ partners share its head), and on non-negative atoms a between-head cosine of zero is again disjoint support, so ω , weighted by usage and computed per head pair, replaced it too.

B.5 EVALUATION PROTOCOLS

Each protocol is a standalone script sharing the eval split and activation conventions; the root-level ones are unit-tested against planted-structure synthetic data. Figure 14 shows what each script reads and the question it answers.

Capacity Pareto (pareto.py). Capacity is continuous coefficients transmitted per sample, plus index bits. Ontologizer code points are *top- ν deviation truncations*: per head, keep the ν entries with largest $|p_j - o_j|$, pin the rest to the origin distribution $\mathbf{o}$, renormalize,

$$\mathcal{S}_\nu := \text{top-}\nu_j |p_j - o_j|, \quad p'_j := \frac{\mathbf{1}[j \in \mathcal{S}_\nu] p_j + \mathbf{1}[j \notin \mathcal{S}_\nu] o_j}{\sum_{j' \in \mathcal{S}_\nu} p_{j'} + \sum_{j' \notin \mathcal{S}_\nu} o_{j'}}, \quad \mathbf{o} := \mathbb{E}_n[\mathbf{p}_n], \quad (24)$$

applied to every head, so the code transmits ν coefficients per head. The origin is the corpus-mean classification $\mathbb{E}[\mathbf{p}]$ — measured, not uniform: the model output is jointly *linear* in the flattened classifications (zero residual seed, linear decoder), so $\text{decode}(\mathbb{E}[\mathbf{p}]) = \mathbb{E}[\hat{\mathbf{x}}]$, which decodes to the corpus-mean generic text, while the uniform code lands off-manifold (cosine 0.067 to $\mathbb{E}[\mathbf{x}]$) — despite KL_m being included in the loss precisely to make the uniform code the generic one (Section 3.1.3). Endpoints: $\nu=0$ (input-independent origin), $\nu=k$ (exact soft model), and end-to-end argmax (“hard”). An alternative head-level code mean-ablates all but the h' lowest-entropy heads per layer. SAE points use measured eval L0. A hard code’s index bits are $l \cdot h \log_2 k$; we compare each against a *Gaussian reference*: the least distortion a code of that rate can achieve on a Gaussian source with the eval tail’s whitened covariance, given by reverse water-filling [30]. With $\lambda_1, \dots, \lambda_d$ that covariance’s eigenvalues, a rate of n_{bits} bits per sample gives

$$\text{FVU}_{w,G}(n_{\text{bits}}) = \frac{\sum_j \min(\theta, \lambda_j)}{\sum_j \lambda_j}, \quad n_{\text{bits}} = \sum_j \max\left(0, \frac{1}{2} \log_2 \frac{\lambda_j}{\theta}\right), \quad (25)$$

with the water level θ solved from the rate. It is a reference, not a floor, for two reasons. Among sources of a given covariance the Gaussian has the largest




FIGURE 14 • What validates what. Columns are the artifacts each evaluation script reads: the embedding or activation cache, whose tail rows are the eval split sae.py excludes from training, with its sidecars; Ontologizer checkpoints; sae.py baselines; and the frozen SONAR encoder and M2M100 decoder. A filled mark is a required input, an open mark an optional input or one of several alternatives (most scripts take either model type; splitting.py takes two runs, such as seed replicas). Dashed arrows are one script's output read by another. selection.py, enrich.py, assignmap.py, entrydrift.py, bilinspec.py, headnmi.py and effinfo.py (in experiments/effective-information/) are diagnostics with no protocol below; no result in this document comes from them.

rate–distortion function under squared error, so for the real embeddings the least achievable distortion at rate n_{bits} is *at most* $\text{FVU}_{w,G}(n_{\text{bits}})$ and may be well below it [30]. And the rate–distortion function is approached by coding long blocks of samples jointly with variable-length codes, where our codes spend a fixed n_{bits} bits on each sample alone. The reference is therefore safe to use in one direction only: a code *above* it is at least that far from optimal, and nothing we report can establish that a given rate is too small.

Categorical form (headstruct.py). An Ontologizer head is exhaustive (always fires) and exclusive (one winner). We test whether SAE latents form such units: discovery clusters latents by *paradigmatic* affinity — similar co-fire profiles with the rest of the dictionary $\times$ anti-correlation with each other ($\text{PMI} < 1$) — under a group-size cap, on one half of the rows; scoring on the held-out half computes per-group exhaustiveness, exclusivity and the coefficient of variation of the group’s summed activation,

$$\begin{aligned} \text{exh}(\mathcal{G}) &:= \Pr\left(\sum_n \sum_{j \in \mathcal{G}} \mathbf{1}[z_{nj} > 0] \geq 1\right), \\ \text{exc}(\mathcal{G}) &:= \mathbb{E}_n\left[\sum_{j \in \mathcal{G}} \mathbf{1}[z_{nj} > 0] \mid \sum_{j \in \mathcal{G}} \mathbf{1}[z_{nj} > 0] \geq 1\right], \end{aligned} \quad (26)$$

against size-matched random partitions of the live-latent pool. A head scores $\text{exh} = \text{exc} = 1$. Head-like structure requires exhaustiveness *above* null with exclusivity *below* null; topical co-firing clusters show the opposite signature. -trained-groups scores a run’s own trained heads instead.

Semantic coherence vs. contrast (headcoh.py). The complement to categorical form: do a head’s entries *mean* similar things, or do they form a contrast set? Every entry has an exact decode direction in SONAR output space — for the Ontologizer, the columns of the linear map W_{flat} with $\hat{\mathbf{x}} = W_{\text{flat}} \mathbf{p}_{\text{flat}}$; for SAEs, decoder directions — and SONAR space is a semantic sentence space, so within-head direction similarity is semantic similarity of what entries decode to. Per head or group $\mathcal{G}$ we report mean pairwise cosine and top-singular-value energy fraction, z-scored as ζ against size-matched random subsets of the same pool (same layer for Ontologizer heads),

$$\kappa(\mathcal{G}) := \mathbb{E}_{j \neq j' \in \mathcal{G}} \cos(\mathbf{v}_j, \mathbf{v}_{j'}), \quad \zeta_{\cos}(\mathcal{G}) := \frac{\kappa(\mathcal{G}) - \mathbb{E}_{\mathcal{G}'}[\kappa(\mathcal{G}')] }{\text{sd}_{\mathcal{G}'}[\kappa(\mathcal{G}')]}, \quad (27)$$

with $\mathbf{v}_j$ entry j ’s decode direction and $\mathcal{G}'$ ranging over 500 random same-size subsets; as a metric-hygiene control the same statistics are recomputed under the training objective’s inverse-variance inner product (-metric whitened). $\zeta_{\cos} \gg 0$ means a head clusters semantically; $\zeta_{\cos} \ll 0$ means its atoms repel — the signature of a *contrastive classifier*, whose entries’ atoms are pushed apart. A

second view SONAR-encodes autointerp description texts and tests pooled within-head similarity against a label-permutation null. -assignment scores discovered (post-hoc) groups.

Between-head support (headsupport.py). The complement on the other axis: do distinct heads occupy distinct coordinates, as Section 3.1.4 argues non-negative heads should? On a trained checkpoint, each head’s usage-weighted coordinate profile and their mean pairwise cosine are those of Eq. (6), with $\bar{\mathbf{p}}_i$ the head’s mean assignment over the evaluation tail at the run’s final temperature, from the model’s own forward. The *decoded* variant replaces each atom by $|\mathbf{a}_{ij}^{\text{dec}}|$, with $\mathbf{a}_{ij}^{\text{dec}} := \text{dec}(\mathbf{a}_{ij}) - \text{dec}(\mathbf{0})$ the output-space direction steering acts in; for an `sae.py` run, which has no dictionary space, the profile of a head or group $\mathcal{G}$ is $\pi_{\mathcal{G}} = \sum_{j \in \mathcal{G}} \mathbb{E}[z_j] |\mathbf{w}_j^{\text{dec}}|$. Non-negative profiles overlap heavily by chance, so ω is scored against a null that permutes each head’s profile coordinates independently (50 draws), keeping every head’s own distribution and destroying only the alignment between heads. A head’s *spread* is its profile’s participation ratio as a fraction of the dimension, $(\sum_q \pi_{iq})^2 / (e \sum_q \pi_{iq}^2)$, with d in place of e when decoded; a partition into h blocks needs it near $1/h$.

Feature splitting and seed stability (splitting.py, partition.py, headcontrib.py). For a coarse model (1) and a fine model (2) on shared rows, each latent j' of model (2) is classified per parent latent j of model (1) by two-way firing containment at threshold τ : a *child* satisfies $\Pr(j \mid j') \geq \tau$ and $\Pr(j' \mid j) < \tau$, with $\Pr(j \mid j')$ the fraction of the rows where j' fires on which j fires too; a *match* satisfies both. Splits (≥ 2 children) are scored by union coverage of the parent’s firing set and pairwise child overlap. Decoder geometry adds absorption candidates (decoder-cosine twins that do not co-fire) and each model’s nearest-neighbor decoder-cosine distribution vs. a random baseline. For two same-architecture runs differing only in seed, the match fractions are the reproducible-feature fractions. For hard-code Ontologizers, `partition.py` matches heads across runs by the NMI of their partitions (what a head separates), and `headcontrib.py` by the cosine of their contributions (what a head writes): each head’s decoded output contribution in the whitened frame, less $\text{dec}(\mathbf{0})$, centered over rows — uncentered, the non-negative dictionary’s shared offset reads as agreement. With $q_{ni}^{(1)}$ the entry head i of run (1) selects on row n , $\mathbf{q}_i^{(1)}$ those selections over all rows, and $\phi_{ni}^{(1)}$ the head’s whitened decoded contribution on row n ,

$$\begin{aligned} [M^{\text{part}}]_{ii'} &:= \text{NMI}(\mathbf{q}_i^{(1)}, \mathbf{q}_{i'}^{(2)}), \\ [M^{\text{contr}}]_{ii'} &:= \cos(\text{vec}_n[\phi_{ni}^{(1)} - \bar{\phi}_i^{(1)}], \text{vec}_n[\phi_{ni'}^{(2)} - \bar{\phi}_{i'}^{(2)}]), \end{aligned} \tag{28}$$

where $\text{NMI}(\mathbf{q}, \mathbf{q}') := 2I(\mathbf{q}; \mathbf{q}') / (H(\mathbf{q}) + H(\mathbf{q}'))$ over the rows’ empirical distribution, $\bar{\phi}_i$ is the mean over rows, and each cosine runs over the contributions concatenated

across rows. Each matrix is matched one-to-one (Hungarian) by the pairing of heads that maximizes its summed entries, and the reported agreement is the mean over matched pairs. The null repeats this with run (2)’s rows shuffled, and the positive control matches the run against its own earlier checkpoint. `headcontrib.py` also splits each head’s contribution size into its layer’s input gain and its own decoded atoms. Since a head’s contribution on row n is its layer’s gain g_n times the decoded atom $\mathbf{a}_{i,q_{ni}}^{\text{dec}}$ of the entry it selected,

$$\text{size}_i := \frac{\sum_n \|\phi_{ni} - \bar{\phi}_i\|^2}{\sum_n \|\mathbf{w}^{1/2} \odot (\mathbf{x}_n - \bar{\mathbf{x}})\|^2} \approx \frac{\mathbb{E}_n[g_n^2] \mathbb{E}_{j \sim \bar{\mathbf{p}}_i} \|\mathbf{a}_{ij}^{\text{dec}} - \bar{\mathbf{a}}_i^{\text{dec}}\|^2}{\mathbb{E}_n \|\mathbf{w}^{1/2} \odot (\mathbf{x}_n - \bar{\mathbf{x}})\|^2}, \quad (29)$$

with $\bar{\mathbf{p}}_i$ head i ’s usage over entries (its mean assignment) and $\bar{\mathbf{a}}_i^{\text{dec}}$ its usage-weighted mean decoded atom: gain squared times atom spread squared, up to the gain’s small variation. With a reference seed pair (a flat arm) it compares agreement at matched head size.

Text fidelity (`textfid.py`). The “loss recovered” analogue: decode $\mathbf{x}$ and $\text{recon}(\mathbf{x})$ through M2M100 (forced `eng_Latn`, both rescaled to the reference SONAR norm, so the metric sees direction error) and report chrF [31] and exact-match between the decodes, plus the decoder’s teacher-forced per-token NLL of the *original*’s decode under the original (ceiling), the reconstruction, and the corpus-mean embedding (floor — which decodes to the generic text). Loss recovered places the reconstruction between the two,

$$\text{recovered} := \frac{\text{NLL}_{\text{floor}} - \text{NLL}_{\text{recon}}}{\text{NLL}_{\text{floor}} - \text{NLL}_{\text{ceil}}}, \quad (30)$$

the form GPT-2’s loss recovered takes with cross-entropy in place of NLL (Section 6). Under the steering frame this doubles as the collateral ceiling: a code that cannot preserve text cannot steer text with low collateral.

Language probing (`langprobe.py`, `headlang.py`, `headscript.py`). SAE Bench-style k_{probe} -sparse probing with the free 86-language labels: top latents per language by two-sample t-statistic on training rows; held-out one-vs-rest F1 at $k_{\text{probe}}=1$ (single latent, fitted threshold), $k_{\text{probe}}=4, 16$ (logistic probes), and a dense probe on the raw embedding as ceiling. Macro-F1 over languages. *Frame caveat: this measures whether a factor lives in individual latent magnitudes. A classifier head carries its factor in its assignment (partition), which single-latent probes cannot see.* `headlang.py` is the matching instrument: per head, the NMI between its partition and the label, null-corrected by a label shuffle, plus a held-out probe on the best h' heads’ one-hot codes against a dense ridge ceiling. NMI here normalizes by the arithmetic mean of the two entropies, so its attainable maximum is $2 \min(H_{\text{head}}, H_{\text{lang}}) / (H_{\text{head}} + H_{\text{lang}})$ — 0.875 for a 5-bit head against the 6.43-bit language label. `headscript.py` repeats

this for script and scores per-label cells and cross-head conjunctions; `headeta.py` and `topichead.py` measure each partition’s share of embedding variance,

$$\eta^2(C) := \frac{\sum_{C_c \in C} |C_c| \|\bar{\mathbf{x}}_{C_c} - \bar{\mathbf{x}}\|^2}{\sum_n \|\mathbf{x}_n - \bar{\mathbf{x}}\|^2}, \quad \mathbb{E}[\eta_{\text{random}}^2] \approx \frac{|C| - 1}{n_{\text{rows}} - 1}, \quad (31)$$

for a partition C of n_{rows} rows into cells C_c , with $\bar{\mathbf{x}}_{C_c}$ the mean of the rows in cell C_c , and report it raw and against the baseline a random partition into $|C|$ cells scores.

Shrinkage (refit.py). Gao-style refinement: freeze each sample’s support $\mathcal{S}_n$ and solve the whitened least-squares coefficients on it,

$$\mathbf{z}_n^\star := \arg \min_{\mathbf{z}: z_j=0 \ \forall j \notin \mathcal{S}_n} \|\mathbf{w}^{1/2} \odot (\mathbf{x}_n - \mathbf{b} - W\mathbf{z})\|^2, \quad (32)$$

where the columns of W are the code’s decode directions and $\mathbf{b}$ its offset, by masked normal equations (silent head slots forced to zero), and compare FVU_w . Raw–refit is magnitude (shrinkage) error; refit is selection error. The Ontologizer shares the machinery via its exact linearity $\hat{\mathbf{x}} = W_{\text{flat}} \mathbf{p}_{\text{flat}}$, with W_{flat} measured by decoding one-hot constant codes; its supports are the top- ν deviation entries. Points with support $\geq d$ are excluded (underdetermined). Note the raw convention: `refit.py` truncates the soft forward’s P (non-adaptive), whereas `pareto.py` truncates inside the forward, letting later layers reclassify against the truncated prefix residual.

Steering fidelity in text (steerfid.py). Unit-normalized native steering directions — SAE decoder directions, bilinear eigenfeatures, or the Ontologizer’s withArgs set-intervention delta — applied at matched embedding magnitudes. Effect is cycle consistency: the steered decode is re-encoded through SONAR and the target feature’s activation gain is normalized by its firing-conditioned quantiles, with the baseline taken from the *cycled* base decode so both sides share roundtrip degradation. Collateral is $1 - \text{chrF}$ against the base decode. Random unit directions give the control curve, and *net* effect is feature effect minus that control.

Steering fidelity in embedding space (steerembed.py). Drops the text round trip and asks the causal question directly: move the input along a direction, re-run the model, and record whether the intended entry is now selected (*realization*) and what fraction of the *other* heads changed their selection (*collateral*). Directions: decode (the decode delta of forcing the entry), grad (the gradient of the target logit), margin (the gradient of the target’s margin over the head’s current winner), adjoint (the bilinear classifier’s top eigenvector via `BilinearBlock.rev`), and random. Strength is a fraction of each row’s norm. The *native* step is the model’s own intervention size: the output change of forcing the entry, or for an SAE latent its decoder direction at the latent’s mean nonzero activation. Algorithm 4 gives the procedure.

Algorithm 4: Steering scored in embedding space (`steerembed.py`). For an Ontologizer a target j^* is an entry, of some head i in some layer ℓ , and a choice is a head's argmax; for an SAE with heads a target is a latent and a choice is its head's winner (a top-1 head that abstains counts as its own outcome); for an SAE without heads a target is a latent, it is selected when active, and collateral is the share of the row's other active latents that drop out.

Input: model; held-out rows X ; strengths $\alpha \in \{0.25, 0.5, 1, 2, 4\}$; 64 targets per model, stratified by layer, each selected on at least 0.5% of reference rows

Output: realization and collateral per direction and strength, and at the native step

begin

foreach *target* j^* **do**

$X_{j^*} \leftarrow$ 32 rows on which j^* is *not* already selected

 record every head's choice on X_{j^*}

foreach *direction* $\text{dir} \in \{\text{decode}, \text{grad}, \text{margin}, \text{adjoint}, \text{random}\}$ **do**

$\mathbf{v}_n \leftarrow \text{dir}(\mathbf{x}_n, j^*) / \|\text{dir}(\mathbf{x}_n, j^*)\|$ // unit direction per row

foreach *strength* α **do**

$\mathbf{x}'_n \leftarrow \mathbf{x}_n + \alpha \|\mathbf{x}_n\| \mathbf{v}_n$; $\mathbf{x}'_n \leftarrow \mathbf{x}'_n \|\mathbf{x}_n\| / \|\mathbf{x}'_n\|$ // rescale to $\|\mathbf{x}\|$

 recompute every head's choice on the steered rows

 realization $\leftarrow \mathbb{E}_n \mathbf{1}[j^* \text{ selected on } \mathbf{x}'_n]$

 collateral $\leftarrow$ share of (row, other head) pairs whose choice changed

end

end

$\delta_n \leftarrow \hat{\mathbf{x}}_n(j^* \text{ forced}) - \hat{\mathbf{x}}_n$, or, for an SAE, the latent's decoder direction times $\bar{z}_{j^*}$, its mean activation when it fires // native step

 score $\mathbf{x}_n + \delta_n$, unrescaled, and the control $\mathbf{x}_n + \|\delta_n\| \xi_n$ for a random unit ξ_n

end

return *the means over targets*

end

Eigenvector orientation. The adjoint direction of a bilinear classifier is an eigenvector of its interaction matrix, which `eigh` returns only up to sign. Orienting it by the sign of its eigenvalue, with a leading-coordinate tie-break (`nlinear.orient`), realizes 0.427 of targets on the softmax model at strength 0.25, against 0.224 with the sign `eigh` happens to return and 0.446 when each sample is oriented separately. A weight-only convention should not be able to do that: a step of α along $\mathbf{v}$ changes $\mathbf{x}^\top B \mathbf{x}$ by $2\alpha \lambda(\mathbf{x} \cdot \mathbf{v})$, a property of the sample. It can on SONAR because the embedding is anisotropic, so $\mathbf{x} \cdot \mathbf{v}$ is dominated by $\mathbb{E}[\mathbf{x}] \cdot \mathbf{v}$, which has a consistent sign; on isotropic Gaussian inputs a weight-only direction raises the logit on about half the samples. On the hard code the oriented eigendirection is below random at every strength.

Linearized collateral (`steergeom.py`). Separates the direct effect of a step on the other heads of the target’s layer from the cascade through later layers, by linearizing that layer’s logits with every upstream code held fixed. Under residual forwarding an input step $\mathbf{v}$ moves the layer’s residual input by exactly $\mathbf{v}$, so with J_i the Jacobian of head i ’s logits at the layer input $\mathbf{e}_n$, head i ’s winner–runner-up margin and its first-order change along a step of $\alpha \|\mathbf{x}_n\|$ are

$$\mu_{ni} := [\mathbf{K}]_{nij_1} - [\mathbf{K}]_{nij_2}, \quad \Delta\mu_{ni} := \alpha \|\mathbf{x}_n\| \left((J_i \mathbf{v}_n)_{j_1} - (J_i \mathbf{v}_n)_{j_2} \right), \quad (33)$$

with $\mathbf{K}$ the logits and j_1, j_2 the head’s top two entries on row n . With i^* the target’s head, the linearized flip share is $\mathbb{E}_{n, i \neq i^*} \mathbf{1}[\mu_{ni} + \Delta\mu_{ni} < 0]$ at $\alpha=0.25$, reported for the decode, grad and random directions with the median margin μ . Two further measures use the same Jacobians: *off*, the mean $\|\mathbf{x}_n\| \|J_i \mathbf{v}_n\|$ over other heads and entries as a ratio to the random direction’s, which is below 1 when a direction avoids the other heads’ detectors; and *own*, the median first-order gain of the target over the incumbent per unit of $\|\mathbf{x}_n\|$, divided by the gap it has to close.

Intervention composition (`compose.py`). Treating heads as candidate causal variables: for pairs of forced assignments $(\ell_1, i_1) \rightarrow j_1, (\ell_2, i_2) \rightarrow j_2$ (8 random heads per layer, one random entry each; all same-layer pairs plus 60 cross-layer pairs), compare the joint intervention’s decode delta δ_{12} against the sum $\delta_1 + \delta_2$ of the single-intervention deltas on 512 held-out rows: median per-row cosine and relative error $\|\delta_{12} - (\delta_1 + \delta_2)\| / \|\delta_1 + \delta_2\|$, raw and whitened. Because the decode is exactly linear in the flattened classification, final-layer pairs compose exactly by construction; that stratum doubles as the measured numerical floor (the $\|\mathbf{r}\| \approx 31 \rightarrow \|\hat{\mathbf{x}}\| \approx 1$ decoder cancellation amplifies float32 rounding to ~ 0.02 relative error between differently-compiled variants — verified exact in float64). All other non-additivity is downstream reclassification: later layers re-reading the intervened residual.

Autointerp (autointerp.py). Four stages: harvest (top-activating rows, negatives, frequency-stratified feature selection, parameter-based decode embeddings), texts (deterministic mC4 stream replay recovers text by row index, verified against the language sidecar), describe, and score. Description modes: params, the M2M100 decode of a parameter-derived embedding (for entries, the $\mathbb{E}[\mathbf{p}]$ -origin code with one head forced one-hot; for SAE latents, $\mathbf{b}_{\text{dec}} + \alpha \mathbf{w}_j^{\text{dec}}$); pdev, the same decode with the pure-origin decode subtracted; eig, the bilinear eigenfeature decode; acts, an LLM summary of top-activating snippets; and cacts, a *contrastive* summary — the judge sees top snippets next to random corpus draws and must name what distinguishes them, blocking the degenerate corpus-level answer that plain summaries invite on noisy multilingual web text. Scoring is blinded detection over held-out positives and negatives, identical for all modes; per-feature F1, to be read against the always-match null of $F1 \approx 0.67$ on balanced items. The judge is Haiku (claude-haiku-4-5) via the Anthropic Messages API, rate-metered, with bounded reply length; failed judgements are never recorded (an earlier outage mode scored empty replies as $F1 = 0$). Six models: the Ontologizer, rungs k32/k32_b1/g160top1/g160softmax, and the dense control k5120.

Diagnostic scripts (experiments/ste-arm/). The straight-through and geometry results rest on further single-purpose instruments, described where used: inputgeom.py (geometry of the vector each classifier sees), blendablate.py (severing the pairing between a layer’s input and the sample’s own earlier reconstruction), gainablate.py (the norm side-channel), dictgeom.py (row collinearity and effective rank read from the weights), lastlayer.py (final-layer ablations against the deep-supervision prefixes), decodehead.py (decoding a head’s atoms as deviations from their mean), embedgeom.py (SONAR cloud geometry), entryshare.py (how evenly a layer’s entries contribute to the output, and whether usage, gain or atom norm drives the inequality), and bigatoms.py (how a layer’s large atoms are distributed among its heads, and whether the heads holding them reproduce).

C SUPPLEMENTARY RESULTS: THE SOFTMAX ONTOLOGIZER AGAINST SAEs

C.1 CATEGORICAL FORM

At 1M rows with split-half scoring and 20 size-matched nulls: g160top1’s trained partition is *perfectly* head-like (exhaustiveness 1.0000 vs. null 0.64; exclusivity 1.0000 vs. 1.51; sum-CV 0.25 vs. 1.03). Discovery on every unstructured run (m11264_k32, m5120_k32, k32_p5, m11264_k160) finds the *opposite* signature — exhaustiveness below null (ζ from -30 to -50), exclusivity above (ζ $+13$ to $+31$): spontaneous SAE group structure consists of co-firing topical clusters, not categorical slots, even when the affinity explicitly selects for anti-correlated same-context pairs. And g160softmax never fires distinctly at all: no coefficient ever exceeds $2/32$ on any test row — soft per-head codes converge to dense near-uniform mixtures,

exactly as the Ontologizer’s soft code does ($\mathbb{E}[\mathbf{p}]$ sits 0.06–0.16 bits from uniform in healthy layers; matching soft reconstruction requires ~ 16 of 32 deviations per head). **Categorical structure exists precisely when trained in as hard selection, never emerges from softmax competition under a reconstruction objective, and costs $\sim 60\times$ in FVU_w (0.274 vs. 0.0044) at the SAE’s 160 heads.** For steering this means the *soft* architectures’ classifier reading must be cashed out in deviations around $\mathbb{E}[\mathbf{p}]$, not in discrete winners — or the selection must be made hard during training, which Section 5 does at a far smaller cost once the head count is raised. Figure 15 shows one such soft head of `resid_nc`.

C.2 SEMANTIC COHERENCE IN DETAIL

Under the whitened metric the training objective optimizes, the Ontologizer’s mean ζ_{cos} is -24.4 (-17.6 raw; the fraction of heads below -2 is unchanged at 0.96), the single-layer soft variant’s -14.7 (unchanged) and the hard variant’s $+179.3$ ($+161.3$ raw).

Two heads of the reference softmax model cluster instead of contrasting: layer 2’s head 25 ($\zeta_{\text{cos}}=+346$) and layer 3’s head 28 ($+280$), with the next candidates in the partially frozen layer 4. Their sampled entries also decode to generic text, which suggests partial head collapse (entries that never differentiated) rather than semantic coherence; they have not been inspected further.

The description-text view of coherence is uninformative on plain summaries, which are generic for about 77% of entries and 49% of $k_{\text{SAE}}=32$ latents and so inflate the permutation null (pooled $\zeta=-0.25$). On the contrastive descriptions it acquires signal. Within-head description similarity clears the head-membership permutation null for the single-layer hard variant ($\zeta=+2.42$, 62 within-head pairs), is marginal for Ontologizer heads ($+1.92$, 155 pairs), and is absent for the single-layer soft variant ($+0.79$, 217 pairs). Against the geometry, this is the pattern that alternative values of one shared dimension predict: Ontologizer entries’ decode directions repel while their described contents weakly cohere. It is also the first measure to separate the Ontologizer from its single-layer soft rival, though only suggestively at this coverage (one global ζ per model over the campaign’s sampled features).

C.3 FEATURE RELIABILITY

Seed stability ($k32$ vs. $k32_s43$; identical architecture and data, FVU_w 0.4994 vs. 0.4967): matched-feature fractions are **0.31 / 0.24 / 0.15 / 0.03** at $\tau=0.3/0.5/0.7/0.9$. Roughly three quarters of individual features are not reproducible across initializations at moderate containment. Cross-width splitting ($k32 \rightarrow m11264_k32$): 190 of 5114 live parents split at $\tau=0.7$ (median 3 children; union coverage 0.37; child overlap 0.42), with 441 absorption candidates and a heavy nearest-neighbor decoder-cosine tail ($p50$ 0.234 vs. random 0.119): splitting is real but is not the




FIGURE 15 · One softmax head, block by block. The blocks of Figure 4 for head 0 of layer 0 of resid_nc (step 372,600, $t=0.03$) on $n=384$ eval-tail rows; one seed. Usage is spread across entries (use 29.9 of $k=32$), but each row is a near-uniform mixture (row 23.5; a), so the neighbor graph (b) is close to the heat kernel alone, with hub columns rather than blocks. $\kappa_i=0.138$, $\omega=0.901$, $e=4096$. As in Figure 4, the model was not trained with the statistics that use the graph.

dominant organization; topical co-firing is. For a steering generator the unit that must reproduce is the head’s partition rather than the individual feature vector; Section 5.5 measures that on the hard-code Ontologizer and finds it reproduces far *worse* than SAE features do.

C.4 TEXT-SPACE FIDELITY (COLLATERAL CEILING)

model	chrF	exact	NLL (ceil. ≈ 1.56 – 1.63 , floor ≈ 3.8)
corpus-mean floor	≈ 0.20	0	3.60–3.89
k32 tier (incl. p5, m11264)	0.23	0.00	3.00–3.08
k32_bl	0.21	0.00	3.19
k160 / g160top1	0.28–0.29	0.00	2.34–2.37
onto soft	0.538	0.090	1.658 (gap 0.026)
g160softmax	0.592	0.119	1.563 (gap 0.007)
k5120	0.856	0.592	1.587 (gap 0.000)

TABLE 14 · Downstream text fidelity, 512 tail rows, forced English decodes.

The FVU–text mapping is sharply nonlinear: $FVU_w \approx 0.5$ decodes to same-topic-different-facts text barely above the generic floor, while the Ontologizer’s soft code and g160softmax are near the likelihood ceiling — almost all decodable content survives. Across the eleven Ontologizer runs with both numbers (the selection sweep and the straight-through arms; Table 16), “loss recovered” is a monotone function of FVU_w (rank correlation essentially -1), so text fidelity is a check that the whitened objective transports to text, not an independent axis — and it does. This matters doubly under the steering frame: a steering primitive that routes through a code which cannot preserve text (the entire $k_{SAE} \leq 160$ sparse tier) starts from ~ 0.75 collateral before any intervention, whereas the dense-mixture classifier codes start near zero.

C.5 CONCEPT (LANGUAGE) LOCALIZATION

Language remains distributed in every model (best single-latent macro-F1 0.119 vs. dense ceiling 0.676; script-distinctive languages such as Chinese reach ~ 0.98 from one latent, same-script families smear), and localization *anticorrelates* with reconstruction across the family. The Ontologizer’s entry probabilities are decisively the worst localizers ($F1@1 = 0.022$). Read on SAE terms this is an indictment; read on classifier terms it is expected: a head encodes a factor in its *assignment*, not in any single entry’s probability, and partition-level analysis of `resid_nc` found layer-1 heads whose partitions carry script/family mutual information, one of which ((L1,h8,k17)) was causally verified to force Burmese output. The partition-level instrument (`headlang.py`) was later run on a config-matched softmax reference and on the hard-code arms (Section 5.6): the softmax reference’s best head reaches

model	F1@1	F1@4	F1@16	dense
k160	0.119	0.393	0.473	0.676
k32 ($m=5120$)	0.089	0.347	0.465	0.676
m11264_k32	0.078	0.364	0.471	0.676
k32_bl	0.084	0.228	0.368	0.676
g160softmax	0.068	0.162	0.328	0.676
g160top1	0.066	0.200	0.302	0.676
k32_p5	0.061	0.322	0.454	0.676
k5120	0.064	0.130	0.237	0.676
onto (entries)	0.022	0.021	0.117	0.677

TABLE 15 · Macro one-vs-rest F1 over 86 languages, held-out tail.

a null-corrected NMI of only 0.043, a quarter of the hard code’s. Language factors are carried by a *few* heads’ partitions, weakly, and most heads classify something else.

C.6 SHRINKAGE DECOMPOSITION

TopK SAEs have small magnitude error (share of raw FVU_w removed by support-frozen least squares: 2.2–2.4% for k32 variants, 8.5% for k160, 0.6% bilinear, 1.9% p5); g160top1 is essentially perfectly calibrated (0.3%) — its reconstruction deficit is pure selection. The Ontologizer’s truncated deviation codes are magnitude-heavy: 16% at $v=1$, 33% at $v=2$, i.e. the pin-and-renormalize convention mis-scales what it keeps (supports of $\geq d$ coefficients are excluded as underdetermined). Separately, the gap between non-adaptive truncation (freeze soft P , truncate, decode linearly) and in-forward truncation (later layers reclassify against the truncated prefix residual) is large — 0.268 vs. 0.094 at $v=8$: downstream layer compensation is worth $\sim 3\times$ under truncation. This was the first quantified reconstruction credit for the residual stack (Section 5.3 measures a larger one), and it is steering-relevant: the stack actively re-routes around code edits, which cuts both ways — robustness to small interventions, and, as Section 5.4 shows, widespread collateral.

C.7 STEERING THROUGH THE TEXT ROUND TRIP

On the SAE side only a smoke-scale run exists (k32, 3 features): cycle-consistency effect rises with magnitude ($1.2 \rightarrow 1.8$ over magnitudes $0.5 \rightarrow 1.0$ relative to the unit input norm; hit rate $0.67 \rightarrow 0.89$), while collateral at these magnitudes (≈ 0.75) is near the random-direction control. Across the Ontologizer selection sweep and the straight-through arm, at 64 features per run:

Steerability does not track reconstruction. The best steerer is top8 at FVU_w 0.037, the most faithful model (softmax_shm) is among the weakest, and every _shm variant steers about three times worse than its twin (top2 0.474 vs. top2_shm 0.141; top8 0.580 vs. 0.193). The straight-through arm’s effect is positive but nearly flat

run	FVU _w	chrF	loss rec.	net @0.25	@0.5	@1.0	hit @1.0	rand. hit
top8	0.0370	—	0.940	0.131	0.303	0.580	0.713	0.277
top2	0.3516	0.248	0.532	0.125	0.289	0.474	0.564	0.217
top16	0.0052	—	—	0.065	0.119	0.243	0.313	0.139
top8_shm	0.0287	0.416	0.956	0.072	0.096	0.193	0.291	0.142
softmax_shm	0.0007	0.694	0.998	0.032	0.072	0.146	0.150	0.064
top2_shm	0.1911	0.303	0.738	0.022	0.070	0.141	0.154	0.041
top4_shm	0.0761	0.360	0.897	—	—	—	—	—
ste_h76	0.2293	0.282	0.702	0.072	0.096	0.107	0.068	0.062
top1	0.7190	0.180	0.123	0.011	0.026	0.048	0.069	0.020

TABLE 16 · Text fidelity and text-cycle steering (steerfid.py, -n-features 64) across selection rules. Net effect is feature effect minus the random-direction control. A further run (topk4_rev, FVU_w 0.164) reads net 0.520 at 1.0.

in magnitude, and its hit rate (0.068) barely clears its own random control (0.062). Two cautions on this instrument. Sixteen features is too few to quote: at 16 features the straight-through arm’s net effect was negative at every magnitude and at 64 it is positive throughout. And the hit rate is not comparable across hard and soft codes: under a hard code an activation is exactly 0 or 1, so a “hit” demands the argmax land exactly on the target entry, where a soft code need only cross a fractional threshold.

The text cycle is too weak to support conclusions. Measured on an *unsteered* round trip (512 held-out rows), the cycled embedding retains barely more than the shared corpus direction: $\cos(\mathbf{x}, \text{cycle}(\mathbf{x}))$ is 0.378 for ste_h76 and 0.380 for softmax, against 0.314 between random pairs, and a head’s argmax survives the round trip on 12.1% and 10.5% of rows against 3.1% chance (34.8% at layer 0, 4.6% by layer 4 for the hard code). A perfect intervention could therefore be observed at most ~12% of the time — the same order as the hit rates the instrument reports. Reconstruction into embedding space is fine; generation is where the signal goes.

D SUPPLEMENTARY RESULTS: TRAINING, REPRODUCIBILITY AND STEERING OF THE SOFTMAX ONTOLOGIZER

These results come from the project proposal’s evaluation and concern the softmax Ontologizer. Two models recur. resid_nc (372,600 steps) is the reference model behind the softmax tables above. resid_nc_hm is the same configuration plus the batch mean-entropy bonus s_{KL_m} , trained to 4.38M steps; results here use its final checkpoint unless stated otherwise.

D.1 HEAD FREEZING AND THE MEAN-ENTROPY BONUS

A head is counted frozen when its top classification share is at least 0.95 and its churn at most 0.02. In `resid_nc` the top layer (layer 4 of 0–4) begins freezing near step 190k, well after the temperature anneal, reaches 27 of 32 heads by step 210k and 28 at its peak, and is still frozen at the last checkpoint (Figure 16). Section G.4 shows the same layer to be inert once the prefix error is small, which is why the freeze costs so little reconstruction. In `resid_nc_hm` the late freeze never occurs. What remains is an early annealing transient — 22 of 32 heads at layer 1 at step 10k, with 16, 11 and 4 of 32 at layers 2, 3 and 0 — that recovers completely by step 30k, leaving the final checkpoint unfrozen everywhere. **The mean-entropy bonus prevents the late freeze.**




FIGURE 16 · Frozen heads per layer over training. (a) `resid_nc_hm`, with the mean-entropy bonus, has only an early annealing transient, peaking at 22/32 in layer 1 at step 10k and recovering to zero everywhere by step 30k. (b) `resid_nc` develops a post-anneal layer-4 freeze from ~190k that persists through its last checkpoint.

Four retrains of the `resid_nc_hm` configuration (372k steps each) isolate the transient: a zero-change control and three single-knob variants. The control reproduces it head for head (22/32 at layer 1 at step 10k, zero everywhere by 30k) and shows no late freeze. Stretching the temperature anneal from 50k to 150k steps attenuates and spreads it (peak 14/32 at layer 1 at step 30k, clear by 50k). Raising winner dropout from 0.1 to 0.25 tracks the control to within two heads at the peak, and holding the classification-noise floor at its starting value changes nothing (Figure 17). The transient belongs to the temperature schedule alone, and every retrain ends unfrozen.

An HSIC bottleneck in place of the auxiliary losses. The candidate objective maximizes dependence with the reconstruction target while penalizing dependence between heads. It is distinct from the pairwise head-CKA penalty of Section E.6, which was added to the straight-through arm’s losses rather than replacing them. Both arms were retrained on the `resid_nc_hm` configuration to 372k steps.




FIGURE 17 · The early transient under the four retrains, layer-1 frozen count. The noise-floor variant overlaps the control exactly and the dropout variant tracks it to within two heads at the peak; only stretching the temperature anneal moves the curve, trading peak height for duration.

Replacing the four auxiliary terms with the HSIC objective destroys head health: a late, persistent freeze takes 77 of 160 heads by the final checkpoint, one layer entirely (layer 3, 32/32 from step 80k), while reconstruction lands slightly behind the control (tail-mean training MSE 3.7×10^{-5} against 3.4×10^{-5}). Adding HSIC on top of the auxiliary terms is worse on both counts: MSE 4.4×10^{-5} , and 41 of 160 heads still frozen at the end, layer 3 fully frozen by 80k with only partial thaw (Figure 18). The failure is `resid_nc`'s late freeze rather than the transient, which puts the bottleneck in direct opposition to the mean-entropy bonus: the bonus prevents exactly what HSIC induces. As formulated it is neither a replacement nor an addition.




FIGURE 18 · Total frozen heads over training for the three loss configurations, all on the `resid_nc_hm` configuration at 372k steps. The auxiliary-loss control's freeze is the early transient and resolves; both HSIC arms develop the late, persistent freeze instead, and the auxiliary terms only partially thaw it.

D.2 DEVELOPMENT OVER TRAINING

Tracking `resid_nc_hm`'s full 448-checkpoint history applies the z-score machinery of Section 4.3 to the refit end-to-end direction matrix rather than to raw dictionary decode directions. The two are not directly comparable: raw decode directions on `resid_nc` repel ($\zeta_{\text{cos}}=-17.6$), while refit directions on `resid_nc_hm`'s final checkpoint cohere ($\zeta_{\text{cos}}=+72$); they differ in both instrument and substrate, and the refit frame folds in the shared decoder and cross-layer interactions. On that statistic the shape is rise, peak and partial relaxation (Figure 19): mean within-head ζ_{cos} climbs from 1.8 at step 10k to a peak of 119 near step 370k (median 123, every head above $\zeta_{\text{cos}}=2$), then relaxes over the remaining 4M steps to 72 at the final checkpoint (median 105, 0.79 above $\zeta_{\text{cos}}=2$). The code itself moves the other way after the annealing transient: per-sample entropy returns to 4.5 of a 5-bit maximum and usage KL falls toward uniform. Relative to initialization, the consolidation is in the geometry the classifications write, while the classifications stay soft.




FIGURE 19 · The 448-checkpoint training history of `resid_nc_hm`. (a) Within-head coherence of refit end-to-end directions, z-scored against size-matched nulls; rise, peak near step 370k, partial relaxation. (b) Mean per-sample classification entropy returns toward the 5-bit maximum after the annealing transient.

D.3 SCALING: A STAR SWEEP

Seven configurations around the shared $h=32, k=32, l=5$ shape, varying one axis at a time, each trained 4 epochs (62k steps against the retrain recipe's 372k), test whether head coherence rises with parameter count. No measured axis delivers the rise: coherence shows no positive relation to parameter count (Spearman +0.14 for mean ζ_{cos} , -0.54 for the fraction above $\zeta_{\text{cos}}=2$), usage imbalance grows with size (+0.86), and the cleanest per-axis trend runs opposite to the prediction, coherence falling monotonically as entries per head grow ($k=16/32/64$: mean ζ_{cos} 20.2/6.3/3.7). The widest model ($h=64$) has not converged at this budget (tail MSE 20× the others), so its collapse is not evidence (Figure 20). Seven points at short

training are weak evidence either way; full-length retrains are the missing control.




FIGURE 20 · Head coherence against each parameter axis of the star sweep: heads (a), entries (b) and layers (c), all runs at 4 epochs (62k steps). No axis delivers the predicted rise; the entries axis falls monotonically. The open marker in (a) is the unconverged $h=64$ run.

D.4 SEED STABILITY OF SOFT PARTITIONS

A second run of the `resid_nc_hm` configuration from a different seed, compared at matched training (370k steps on both sides), matches essentially nothing (Figure 21): no head pair and at most 0.3% of atoms clear a 0.3 cosine threshold on dictionary geometry, and activation-frame containment over a shared evaluation batch is at most 0.001 at any threshold, against 0.31/0.24/0.15/0.03 for the `k32` SAE and its seed replica across the four thresholds (0.3–0.9). The metric has headroom: the same run against itself 2,600 steps apart scores 1.000 on geometry at every threshold and 0.65 at the loosest activation threshold (0.20 at the strictest). At matched training the softmax Ontologizer’s partitions are less reproducible across seeds than SAE latents, the same verdict Section 5.5 reaches for the straight-through arm with partition and contribution measures. A replica carried to the full 4.4M steps is the remaining control.

D.5 STEERING AGAINST SUPERVISED BASELINES

A matched-budget overlay on `resid_nc_hm`’s final checkpoint compares forced classifications against difference-of-means and linear-probe directions on language steering, with a random-direction floor (Table 17). Each arm wins its own effect frame, as expected. On the shared difference-of-means frame, forcing a classification moves decodes along the target direction 0.43× as far as supervised steering does, and well ahead of probe directions. Every arm’s collateral, supervised steering included, sits at the random-direction floor (0.72–0.78), the same floor the text round trip shows for SAE steering (Section C.7). On naturalness the arms are so far indistinguishable (Figure 22): `chrF` against unsteered decodes falls from ~ 0.35 at magnitude 0.25 to ~ 0.15 at 1.0 for the difference-of-means and forced-classification




FIGURE 21 · Cross-seed matched-entry fraction on the activation frame by threshold. The same-run pair bounds what the metric can detect (0.65 at $\tau=0.3$); the k32 SAE pair reproduces partially; the softmax Ontologizer’s partitions reproduce not at all.

arms, with probe and random decaying in parallel from ~ 0.38 to ~ 0.20 . A blinded human-rating sheet (120 items) is prepared but unscored.

arm	effect (native)	effect (dm frame)	hit rate	collateral
difference of means	0.337	0.337	0.250	0.757
classification forcing	0.394	0.146	0.125	0.778
linear probe	0.293	0.028	0.000	0.720
random direction	—	—	—	0.721

TABLE 17 · Language steering under matched intervention budgets on `resid_nc_hm`’s final checkpoint; medians over 24 language–magnitude conditions per arm. Effect is displacement along the arm’s own direction (native) and along the shared difference-of-means direction (dm frame); hit rate is the fraction of steered decodes reaching the median projection of genuine target-language rows on the shared frame; collateral is content change on unrelated dimensions.

A classification as a training objective. A REINFORCE policy over steering directions on frozen `resid_nc_hm`, rewarded by one classification’s quantile-normalized score at a matched intervention budget, converges cleanly in the differentiable in-embedding tier, saturating within roughly 40 steps at flat collateral (Figure 23). Neither the decode-cycle tier nor a convergence race against paired steering has been run.

D.6 HEAD-LEVEL AUTOINTERP

The per-latent lenses of the description-specificity results mis-measure classifier codes, whose content lives in partitions, so a head-level pass over `resid_nc` (step 372,600) tests the claim directly; its frozen top layer (Section D.1) is a caveat for any




FIGURE 22 · chrF between steered and unsteered decodes by intervention magnitude. All four arms, the random direction included, damage the decode at nearly the same rate; no arm buys naturalness at matched magnitude.




FIGURE 23 · REINFORCE against a single frozen classification's quantile-normalized score, in-embedding tier. The policy saturates the reward within roughly 40 steps (a) while collateral stays at its starting level (b).

layer-4 heads in the sample. Descriptions decoded from dictionary parameters fail at the partition level as they do per latent (median self accuracy 0.02). Activation-grounded head descriptions do no better than a null pairing at this sample size (20 heads per condition, median accuracy 0.083 against 0.125), and a head’s detection accuracy is uncorrelated with the mean describability of its latents (Spearman +0.07 over 30 head–lens pairs, 15 heads under each lens; neither per-lens value distinguishable from zero). The two lenses disagree and neither shows above-chance head-level signal, so the mis-measurement argument predicts an advantage this sample does not deliver.

Describability against density. Across the six runs of the description-specificity table, the correlation between log nominal density and activation-description F1 is -0.894 , so density alone carries about 80% of the variance. Fitting contrastive-description F1 linearly in log firing rate on the four `sae.py` runs with a nonzero firing rate (slope -0.069 , correlation -0.766), and holding the Ontologizer out, predicts 0.421 against an actual 0.392, a residual of -1.8 standard errors: the Ontologizer sits on the trend, and `g160top1` lies 0.120 above it. The fit is thin. It rests on four runs and a weak correlation, it excludes `g160softmax`, whose thresholded firing rate is exactly zero, and the $2/k$ threshold is not monotone in true density.

E SUPPLEMENTARY RESULTS: THE STRAIGHT-THROUGH ARM

E.1 CALIBRATION

A naive port (live configuration, `select="ste"`) trains but the classifier learns nothing. The root cause is one property: **the argmax forward is scale-invariant, so nothing pushes the logits up**; they sit at their initialization scale (within-head standard deviation 0.0015) for the whole run. Three settings follow from it. *Logit noise must be relative*: $\sigma_K=0.02$ is thirteen times the logit spread, so the argmax is chosen by noise, while $\sigma_K=0$ removes the only exploration and the codebook collapses; `-noise-k batchnorm` (Eq. (17)) scales noise by the logit RMS. *The surrogate temperature must track the logit scale*: at the live floor $t=0.03$, spread/t is 0.05 and the surrogate is flat, so what improves is the dictionary adapting to a fixed random partition. *s_{KL_m} must be rescaled*: 10^{-6} is calibrated against a whitened MSE of $\sim 3 \times 10^{-6}$, and a hard code trains at an MSE orders of magnitude larger, where the term is inert and heads collapse.

t	spread/t	$\max p$	effective entries/head	raw FVU
0.01	0.15	0.043	23.4	178
0.0015	0.78	0.135	26.4	193
0.0005	3.38	0.681	14.0	59.2
0.00015	16.84	0.974	3.8	7.3

(4000 steps, lr 10^{-5} , -noise-k batchnorm -sd-k 0.3; one seed per arm, which ranks configurations but does not predict final quality.) Sharpening the surrogate is what makes the classifier train at all and improves reconstruction 25-fold, but it trades directly against head usage, which s_{KL_m} exists to hold. At sonar.py’s learning rate 5×10^{-5} every sharp-surrogate arm collapsed abruptly after ~ 1700 steps; 10^{-5} is stable. The shipped arm is $t=1.5 \times 10^{-4}$, lr 10^{-5} , -noise-k batchnorm -sd-k 0.3, $s_{\text{KL}_m}=10^{-4}$, trained 24 epochs (369,500 steps). The temperature is hand-matched to a logit scale measured once; the principled version normalizes logits per head or controls t against a spread/ t setpoint.

E.2 CAPACITY BEYOND BITS, AND TRAINING DYNAMICS

Distance from the Gaussian reference. Measured against the Gaussian reference at each code’s own rate, the full stack sits further from it than one layer’s prefix does. Each ratio bounds that code’s distance from the optimum at its rate from below, since the optimum for real data may lie below the reference:

code	Gaussian reference	achieved	ratio
380 bits (one layer’s prefix)	0.4208	0.6427	1.53×
1900 bits, 5×76	0.0499	0.1535	3.08×
1900 bits, 5×380 , $k=2$	0.0499	0.1844	3.70×

Bits are not the whole capacity. At identical bits, 380 binary heads reach 0.1844 against 76 thirty-two-way heads at 0.1535: bits go as $h \log_2 k$ but the codebook goes as $h \cdot k$ (3800 atoms against 12,160). The binary arm gets within 17% on $2.9\times$ fewer parameters, with features that are one bit each — the form that is actually nameable.

Realized bits, and what they bound. This paragraph follows information theory in writing random variables in uppercase. For a hard code and a random input $\mathbf{X}$, write $A_{\ell i}$ for the entry head (ℓ, i) chooses and $\mathbf{A}$ for all of them. *Realized bits* are the summed entropy of each head’s usage over the data,

$$H_{\text{real}} := \sum_{\ell, i} H(A_{\ell i}) = \sum_{\ell, i} (\log_2 k - \text{KL}(\bar{\mathbf{p}}_{\ell i} \parallel \text{uniform})) \geq H(\mathbf{A}) = I(\mathbf{X}; \mathbf{A}), \quad (34)$$

with $\bar{\mathbf{p}}_{\ell i}$ head (ℓ, i) ’s usage distribution, its mean assignment; summed over layers this is $h(l \log_2 k - \text{KL}_m)$ in the logged statistic’s terms. Because the code is a deterministic function of the input, its joint entropy is exactly the information it carries about the input, and subadditivity puts the sum of marginal entropies above it, with equality only when the heads are independent. Realized bits therefore measure how evenly each head uses its entries, not how much the code says: two codes at the same H_{real} can carry different information if their heads are redundant




FIGURE 24 · Realized code bits over training, as a share of the nominal 1900, from the logged KL_m less its one-batch sampling floor. The shipped initialization’s head usage collapses to under half before recovering to 90% (the curve’s 91% over-corrects heads with dead entries); at the fixed initialization every head in the stack and the flat arm reaches near-uniform usage within 5k steps. An upper bound on the information carried, not a measure of it (Eq. (34)).

to different degrees, and the joint entropy, which needs a density estimate over k^{lh} outcomes, is not measured here.

Training dynamics. On the shipped arm the training-log MSE (the deep-supervision prefix mean) is flat from step 237k while within-head row collinearity keeps falling ($0.812 \rightarrow 0.686$, against ~ 0.64 for random non-negative rows). Its heads’ usage collapses early and never evens out (Figure 24): realized bits fall to 43% of nominal near step 6k, recover to 85% by 20k as the s_{KL_m} rescale turns the usage collapse around, and end at 90%. At the fixed initialization every head’s usage is near-uniform by step 5k (99% of nominal), in the stack and the flat arm alike, and stays there. Realized bits come from the logged KL_m , $h(l \log_2 k - KL_m)$, less its sampling floor: KL_m is a plug-in estimate on one batch of 256 rows, which reads $(k - 1)/(2b \ln 2) = 0.087$ bits per layer even for perfectly balanced usage, and on the converged fixed-init arms the logged value is almost entirely that floor. Corrected, the endpoints match the held-out measurement of Table 5 to within a bit; the correction assumes every entry is live, so it overstates the shipped arm by 14 bits.

E.3 DEPTH: ELIMINATIONS AND MECHANISM

Over training. Scored held out at every 10k checkpoint (Figure 25), the stack leads the flat arms from the first checkpoint on: the ratio to the matched flat arm is $1.59\times$ at step 10k, peaks at $1.83\times$ near 200k and ends at $1.79\times$. Neither shape has converged. Over the last quarter the stack improves 0.86% and the flat arms 2.3–2.75%, so the gap is narrowing, slowly; where it ends needs longer runs. Seed




FIGURE 25 · Held-out FVU_w over training for the 5×76 stack and the 1×380 flat arms at matched bits and parameters, with seed replicas. The flat arm with the stack’s entropy pressure ($s_{KL_m}=10^{-4}$) is the one compared in Section 5.3.

replicas overlap at every checkpoint.

The comparison went through two wrong versions first, and the eliminations are informative. On the shipped initialization the flat arm read 0.5822 at 484 realized bits (25.5%), a $2.5\times$ gap. Correcting the per-head entropy pressure — KL_m averages over a layer’s heads and the loss sums over layers, so per-head pressure scales as s_{KL_m}/h per layer and the flat arm had trained at a fifth of the stack’s — took it to 1199 bits and 0.4307. A frozen ZCA whitener (`make_zca.py`) took the classifier’s input from effective dimension 107 to 901 and the code to 1900 of 1900 bits, every head at 4.999 of 5 bits, more than the stack itself realized — and still reconstructed $1.7\times$ worse (0.3842). Two survivors from that line: **usage is not what depth buys** (bits are not the objective either: dropping s_{KL_m} $100\times$ costs the stack 176 bits and *improves* its error to 0.2095), and **a learned input transform collapses where a supplied one does not** — a jointly trained Linear encoder ends with every head on exactly two entries at FVU_w 0.6853, worse than no encoder.

The input each classifier sees. `inputgeom.py`, on the exact vector handed to each classifier (shipped arm):

layer	mean direction	mean pairwise cos	effective dim / 1024
0	0.5638	0.3146	107.5
1	0.0932	0.0087	611.3
2	0.1522	0.0233	596.3
3	0.1983	0.0396	548.8
4	0.2348	0.0553	507.1

Layer 0 (and every flat arm) classifies the raw cap; layers 1–4 classify residuals that span over half the space. On the shipped arm, layer 0 was the only layer whose

heads left most entries unused (10.6 live entries, 216 bits, against ~ 375 for each later layer) — a within-model control on everything except the input.

What the cascade contributes: pairing. `blendablate.py` blends the subtracted term $\text{dec}(\mathbf{r}_\ell)$ in layer ℓ 's input toward a reference with the same marginals and none of the pairing — another row's reconstruction (shuffle), the mean, or Gaussian noise rescaled to the shuffle's per-sample magnitude — without retraining:

blend	shuffle	mean	noise	eff. dim L1 (shuffle)	(noise)
0.000	0.2292	0.2292	0.2292	499.1	499.1
0.125	0.2460	0.2374	0.2426	476.0	504.7
0.250	0.3084	0.2652	0.2838	394.2	518.3
0.500	0.6910	0.4100	0.4665	193.4	539.6
0.750	1.8154	0.7497	0.8502	103.0	527.6
1.000	4.6613	1.4302	1.5992	68.5	496.2

The pairing is load-bearing well below full severance — a quarter blend costs 35% of FVU_w — and another row's reconstruction does more damage than matched-magnitude noise at every level, so the damage is specific rather than generic perturbation sensitivity. The sweep cannot separate pairing from conditioning: effective dimension falls in lockstep with the error, because a residual is small and nearly isotropic *because* it is this sample's own error. That resolves the whitened arm rather than conflicting with it: whitening is conditioning without pairing, and it bought uniform usage of every entry and $\sim 7\%$ of FVU_w . (`forward="labels"`, which would test this by removing the residual entirely, has not been run at the fixed initialization. Forwarding the code *alongside* the residual, `forward="resid_labels"`, was a clear negative in a three-epoch pilot: held-out FVU_w 0.826 against 0.249, after leading early by $50\times$ — the signature of a parameter advantage without an information advantage, since the residual is by construction what the code failed to explain.)

The code is not the whole channel. Each layer splits its input into a unit direction (all the classifier sees) and a norm that multiplies the layer's output, so one real number per layer per sample reaches the decoder outside the bottleneck. Pinning each gain to its held-out mean (`gainablate.py`) costs the stack +2.1% in total (layer o 's gain is identically 1, since SONAR is unit-norm), against a 68% gap: the channel is real and is not the explanation. It does matter as a confound for the whitened flat arm, whose whitener creates a varying layer- o gain worth +4.5% — about 40% of whitening's apparent benefit.

The decoder is dense, and not because of the offset. Magnitude-pruning the trained decoder's smallest half takes FVU_w into the thousands, and its weight

distribution looks like a Gaussian matrix's. The non-negative dictionary makes the accumulator's constant part dwarf its varying one ($\|\mathbb{E}\mathbf{r}\|/\mathbb{E}\|\mathbf{r} - \mathbb{E}\mathbf{r}\| = 11.6$), and the bias-free decoder does attenuate that constant to the embedding mean (norm 0.581 against 0.564). But a decoder bias (–biased-dec) is worth only –0.62% of training MSE at matched step (the run died at 23% of its schedule, so its held-out 0.2464 is not comparable to the converged arm's): pinning $W_{\text{dec}}\mathbb{E}\mathbf{r}$ is 1024 linear constraints on a 1024×2048 matrix, 0.05% of its freedom. The decoder is dense because a full-rank map carrying a ~ 900 -dimensional signal is dense.

E.4 STEERING BY TARGET LAYER, AND HEAD-HEALTH LEVERS

Splitting the stack's steering by target layer (strength 1.0, 12 entries per layer) reverses the aggregate ranking of directions:

layer	decode	grad	margin	adjoint	random
0	0.216	0.583	0.547	0.237	0.034
1	0.464	0.089	0.367	0.000	0.023
2	0.742	0.130	0.216	0.000	0.055
3	0.966	0.130	0.091	0.000	0.034
4	0.362	0.047	0.036	0.005	0.047

Where the input reaches the classifier directly (layer 0), the classifier's own geometry (grad, margin) beats decode three- to tenfold. With depth it fails: a direction computed against layer ℓ 's classifier input is only valid there, and the perturbation arrives after earlier layers have re-encoded it. The decode direction moves the input toward what the model itself emits with the entry active, which is self-consistent down the stack. Layer 4 is hard for everything: its residual is nearly exhausted. The practical reading is to steer layer 0 by gradient and the middle layers by decode.

Head-health levers and steering. Three levers were tried on the shipped stack. *Mean-entropy pressure is not what costs steerability*: a twin trained at $s_{\text{KL}_m} = 10^{-6}$ has raw text-cycle effect 0.039 against 0.029, both an order of magnitude below top2/top8 at the same weight, so hard selection, not the bonus, is the cost (removing the bonus does cost head usage: imbalance 5.00 bits against 2.71). *Head-independence pressure (HSIC/CKA) was chasing estimator bias* (Section E.6). *A row-collinearity setpoint is the one lever that bought steerability*: holding within-head row cosine at 0.5 (against an uncontrolled 0.774) improves embedding-space realization in every direction at lower collateral — decode +15%, grad +17%, margin +52%, oriented adjoint +122%, collateral 0.591 $\rightarrow$ 0.552 — with held-out FVU_w slightly better (0.2157 vs. 0.2293). The controller acts once (Figure 26): as the max row cosine over heads and layers climbs past 0.8, the multiplier rises to 4.9×10^{-3} near step 35k and pulls it down to the setpoint by 50k; from about 60k the multiplier is projected to zero




FIGURE 26 · The row-collinearity setpoint on the shipped stack, from the training log. (a) The max within-head row cosine over heads and layers, the statistic the controller watches, against the uncontrolled twin; the dotted line is the 0.5 setpoint. (b) The applied multiplier $s_{kcossim}$, averaged over 1000-step windows.

and the statistic stays at or below 0.5 without pressure, ending at 0.473 against the uncontrolled 0.774. Its mechanism is unexplained: the statistic is computed in the e -dimensional dictionary space, where 74–78% of each atom’s energy is in the decoder’s null space (Section G.1), and decoded the atoms were already orthogonal. That the classifier-geometry directions gain most suggests an effect on the classifier rather than the dictionary, but nothing establishes it. The setpoint does not transport to the flat arm, whose steering is already saturated (0.992 vs. 0.966).

E.5 SUPERVISED VECTORS AT THEIR OWN STEP SIZE

Table 3 compares directions at an equal step. A real intervention has its own length: for a model, the output change of forcing the entry; for an SAE latent, its decoder direction at its mean activation; for the difference of means, the difference vector itself, which is the step activation addition takes. At those lengths (realized / collateral, median step as a fraction of $\|\mathbf{x}\|$, and each equal-length random step’s collateral):

model	own (step)	diff. of means (step)	random: own / dm
softmax stack	0.19 / 0.25 (0.064)	0.18 / 0.25 (0.079)	0.25 / 0.30
hard-code stack	0.08 / 0.54 (0.129)	0.27 / 0.47 (0.064)	0.53 / 0.48
hard-code flat	0.17 / 0.04 (0.028)	0.54 / 0.12 (0.087)	0.05 / 0.17
g160top1	0.41 / 0.01 (0.186)	0.57 / 0.04 (0.053)	0.18 / 0.07
g160softmax	0.54 / 0.07 (0.050)	0.57 / 0.13 (0.107)	0.08 / 0.16
TopK SAE, $k_{\text{SAE}}=32$ ($m=5120$)	0.46 / 0.06 (0.072)	0.77 / 0.15 (0.180)	0.07 / 0.17
TopK SAE, $k_{\text{SAE}}=32$ ($m=11264$)	0.50 / 0.08 (0.076)	0.82 / 0.17 (0.172)	0.09 / 0.19
SAE, bilinear encoder	0.82 / 0.04 (0.078)	0.88 / 0.08 (0.121)	0.04 / 0.07

The difference of means realizes more in most models, and in the SAEs and the flat arm partly by taking 2–3 times longer steps. The hard-code stack is the clear case where direction, not length, decides: the supervised vector reaches its entries three times as often with half the step. The softmax stack’s own intervention and the supervised vector tie. g160top1’s own intervention is the most precise step any method takes, disturbing 1.4% of other heads against 18% for random.

E.6 HEAD INDEPENDENCE: A MEASUREMENT ARTIFACT

A head-independence penalty (`-s-hsic-heads`, pairwise head CKA [32]) appeared to find substantial dependence. It was measuring bias:

layer	0	1	2	3	4	mean
unbiased	0.0041	0.0005	0.0005	0.0004	0.0002	0.0011
biased	0.0315	0.1062	0.1018	0.0963	0.0976	0.0867

The biased estimator reads chance co-occurrence between independent categorical heads as dependence, with a floor that grows as the batch shrinks (0.109 at $b=256$, 0.029 at 1024, against an unbiased 0.0001 at every size). At the training batch the floor exceeded any real dependence, so the penalty had nothing to descend but the bias. `hsic.pairwise_head_cka` now defaults to the unbiased estimator [33], which reads 0.06 at 0.5 planted factor correlation and 0.41 at 0.8. Conditioning on language or script explains none of the residue. **This architecture produces near-independent heads with no pressure to do so**; trained out under the corrected estimator the penalty changes nothing and costs localization (best-head language NMI 0.199 against 0.269). Two lessons generalize: before trusting a statistic as a training signal, measure it on a null it should score zero on; and loss-value share is not gradient share — at $s_{\text{HSIC}}=10^{-4}$ the penalty was 12% of the loss value, while the measured MSE-to-CKA gradient ratio (1250 at the classifier, 221 at the dictionary) put equal pull at 8×10^{-4} .

E.7 SEED REPRODUCIBILITY, MEASURE BY MEASURE

Partitions. The best-matched pair of heads reaches NMI 0.42, but the mean is 0.012 against a 0.005 null, and `splitting.py`'s entry containment reads 0.000. Neither is a metric artifact: comparing one run against its own checkpoint 19,500 steps earlier, the same partition measure scores 0.38 overall and 0.87 on layer 0. *Contributions.* Two heads could carve differently and add the same thing; they do not (centered cosine 0.0074 vs. 0.0015 null; uncentered it reads 0.205, which is the non-negative dictionary's shared offset). `headcontrib.py`, which fixes the measurement (decoded, whitened, centered contributions, Hungarian-matched), reads 0.019 against a 0.0017 null on the same pair at 4,096 rows; the earlier one-off figure's exact method was not recorded, and the conclusion is unchanged under either. *Subspaces.* Top-8 eigenspaces of per-layer contributions agree 25–33× above chance — but each model agrees with the *cache's own* principal directions by the same amount (e.g. layer 1: 123.7× between seeds, 101.3× and 101.2× to the data), so the agreement is the data's and nothing is left for the architecture. At layer 4 each seed tracks the data ~23× above chance while agreeing with the other only 9×: they pick *different* data-aligned directions. *Grams.* Per-layer row collinearity agrees to four digits (0.2880 / 0.2891 at layer 0), but the Gram spectrum is nearly one-parameter — the top eigenvalue is exactly $1 + (k - 1) \kappa_i$, carries 82% of the spectrum's norm, and is forced by non-negativity. Compared entrywise in canonical order, Grams agree at chance across seeds (1.05–1.11×) and *equally* at chance against the same run 19,500 steps earlier (1.07–1.15×): a head's Gram is close to a fixed property of the configuration rather than something learned. Decoding atoms to output space first changes nothing (1.18× at layer 0 to 1.03× at layer 4).

E.8 WHAT MAKES LAYER 0 REPRODUCE: GAIN, ATOMS, AND DEPTH

Layer-0 heads are the largest and the most reproducible, so “layer 0 reproduces” might be “large heads reproduce”. `headcontrib.py` measures a head's size as its centered contribution energy over the whitened target variance, and splits it exactly: a head's contribution on a row is its layer's input gain times the decoded atom it selected, so its size is gain^2 times atom spread^2 (the usage-weighted RMS distance of its decoded atoms from their mean); on both stacks the product reproduces the measured size to a log correlation of 1.000.

Layer 0's larger contributions are all gain. It receives the raw input — GPT-2 activations of norm ~122 against residuals of ~60; unit SONAR embeddings against residuals of 0.36–0.59 — while its atoms are no larger than layer 1's on either dataset, and on SONAR smaller than those of layers 1–3 and equal to layer 4's. Gain is one number per layer per row, shared by all of a layer's heads, so across layers “size” is the layer itself under another name. At equal atom spread, layer 0 still reproduces 14× better on GPT-2 and 23× better on SONAR (log agreement on

layer	GPT-2 concat			SONAR stack		
	agreement	gain	atom norm (10^{-4})	agreement	gain	atom norm
0	0.171	122.2	4.6	0.080	1.000	0.043
1	0.017	67.9	5.3	0.005	0.588	0.054
2	0.007	62.3	4.4	0.004	0.490	0.052
3	0.005	59.2	3.7	0.003	0.417	0.051
4	0.004	57.2	3.2	0.003	0.359	0.043

TABLE 18 · Head agreement (matched centered contribution cosine) and the two factors of head size, medians over each layer’s heads; atoms decoded and whitened. SONAR: ste_h76_init01 against ste_h76_i01_s43, 4,096 rows.

log spread and a layer-0 indicator, coefficient of determination 0.979 and 0.941).

Against the flat arm at matched size. Within a stack, head size and layer do not overlap, so their effects cannot be separated there. The flat arm breaks the confound: all its heads see the raw input, with gain 1 and nothing downstream. On SONAR its own agreement tracks log head size at correlation 0.876 — from 0.008 for its smallest heads (0.057–0.060% of target variance) to 0.150 above 0.071%, and 0.78 and 0.56 for its two largest — so size matters a great deal on its own, and since its gain is constant it is an atom-magnitude curve. But the stack’s heads fall far below that curve at every depth, layer 0 included:

stack layer	size, median	stack agreement	flat heads at that size ($\pm 10\%$)	gap
0	0.31%	0.080	0.67 (2 heads)	~8×
1	0.16%	0.005	0.12 (1 head)	~23×
2	0.11%	0.004	0.19 (3 heads)	~49×
3	0.073%	0.003	0.059 (115 heads)	~19×
4	0.039%	0.003	below the flat range	—

The stack costs reproducibility at matched size, at every layer. Layer 0 sees exactly the flat arm’s input and still reproduces several times worse than an equal-size flat head: in the stack, what a layer-0 head should carve depends on what the layers after it do with its error, so the input alone no longer pins it. Deeper layers lose more. This is the reproducibility face of the entanglement result (Section 5.4): the cascade that doubles a single-head intervention’s collateral, to 66% of the other heads, also leaves each head underdetermined by the data. Read against the flat curve, the gain difference between layers 0 and 3 accounts for roughly a third of layer 0’s lead in log terms; only layer 3’s size match is solid — flat heads are almost all small, so layers 0–2 rest on one to three heads each — so the fraction is rough and the direction firm. The matched-size comparison is SONAR-only; GPT-2 has no flat arm.

Within layer 0 the atoms decide. SONAR’s layer-0 gain is identically 1, so its within-layer size variation is pure atom magnitude, and there log agreement tracks atom norm at correlation 0.62 and spread at 0.64 (cosine against log size, 0.91); GPT-2’s layer 0 shows 0.17 and 0.42. Past layer 1 the atoms predict nothing on either dataset. Spread and norm are one quantity here — their ratio is 0.99–1.08 in every layer, so a head’s usage-weighted mean decoded atom is near zero: the non-negative dictionary’s offset is cancelled by the decoder before it reaches the output.

E.9 HOW ENTRIES SHARE A LAYER

An entry writes only where it wins its head, and then writes its layer’s gain times its decoded atom, so its energy over the data is $\mathbb{E}[1\{\text{selected}\} \text{gain}^2] \|\text{atom}\|^2$. Per layer, over 32,768 held-out rows (entryshare.py):

run	layer	Gini	top 10% carry	max/ median	from atom norm	within heads
GPT-2 concat, $k=256$	0	0.587	42%	38.7	76%	60%
	1	0.155	16%	30.3	44%	100%
	2–4	0.11–0.12	14%	~2	~40%	100%
SONAR stack, $k=32$	0	0.326	32%	46.3	92%	83%
	1–3	0.06–0.10	12–14%	1.5–1.9	16–24%	~100%
	4	0.195	18%	3.2	77%	87%
SONAR flat, 1×380	—	0.190	23%	64.1	90%	76%

TABLE 19 · Inequality of entry energy within a layer. “From atom norm” is the covariance share of the variance of log entry energy (the remainder is usage; gain conditional on selection contributes ~0). “Within heads” is the share of that variance not explained by which head an entry is in. No entry is dead in any layer.

Layers 1–3 are near-equal; layer 0 is not, and its atoms are why. Deep entries have an effective number of 95–99.5% of the total and the largest is about twice the median; what inequality they have is mostly usage, itself mild (CV 0.14–0.18). Layer 0’s usage is just as even — on SONAR every entry fires 2.1–4.5% of the time against $1/k = 3.1\%$ — but its atom norms vary five to ten times more (CV 0.34–0.52 against 0.02–0.10), mostly inside heads. The flat arm shows the same pattern more mildly, so it belongs to classifying the raw input. SONAR’s last layer is unequal for the other reason: some atoms have shrunk almost to nothing (smallest/median 0.017), the inert final layer of Section G.4. Whole layers are unequal by gain: layer 0 holds 67% of entry energy on GPT-2 and 49% on SONAR, layer 4 4–5%.

Where the large atoms sit differs completely between the models (bigatoms.py):

	GPT-2 concat, 20 heads	SONAR stack, 76 heads
<i>within a head</i> (atom norm / the head's own median, mean over heads)		
largest	3.7 (1.4–5.1 across heads)	2.0 (1.5–4.3 across heads)
2nd / 8th / 16th largest	3.2 / 2.0 / 1.7	1.7 / 1.2 / 1.0
largest entry's share of head	6.1% (uniform 0.4%)	9.9% (uniform 3.1%)
<i>energy</i>		
<i>across the layer</i> (atom norm / the layer median)		
atoms > 2×	1030 (20%), in all 20 heads	63 (2.6%), in 28 heads
atoms > 3×	58, in all 20 heads (at most 15 each)	11, in 4 heads
atoms > 5×	1	5, all in head 54
largest head's energy share	15.5% (3× uniform)	9.2% (7× uniform)
<i>usage</i>		
Spearman(atom norm, usage) in heads	+0.81	+0.27 (−0.04 to +0.58)
usage of atoms > 3× / of the rest	0.0062 / 0.0037 (1/k 0.0039)	0.033 / 0.031 (1/k 0.031)
<i>direction</i>		
cos between heads' largest atoms	0.022 (random pairs 0.048)	0.023 (random pairs 0.028)
energy in the data's top 10 PCs	4% (all atoms 13%)	9% (all atoms 3%)

On GPT-2 every head has a graded tail — its eighth-largest atom is still twice its median — and the large atoms are its frequent outcomes (1.6× uniform usage). Agreement is fairly even across heads (0.135–0.235) and only loosely tied to atoms (Spearman 0.49 with energy share, 0.34 with the largest atom). On SONAR the large atoms are a few heads' property, and fire at the uniform rate: large, not frequent. Head 54 holds 8 of the 11 atoms above 3× and all 5 above 5×, carries 9% of the layer, and reproduces at **0.844** against 0.249 for the next head; across heads, agreement tracks the largest atom (Spearman 0.62) and how peaked the head's profile is (0.62), not its typical atom (−0.14). That resolves the within-layer-0 atom effect into a few heads with outsized atoms — but not the layer-0 lift: without head 54 the layer-0 mean falls only from 0.080 to 0.070, still twenty times the deep layers. Head 54 is the topic/genre head of Section 5.6, which recurs across seeds.

The two models differ in kind here. On SONAR one head stands out on every atom measure and alone reproduces strongly; the rest of layer 0 is a long, nearly flat tail, and agreement follows how peaked a head's atoms are. On GPT-2 no head stands out: every head has several large atoms, the spread of agreement is narrow, and the head with the largest share of the layer's energy (17) has a flat profile — uniformly large atoms rather than a few outsized ones — yet reproduces no better than its neighbors. Peakedness predicts nothing there (+0.09).

The large atoms are distinct directions: different heads' largest atoms are near orthogonal (median |cos| 0.022–0.023, below random atom pairs). On SONAR they

head	agreement	energy share	largest / layer median	atoms > 3×	largest / head median
<i>SONAR stack, layer 0</i> (uniform share 1.3%)					
54	0.844	9.2%	6.81	8	3.41
46	0.249	2.5%	3.21	1	2.24
6	0.181	1.5%	2.76	0	2.83
18	0.150	1.9%	4.50	1	4.34
lowest four	0.008–0.021	1.0–1.1%	1.51–1.84	0	1.48–1.84
<i>GPT-2 concat, layer 0</i> (uniform share 5%)					
6	0.235	3.1%	4.31	2	4.77
17	0.220	15.5%	4.99	5	1.90
19	0.215	3.4%	3.82	3	4.16
0	0.202	3.1%	3.95	1	4.19
lowest four	0.135–0.141	2.6–3.0%	3.13–4.17	1–2	3.45–4.56

TABLE 20 · Layer-0 heads by seed agreement (`bigatoms.py`; matched centered contribution cosine against the seed-43 twin). Spearman of agreement across heads, SONAR / GPT-2: with energy share +0.53 / +0.49; with the largest atom +0.62 / +0.34; with the head’s peakedness (largest / head median) +0.62 / +0.09; with its median atom −0.14 / +0.16.

lean into the data’s top ten principal directions (9% of their energy against 3% for a typical atom), consistent with cross-seed agreement being the data’s dominant directions; on GPT-2 they lean the other way (4% against 13%). One seed pair per dataset throughout.

E.10 LANGUAGE: THE PROBE IN DETAIL

Held-out language identity over 86 languages (dense ridge ceiling 0.721, chance 0.012):

As expected of a language-agnostic embedding, no head encodes language outright: the best reaches 0.18 of an attainable 0.875, and no head-health lever lifts any arm past 0.29 (a $k=32$ head could not hold 86 languages anyway, $\log_2 32 = 5 < 6.43$ bits). What the probe does show is that partitions carry a known factor well above chance, and that **the hard code carries more of it**: four times the softmax reference on null-corrected best-head NMI, and 0.42–0.43 against 0.28 in a 32-head probe, consistently across seeds. The fixed initialization spreads language over more heads than the unfixed one did (best head 0.17 against 0.26, 32-head probe 0.42 against 0.38). The stack beats the flat arm on probe accuracy at every head budget (0.378 vs. 0.257 at $h'=32$, both unfixed). (`summary.json`’s `info_share_best` favors the flat arm, but it divides by summed plug-in MI whose per-head bias differs eightfold between the shapes; do not read it across architectures.)

arm	shape	best-head NMI	less null	probe $h'=1$	$h'=32$
ste_h76_init01	5×76	0.180	0.170	0.103	0.423
seed 43	5×76	0.165	0.155	0.089	0.428
softmax reference	5×32	0.055	0.043	0.036	0.275
<i>unfixed initialization</i>					
ste_h76	5×76	0.269	0.261	0.165	0.378
+ collinearity setpoint	5×76	0.287	0.279	0.139	0.318
– mean-entropy bonus	5×76	0.085	0.078	0.038	0.322
+ head independence	5×76	0.199	0.189	0.097	0.383
flat	1×380	0.277	0.276	0.102	0.257
+ collinearity setpoint	1×380	0.247	0.246	0.080	0.271

TABLE 21 · headlang.py. The attainable NMI for a 5-bit head against the 6.43-bit label is 0.875. Head-health levers and the flat arm were measured only at the unfixed initialization.

Script, cells, conjunctions. Script behaves the same way (best-head NMI 0.191 of a 0.671 ceiling; the same head, Lo h50, is best for both). A single cell is a strong but permissive detector: the best cell per language reaches recall 0.74–0.98 at FPR 1.6–2.2%, but no language reaches precision 0.8. Conjunctions of top cells from distinct heads do not fix it: F1 peaks at $h'=2$ conjoined cells, 7% above a single cell, and the best precision is 0.598 at 4.8% recall. Recall combines nearly independently (0.234 against 0.262 predicted) but false positives do not (precision 0.43 at $h'=2$ against 0.83 if independent): many heads send the same confusable rows to the same wrong language — Finnish with Estonian, Japanese with Korean. The code encodes language *neighborhood* rather than identity.

Why neighborhoods, from the embedding geometry. The language-agnostic training leaves exactly the residue this pattern implies. Language explains 6.51% of embedding variance and script 2.26%; corrected for the number of classes (headeta.py) language is still a stronger partition than 90% of the heads ($51\times$ its random baseline, against $5\times$ for the median head and $235\times$ for the best), but its languages are weakly separated regions, not directions: within-language mean cosine is 0.3596 against 0.3160 between, and the shared mean direction accounts for 0.3178 of that, leaving a +0.044 excess, 0.30 sd of the pair distribution. A partition code cuts along local residual variance, and at 0.30 sd those cuts pass through the language clusters rather than around them. Layer-4 heads are near-random partitions ($2\times$ baseline), which is Section G.4’s inert final layer seen from a different direction.

The topic head’s entries. Decoding and autointerp were run on the shipped-initialization run’s head 54, the head that holds the same index in the fixed-

initialization run (Section 5.6). Decoded as deviations from the head’s mean atom, its entries read as topic and genre. Contrastive autointerp (null 0.273) groups them as genre or topic (16 entries, F1 0.523), multilingual or code-switched passages (8, 0.476), text artifacts such as duplication and mojibake (3, 0.395), and script or region (5, 0.394); all four beat the null.

F SUPPLEMENTARY RESULTS: GPT-2 HARD CODES

F.1 WHAT EACH DESIGN CHANGE BUYS

One-epoch runs (10M tokens) isolate them:

variant (800 bits unless noted)	FVU _w	loss recovered
STE alone (latent dictionary + decoder)	0.327	93.3%
+ gain	0.290	94.3%
+ gain, private heads	0.288	94.2%
+ gain, direct atoms	0.217	96.95%
seed 43	0.219	96.85%
$k=64$ (960 bits)	0.186	97.6%
$k=128$ (1120 bits)	0.165	97.9%

Gain-shape and direct atoms each help, and direct atoms most. Their advantage is partly speed: over 40M tokens private heads close the gap and finish slightly ahead (0.186 against 0.195). Without direct atoms and gain, more entries per head collapse rather than help (FVU_w 0.75 and 0.80 at $k=64$ and 128), and the other straight-through configurations froze their last layer at 4.8–5.0 bits, every head constant; the direct code’s deepest layer never froze. The base result reproduces across seeds to 0.002 FVU_w. Annealing matters: straight-through at a constant $t=1$ reaches 0.50 and a scheduled soft-to-hard mix 0.39. Sparse encoders in front of the classifier never beat the dense one: at one epoch, 32- and 640-wide front-ends reach 0.25–0.29 against the dense classifier’s 0.217, and a 4096-wide, 64-active encoder per head with dead-latent resampling reaches 0.225 against 0.195 for the dense front-end on the same 40M tokens, with 88% of its layer-1 latents dead at no cost to that layer.

The 40M-token runs, layer by layer. Every layer contributes: prefix FVU_w runs $0.472 \rightarrow 0.347 \rightarrow 0.278 \rightarrow 0.231 \rightarrow 0.195$ for the direct code and $0.451 \rightarrow 0.186$ for private heads. Entry usage stays near uniform in every layer (imbalance 0.01–0.16 bits of 5, and 0.30 at the private code’s last layer), and the observed gain is close to the best per-sample scale for each layer’s contribution (optimal rescale 0.89–1.09, coefficient of variation 0.05–0.08). The runs use two options this branch does not implement, `resid_first` and `gain_clip`, so they reproduce only on the headline branch; its direct atoms, private heads (this branch’s `ConcatDictBlock`), fibers and head dropout have been ported.

G SUPPLEMENTARY RESULTS: DICTIONARY GEOMETRY

G.1 ROW COLLINEARITY IS MOSTLY GAUGE

Within-head row cosine κ_i (`cossim_k`, `rowcos`; $\bar{\kappa}$ its mean over heads) is computed on `dicts()` in the e -dimensional dictionary space. The decoder is $d \times e$ at full rank, so $e - d$ of those dimensions are its null space, and on the SONAR straight-through arm atoms put 74–78% of their energy there, where nothing they do reaches the output and no gradient reaches them. Decoded into output space the same statistic reads 0.023 at layer 0 and -0.015 to -0.020 below it, against 0.12–0.29 in e . **Dictionary statistics should be computed after decoding.**

On GPT-2 layer 8 ($e=1536$, $d=768$) four tempting readings of the null-space mass each failed a check: head null components add coherently (3.3–11 $\times$ over the cancelling expectation), so they are not interference debris; the fraction does not track feature noise; the per-atom visible fraction is flat at 22% across training and the decoder’s row space rotates only 2.6–3.5°, so it is not a swap space; and a mixed-sign visible component needs no non-negative “lift”. Two baselines also had to be corrected: for $|\mathcal{N}(0, 1)|$ atoms the chance visible fraction is 18.26% (not 50%) and the chance alignment with **1** is 63.7% (not ~ 0), so the trained dictionary is *less* null-heavy than chance; and a reported 49% of reconstruction gradient in $\ker(W_{\text{dec}})$ was a coordinate artifact, 3e-8 with respect to the atoms $|D|$. How large “the null part” is depends entirely on the norm asked for: 2.43% of squared L2, 23.20% of L2, 56.55% of L1.

A random decoder’s cone. Wendel’s theorem [34] gives a geometric reading of the width at which non-negativity stops binding: the decoder’s e random columns lie in a common halfspace with probability $2^{1-e} \sum_{q < d} \binom{e-1}{q}$, exactly 1/2 at $e=2d$, so on GPT-2 $e=1536$ sits on the threshold at which a random decoder’s cone $\{W_{\text{dec}} \mathbf{r} : \mathbf{r} \geq 0\}$ covers the output space at all.

G.2 HEADS SHARE COORDINATES: NO PARTITION

Section 3.1.4’s between-head orthogonality — non-negative heads that are orthogonal own disjoint coordinate blocks of the dictionary space, as a property of the weights — is the premise for reading a dictionary statically. Every model here trained before the support overlap ω existed, most with its predecessor penalized, so it had never been measured. Scored on final checkpoints (`headsupport.py`; one seed per row):

model	ω , dictionary (null)	spread	ω , decoded (null)
SONAR resid_nc (372.6k steps)	0.818 (0.815)	0.83	0.917 (0.910)
SONAR resid_nc_hm (8.27M steps)	0.113 (0.118)	0.13	0.892 (0.878)
SONAR softmax_shm	0.731 (0.738)	0.74	0.938 (0.930)
SONAR top2_shm	0.862 (0.860)	0.87	0.972 (0.958)
SONAR straight-through, shipped init	0.942 (0.942)	0.94	0.721 (0.719)
SONAR straight-through, 5×76	0.962 (0.961)	0.96	0.980 (0.965)
SONAR straight-through, 1×380	0.977 (0.977)	0.98	0.974 (0.963)
GPT-2 abs, $e=1536$	0.739 (0.643)	0.65	0.948 (0.925)
GPT-2 abs, $e=768$	0.596 (0.486)	0.52	0.875 (0.852)
GPT-2 signed, $e=1536$	0.937 (0.636)	0.64	0.964 (0.943)
GPT-2 signed, $e=768$	0.955 (0.950)	0.95	0.965 (0.943)
GPT-2 concat, $d_{\text{head}}=32$	o by construction	—	0.846 (0.788)
g160top1 (160 trained heads)	—	—	0.648 (0.644)
g160softmax (160 trained heads)	—	—	0.983 (0.975)
k32 discovered groups (359)	—	—	0.899 (0.892)

No model partitions its dictionary space between heads. ω sits at or above its null everywhere, and each head’s profile covers 52–98% of the coordinates where a partition into 32 heads needs about 3%. Where ω departs from the null it is *above* it: the GPT-2 arms place their heads on the same coordinates, the signed wide arm most. The one exception is resid_nc_hm, the reference configuration trained 22 times longer: its profiles are sparse and ω falls 4% below the null in four of five layers. That is long training sparsifying each head, not a partition — at spread 0.13 a head still uses four times its share — and the same configuration at 372.6k steps is dense and at its null. The predecessor penalty ($s_{\text{cossim}_h}=10^{-6}$, or 10^{-5} for resid_nc) did not produce the partition; no run has trained with ω itself as a loss term. ConcatDictBlock is the only route to disjoint support here, and it is disjoint by construction.

It is not the decoder’s null space. Splitting each atom into its row-space (visible) and null-space parts and scoring each against the same null, the visible part is at chance on SONAR and above it on GPT-2 (the signed wide arm 0.955 against 0.706), and the $e=768$ arms, whose square decoder has no null space, sit at or above the null in the dictionary space itself. The row space is not coordinate-aligned, so the visible part’s support is basis-dependent; the decoded ω is the nearest basis-free check, and it agrees. The decoded measure itself cannot separate architectures: every model’s atoms are dense across SONAR’s output coordinates, and all fifteen rows land within 7% of their nulls, the trained and discovered SAE groups included. The comparison ω was built for exists only in the dictionary space, which the SAE variants do not have.

G.3 CONCATENATED HEADS

ConcatDictBlock at $d_{\text{head}}=32$ gives each head a (768, 32) decoder block: 32 non-negative generators in a 32-dimensional subspace, a simplicial cone, and the per-head null space is zero. The floor on a head's mean row cosine is $(k/d_{\text{head}} - 1)/(k - 1)$, which is 0 at $d_{\text{head}}=k=32$, so concat faces no raised floor; it faces a collapsed tradeoff, since $\kappa_i=0$ at $d_{\text{head}}=k$ needs one coordinate per atom. It took the correlation, the lowest effective *and* numerical rank (rank99 17.1 of 32, against 30.2 summing), and the worst error (Table 7). Its extra cost is reachability, the one place the cone frame binds; it is not explained by $\bar{\kappa}$ alone (at step 90k it had slightly lower $\bar{\kappa}$ than $\text{abs}/768$, 0.571 against 0.578, and much worse held-out error, 0.806 against 0.463).

G.4 SELECTION RULE, TRAINING DURATION, AND THE LAST LAYER

Softer codes do not raise the dictionary's rank. On the SONAR selection sweep, numerical rank is full (rank99 30.7–31.8 of 32) in all ten arms, and effective rank varies 7-fold but is not monotone in softness, with the two families ($\pm$ mean-entropy bonus) ordering oppositely at convergence. Tracing each arm over training resolves it: **training duration dominates the selection rule by 3.8×** (mean $\bar{\kappa}$ falls 0.750 $\rightarrow$ 0.217 from step 10k to 369.5k, while the spread across rules at a fixed step averages 0.140), and the orderings cross near step 100k (Figure 27b, the `_shm` family). The soft arm's $\bar{\kappa}$ plateaus while its error keeps falling, and the hard arms do the reverse: from 150k to 369.5k softmax cuts its FVU_w by 85% while $\bar{\kappa}$ moves 3%, whereas top8 moves FVU_w 3% while cutting $\bar{\kappa}$ by 45%, top16's FVU_w does not improve while $\bar{\kappa}$ falls 29%, and top2 cuts both, 12% and 34% (Figure 27). Under a hard code a head's output is one of its rows, so collinear rows make the output independent of which entry wins; collinearity costs a hard code information directly, and a soft code, with continuous coefficients, much less. At convergence the softer code has the *more* collinear dictionary. Matching arms on held-out FVU_w instead of step reverses that at moderate error: softmax has the lowest $\bar{\kappa}$ at FVU_w 0.60 and 0.40 (0.60 and 0.58, against 0.72–0.78 for the hard arms). Only at 0.25 does top2 drop below it (0.53 against 0.56), and it reaches that level at step 80k against 21–36k for the others, so its lower $\bar{\kappa}$ there is again duration. One seed per arm; these runs trained on the evaluation tail, so their FVU_w is in-sample.

The last layer is inert once the error is small. Collinearity falls monotonically with depth in every hard arm, but soft arms spike at the last layer (softmax_shm layer 4: $\bar{\kappa}=0.799$, effective rank 1.54). `lastlayer.py` ablates the final layer's heads two ways: uniform (emit the mean atom) and zero (emit nothing):




FIGURE 27 · The selection sweep’s `_shm` arms over training: (a) FVU_w on the (in-sample) tail, scored from each checkpoint, and (b) mean within-head row cosine $\bar{c}$. The soft arm’s $\bar{c}$ flattens near 0.36 while its error keeps falling; the hard arms keep decorrelating.

arm	select	L4 $\bar{c}$	L4 eff	prefix FVU_w		
				before L4	L4 adds	zero-ablate
top1	top1	+0.246	10.60	0.7511	+4.3%	+4.5% worse
top2_shm	top2	+0.188	14.89	0.2378	+19.7%	+24.5% worse
top8_shm	top8	—	—	0.0304	+5.5%	+5.8% worse
top16_shm	top16	+0.433	3.46	0.0043	+0.0%	+0.0%
softmax	softmax	+0.263	6.00	0.0023	−1.2%	−1.1% better
softmax_shm	softmax	+0.799	1.54	0.0007	−1.9%	−1.8% better

The last layer’s contribution tracks the error left before it, not its collinearity: softmax’s layer 4 is not collapsed ($\bar{c}=0.263$) and is just as useless. Below a prefix FVU_w of about 0.005 the final layer is inert or mildly harmful — which accounts for `resid_nc`’s “mildly harmful” layer 4 (prefix 0.0051) without appeal to its frozen heads. The two collapses end differently: `top16_shm`’s layer 4 shrank its norms to 4% of the other layers’ and emits nothing; `softmax_shm`’s is collinear at near-equal magnitude, emits a near-constant vector, and removing it helps — a small overfit, not a useful bias.